

CS 61C:  
Great Ideas in Computer Architecture  
*More RISC-V Instructions and  
How to Implement Functions*

Instructors:  
Nick Weaver & John Wawrzynek  
<http://inst.eecs.Berkeley.edu/~cs61c/fa17>

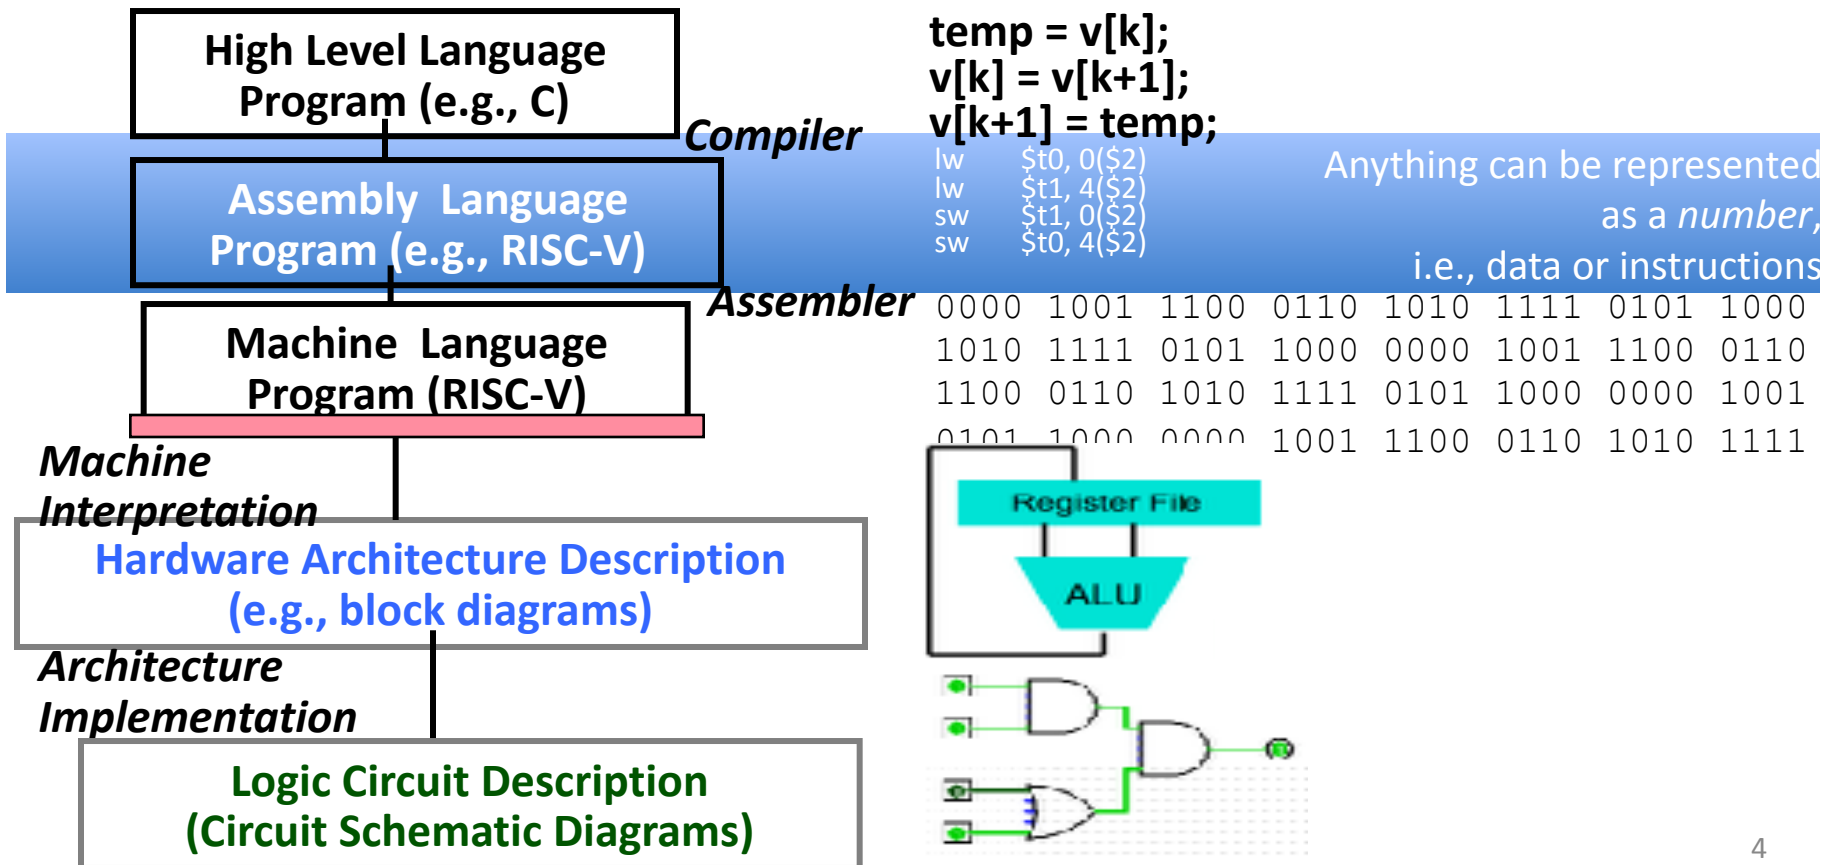
# Outline

- RISC-V ISA and C-to-RISC-V Review
- Program Execution Overview
- Function Call
- Function Call Example
- And in Conclusion ...

# Outline

- RISC-V ISA and C-to-RISC-V Review
- Program Execution Overview
- Function Call
- Function Call Example
- And in Conclusion ...

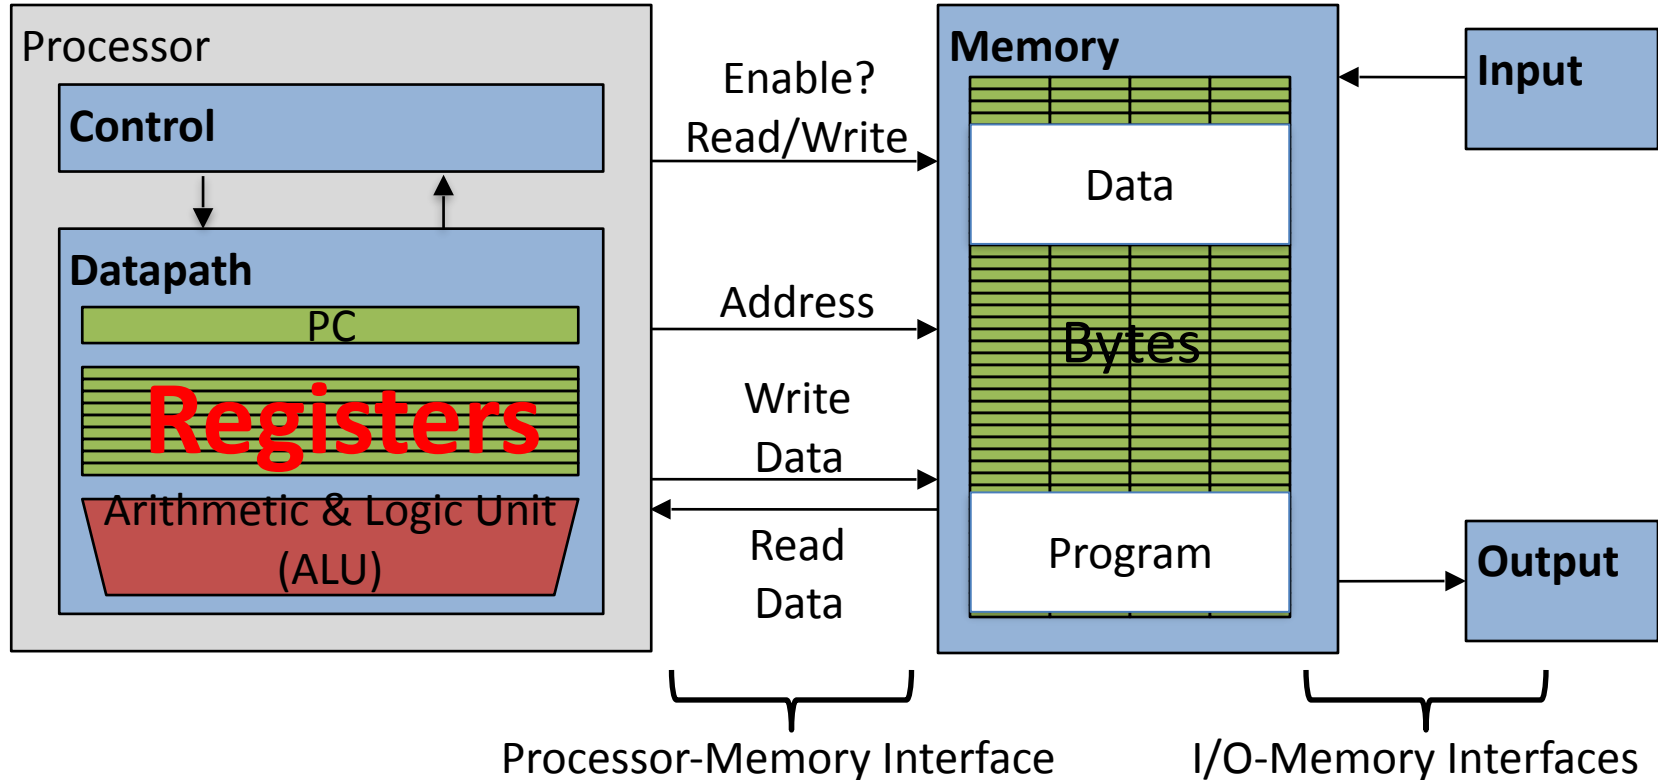
# Levels of Representation/Interpretation



# Review From Last Lecture ...

- Computer's native operations called *instructions*. The *instruction set* defines all the valid instructions.
- RISC-V is example RISC instruction set - used in CS61C
  - Lecture/problems use 32-bit RV32 ISA, book uses 64-bit RV64 ISA
- Rigid format: one operation, two source operands, one destination
  - add, sub
  - lw, sw, lb, sb to move data to/from registers from/to memory
- Simple mappings from arithmetic expressions, array access, in C to RISC-V instructions

# Recap: Registers live inside the Processor



# RISC-V Logical Instructions

- Useful to operate on fields of bits within a word
  - e.g., characters within a word (8 bits)
- Operations to pack /unpack bits into words
- Called *logical operations*

Logical operations	C operators	Java operators	RISC-V instructions
Bit-by-bit AND	&	&	<b>and</b>
Bit-by-bit OR			<b>or</b>
Bit-by-bit XOR	^	^	<b>xor</b>
Shift left logical	<<	<<	<b>sll</b>
Shift right logical	>>	>>	<b>srl</b>

# Logical Shifting

- Shift Left Logical: `slli x11,x12,2` # `x11 = x12<<2`
  - Store in x11 the value from x12 shifted 2 bits to the left (they fall off end), **inserting 0's** on right; `<<` in C

Before: `0000 0002`<sub>hex</sub>

`0000 0000 0000 0000 0000 0000 0000 0010`<sub>two</sub>

After: `0000 0008`<sub>hex</sub>

`0000 0000 0000 0000 0000 0000 0000 1000`<sub>two</sub>

What arithmetic effect does shift left have?

- Shift Right Logical: `srl` is opposite shift; `>>`
  - Zero bits inserted at left of word, right bits shifted off end




# Arithmetic Shifting

- *Shift right arithmetic (srai)* moves  $n$  bits to the right (insert high-order sign bit into empty bits)
- For example, if register x10 contained  
 $1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1110\ 0111_{\text{two}} = -25_{\text{ten}}$
- If execute `sra x10, x10, 4`, result is:  
 $1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1110_{\text{two}} = -2_{\text{ten}}$
- Unfortunately, this is NOT same as dividing by  $2^n$ 
  - Fails for odd negative numbers
  - C arithmetic semantics is that division should round towards 0

# Computer Decision Making

- Based on computation, do something different
- Normal operation on CPU is to execute instructions in sequence
- Need special instructions for programming languages: *if*-statement
- RISC-V: *if*-statement instruction is  
**beq register1, register2, L1**  
means: go to instruction labeled L1  
if (value in register1) == (value in register2)  
....otherwise, go to next instruction
- **beq** stands for *branch if equal*
- Other instruction: **bne** for *branch if not equal*



# Types of Branches

- **Branch** – change of control flow
- **Conditional Branch** – change control flow depending on outcome of comparison
  - branch *if* equal (**beq**) or branch *if not* equal (**bne**)
  - Also branch if less than (**blt**) and branch if greater than or equal (**bge**)
- **Unconditional Branch** – always branch
  - a RISC-V instruction for this: *jump (j)*

# Outline

- Assembly Language
- RISC-V Architecture
- Registers vs. Variables
- RISC-V Instructions
- **C-to-RISC-V Patterns**
- And in Conclusion ...

# Example *if* Statement

- Assuming assignments below, compile *if* block

$f \rightarrow x10$        $g \rightarrow x11$      $h \rightarrow x12$   
 $i \rightarrow x13$        $j \rightarrow x14$

<code>if (i == j)</code>	<code>bne x13,x14,skip</code>
<code>  f = g + h;</code>	<code>add x10,x11,x12</code>
<code>  .</code>	<code>skip: .</code>
<code>  .</code>	<code>  .</code>
<code>  .</code>	<code>  .</code>

# Example *if-else* Statement

- Assuming assignments below, compile

$f \rightarrow x10$      $g \rightarrow x11$      $h \rightarrow x12$      $i \rightarrow x13$      $j \rightarrow x14$

```
if (i == j)                                bne x13,x14,else
    f = g + h;                             add x10,x11,x12
else                                         j done
    f = g - h;                             else:  sub x10,x11,x12
                                           done:
```

# Magnitude Compares in RISC-V

- Until now, we've only tested equalities (`==` and `!=` in C); General programs need to test `<` and `>` as well.

- RISC-V magnitude-compare branches:

“Branch on Less Than”

Syntax: `blt reg1, reg2, label`

Meaning: `if (reg1 < reg2) // treat registers as signed integers  
goto label;`

- “Branch on Less Than Unsigned”

Syntax: `bltu reg1, reg2, label`

Meaning: `if (reg1 < reg2) // treat registers as unsigned integers  
goto label;`

“Branch on Greater Than or Equal” (and its unsigned version) also exists.

# C Loop Mapped to RISC-V Assembly

```
int A[20];  
int sum = 0;  
for (int i=0; i<20; i++)  
    sum += A[i];
```

```
# Assume x8 holds pointer to A  
# Assign x9=A, x10=sum,  
#           x11=i, x13=20  
add x9, x8, x0 # x9=&A[0]  
add x10, x0, x0 # sum=0  
add x11, x0, x0 # i=0  
addi x13,x0,20 # x13=20  
Loop:  
lw x12, 0(x9) # x12=A[i]  
add x10,x10,x12 # sum+=  
addi x9,x9,4 # &A[i++]  
addi x11,x11,1 # i++  
blt x11,x13,Loop
```



# Peer Instruction

Which of the following is TRUE?

**A:** `add x10, x11, 4(x12)` is valid in RV32

**B:** can byte address 8GB of memory with an RV32 word

**C:** `imm` must be multiple of 4 for `lw x10, imm(x10)` to be valid

**D:** None of the above

# Peer Instruction

Which of the following is TRUE?

A: `add x10, x11, 4 (x12)` is valid in RV32

B: can byte address 8GB of memory with an RV32 word

C: `imm` must be multiple of 4 for `lw x10, imm(x10)` to be valid

D: None of the above

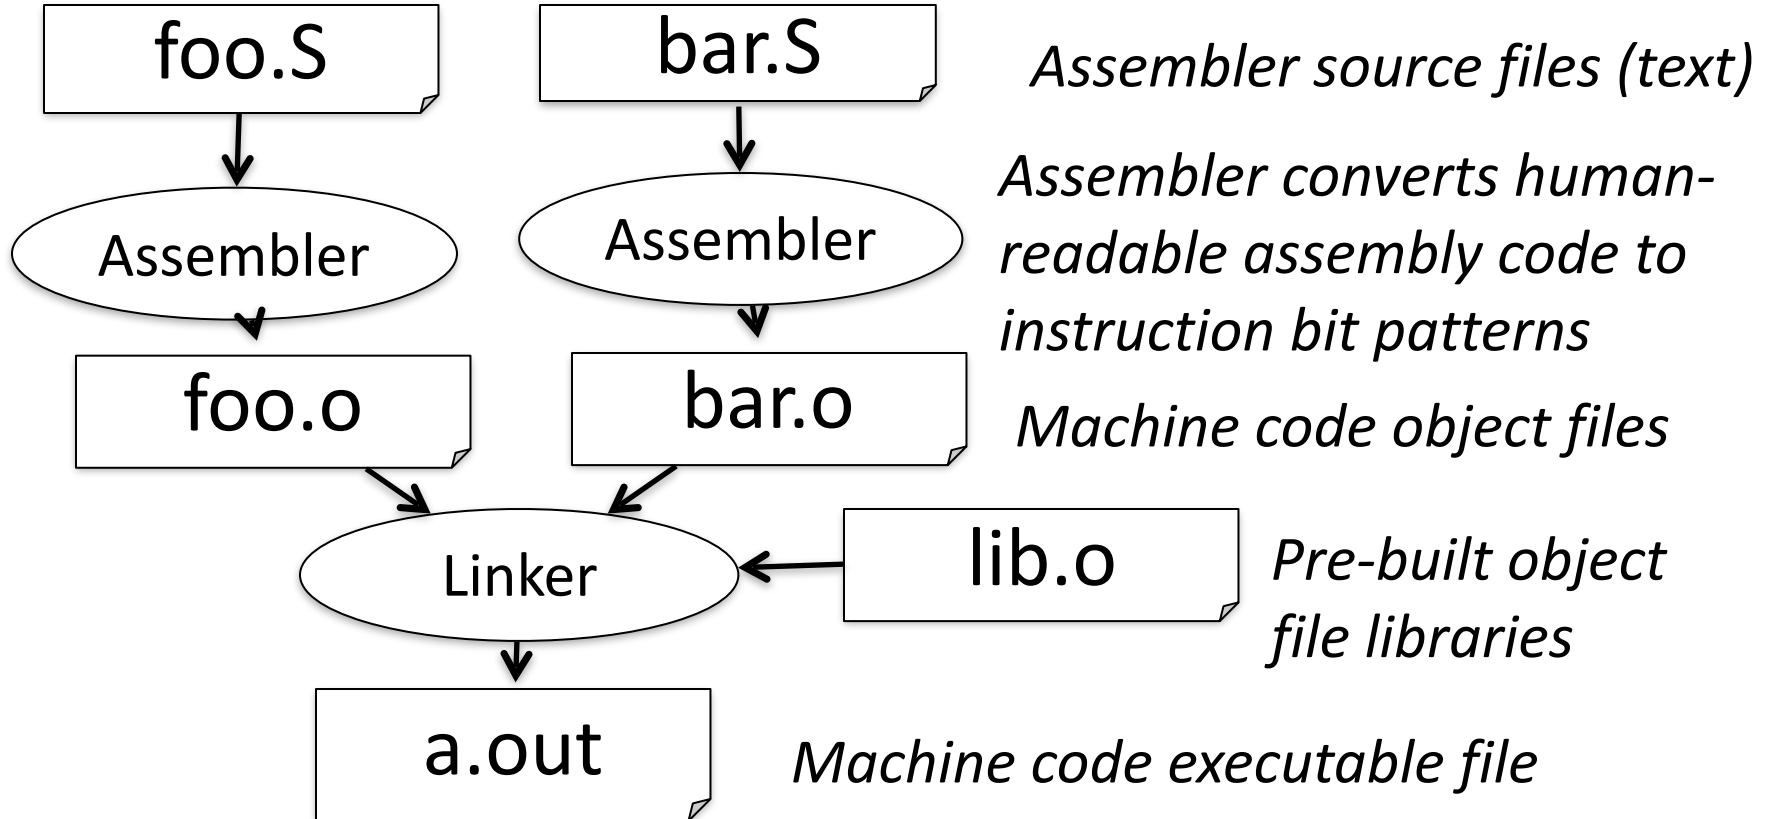
# Administrivia

- The Project 1 deadline extended to Thursday, 11:59pm!
- There will be a guerrilla section Thursday 7-9PM.
- Two weeks to Midterm #1!
- Project 2-1 release later this week or early next, due 2/16.
- Project 2-2 release right after midterm and due 2/23.

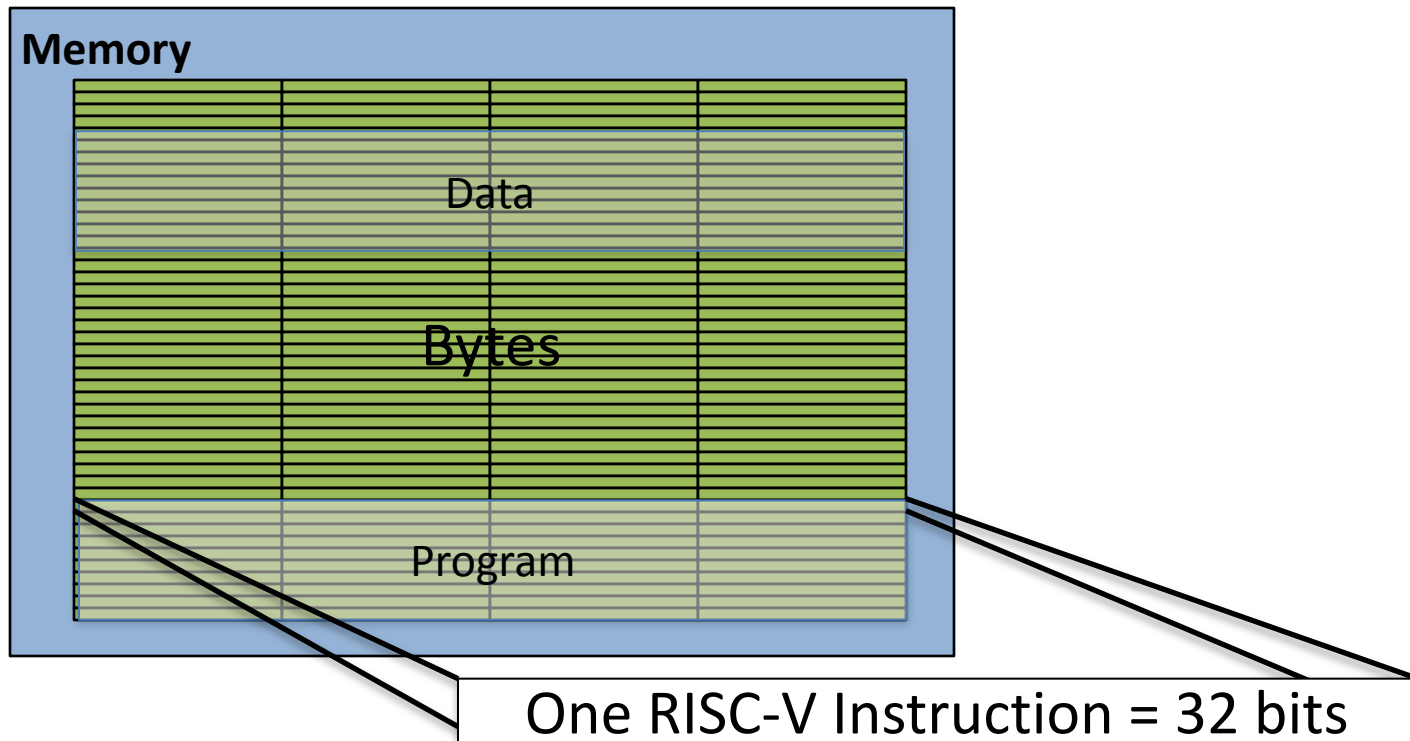
# Outline

- RISC-V ISA and C-to-RISC-V Review
- **Program Execution Overview**
- Function Call
- Function Call Example
- And in Conclusion ...

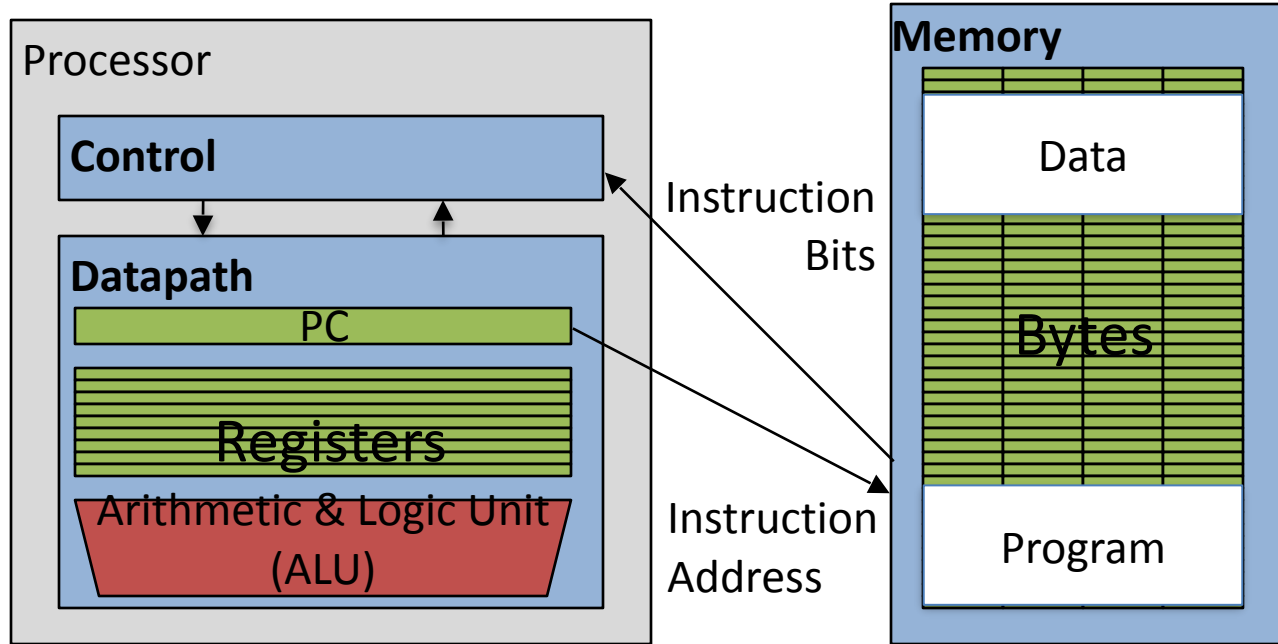
# Assembler to Machine Code (more later in course)



# How Program is Stored



# Program Execution



- **PC** (program counter) is special internal register inside processor holding byte address of next instruction to be executed
- Instruction is fetched from memory, then control unit executes instruction using datapath and memory system, and updates program counter (default is add +4 bytes to PC, to move to next sequential instruction)

# Helpful RISC-V Assembler Features

- Symbolic register names
  - E.g., **a0–a7** for argument registers (**x10–x17**)
  - E.g., **zero** for **x0**
- Pseudo-instructions
  - Shorthand syntax for common assembly idioms
  - E.g., **"mv rd, rs"** = **"addi rd, rs, 0"**
  - E.g., **"li rd, 13"** = **"addi rd, x0, 13"**




# RISC-V Symbolic Register Names

Numbers  
hardware  
understands



Human-friendly  
symbolic names  
in assembly code



Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	—
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5	t0	Temporary/alternate link register	Caller
x6–7	t1–2	Temporaries	Caller
x8	s0/fp	Saved register/frame pointer	Callee
x9	s1	Saved register	Callee
x10–11	a0–1	Function arguments/return values	Caller
x12–17	a2–7	Function arguments	Caller
x18–27	s2–11	Saved registers	Callee
x28–31	t3–6	Temporaries	Caller

# Outline

- RISC-V ISA and C-to-RISC-V Review
- Program Execution Overview
- **Function Call**
- Function Call Example
- And in Conclusion ...

# Six Fundamental Steps in Calling a Function

1. Put parameters in a place where function can access them
2. Transfer control to function
3. Acquire (local) storage resources needed for function
4. Perform desired task of the function
5. Put result value in a place where calling code can access it and maybe restore any registers you used
6. Return control to point of origin. (Note: a function can be called from several points in a program.)

# RISC-V Function Call Conventions

- Registers faster than memory, so use them
- **a0–a7** (**x10–x17**): eight *argument* registers to pass parameters and two return values (**a0–a1**)
- **ra**: one *return address* register for return to the point of origin (**x1**)

# Instruction Support for Functions (1/4)

C

```
... sum(a,b);... /* a,b:s0,s1 */  
}  
int sum(int x, int y) {  
    return x+y;  
}
```

RISC-V

---

address (shown in decimal)

1000

1004

1008


1012

1016

...

2000

2004



In RISC-V, all instructions are 4 bytes, and stored in memory just like data. So here we show the addresses of where the programs are stored.

# Instruction Support for Functions (2/4)

C

```
... sum(a,b);... /* a,b:s0,s1 */  
}  
int sum(int x, int y) {  
    return x+y;  
}
```

RISC-V

---

address (shown in decimal)

```
1000 mv a0,s0 # x = a  
1004 mv a1,s1 # y = b  
1008 addi ra,zero,1016 #ra=1016  
1012 j      sum      #jump to sum  
1016 ...           # next instruction  
...  
2000 sum: add a0,a0,a1  
2004 jr      ra      # new instr. "jump register"
```

# Instruction Support for Functions (3/4)

C

```
... sum(a,b);... /* a,b:s0,s1 */  
}  
int sum(int x, int y) {  
    return x+y;  
}
```

RISC-V

- Question: Why use **jr** here? Why not use **j**?
- Answer: **sum** might be called by many places, so we can't return to a fixed place. The calling proc to **sum** must be able to say "return here" somehow.

```
2000 sum: add a0,a0,a1
```

```
2004 jr     ra     # new instr. "jump register"
```

# Instruction Support for Functions (4/4)

- Single instruction to jump and save return address: jump and link (**j<sub>al</sub>**)
- Before:  

```
1008 addi ra,zero,1016 #ra=1016
1012 j  sum           #goto sum
```
- After:  

```
1008 jal sum # ra=1012,goto sum
```
- Why have a **j<sub>al</sub>**?
  - Make the common case fast: function calls very common
  - Reduce program size
  - Don't have to know where code is in memory with **j<sub>al</sub>**!
- Return from function: *jump register* instruction (**j<sub>r</sub>**)
  - Unconditional jump to address specified in register: **j<sub>r</sub> ra**
  - Assembler shorthand: **ret = j<sub>r</sub> ra**



# Outline

- RISC-V ISA and C-to-RISC-V Review
- Program Execution Overview
- Function Call
- **Function Call Example**
- And in Conclusion ...

# Example

```
int Leaf(int g, int h, int i, int j)
{
    int f;
    f = (g + h) - (i + j);
    return f;
}
```

- Parameter variables **g**, **h**, **i**, and **j** in argument registers **a0**, **a1**, **a2**, and **a3**.
- Assume we compute **f** by using **s0** and **s1**

## Where Are Old Register Values Saved to Restore Them After Function Call?

- Need a place to save old values before call function, restore them when return
- Ideal is *stack*: last-in-first-out queue (e.g., stack of plates)
  - Push: placing data onto stack
  - Pop: removing data from stack
- Stack in memory, so need register to point to it
- **sp** is the *stack pointer* in RISC-V (**x2**)
- **sp** always points to the last used place on the stack
- Convention is grow stack down from high to low addresses
  - *Push* decrements **sp**, *Pop* increments **sp**

# RISC-V Code for Leaf()


```
Leaf: addi sp,sp,-8 # adjust stack for 2 items
      sw s1, 4(sp)  # save s1 for use afterwards
      sw s0, 0(sp)  # save s0 for use afterwards
```

```
      add s0,a0,a1 # s0 = g + h
      add s1,a2,a3 # s1 = i + j
      sub a0,s0,s1 # return value (g + h) - (i + j)
```


```
      lw s0, 0(sp) # restore register s0 for caller
      lw s1, 4(sp) # restore register s1 for caller
      addi sp,sp,8 # adjust stack to delete 2 items
      jr ra       # jump back to calling routine
```

# Stack Before, During, After Function


- Need to save old values of `s0` and `s1`



Before call



During call




After call

# Break!



# New RISC-V book!

- “The RISC-V Reader”, David Patterson, Andrew Waterman
- Available from Amazon
- Print edition **\$19.99**
- Kindle edition to follow at some point
- **Recommended, not required**



# What If a Function Calls a Function? Recursive Function Calls?

- Would clobber values in **a0-a7** and **ra**
- What is the solution?



# Nested Procedures (1/2)

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```

- Something called **sumSquare**, now **sumSquare** is calling **mult**
- So there's a value in **ra** that **sumSquare** wants to jump back to, but this will be overwritten by the call to **mult**

Need to save **sumSquare** return address before call to **mult**

# Nested Procedures (2/2)

- In general, may need to save some other info in addition to **ra**.
- When a C program is run, there are three important memory areas allocated:
  - **Static**: Variables declared once per program, cease to exist only after execution completes - e.g., C globals
  - **Heap**: Variables declared dynamically via **malloc**
  - **Stack**: Space to be used by procedure during execution; this is where we can save register values

# Optimized Function Convention

To reduce expensive loads and stores from spilling and restoring registers, RISC-V function-calling convention divides registers into two categories:

1. Preserved across function call
  - Caller can rely on values being unchanged
  - **sp, gp, tp**, “saved registers” **s0- s11** (**s0** is also **fp**)
2. Not preserved across function call
  - Caller *cannot* rely on values being unchanged
  - Argument/return registers **a0-a7**, **ra**, “temporary registers” **t0-t6**

# Peer Instruction

- Which statement is FALSE?
- **A:** RISC-V uses **jal** to invoke a function and **jr** to return from a function
- **B:** **jal** saves PC+1 in **ra**
- **C:** The callee can use temporary registers (**t*i***) without saving and restoring them
- **D:** The caller can rely on save registers (**s*i***) without fear of callee changing them

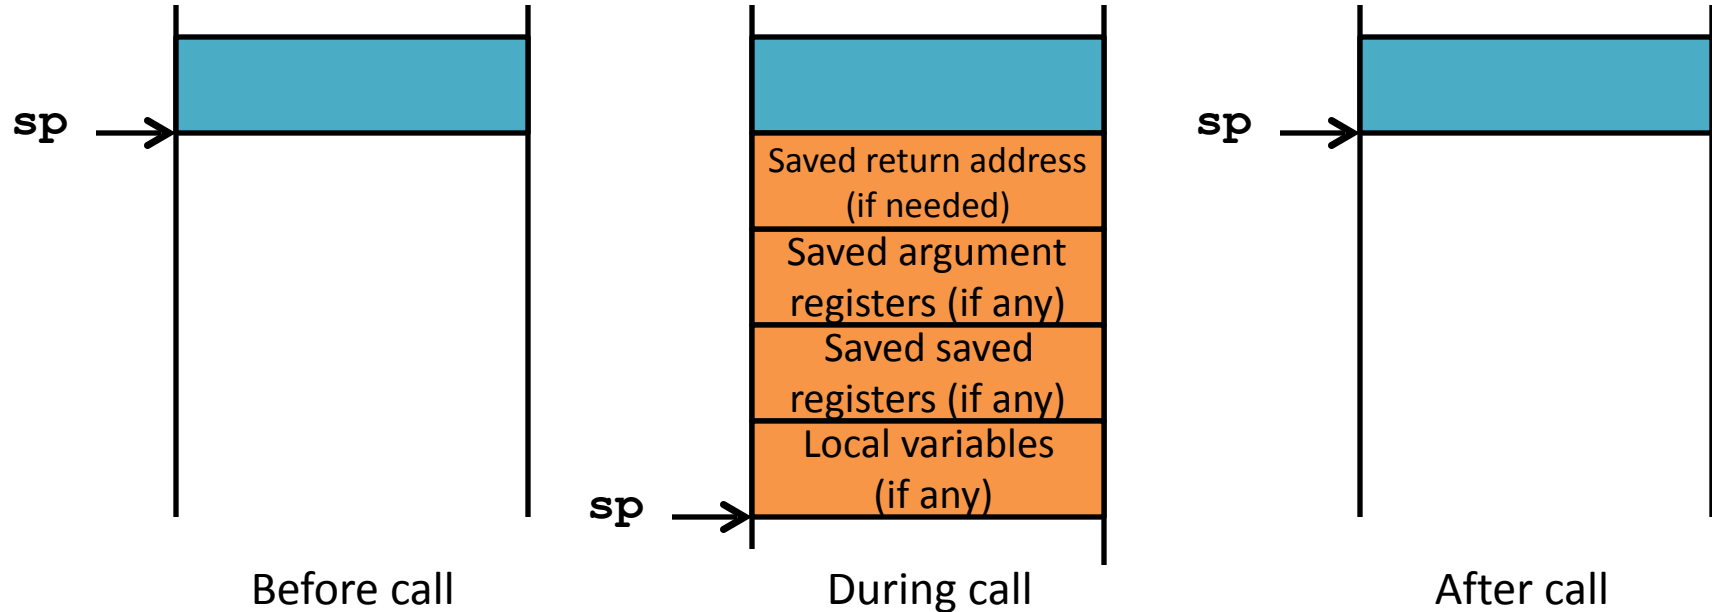
# Peer Instruction

- Which statement is FALSE?
- **A:** RISC-V uses **jal** to invoke a function and **jr** to return from a function
- **B:** **jal** saves PC+1 in **ra**
- **C:** The callee can use temporary registers (**ti**) without saving and restoring them
- **D:** The caller can rely on save registers (**si**) without fear of callee changing them

# Allocating Space on Stack

- C has two storage classes: automatic and static
  - *Automatic* variables are local to function and discarded when function exits
  - *Static* variables exist across exits from and entries to procedures
- Use stack for automatic (local) variables that aren't in registers
- *Procedure frame* or *activation record*: segment of stack with saved registers and local variables

# Stack Before, During, After Function



# Using the Stack (1/2)

- So we have a register **sp** which always points to the last used space in the stack
- To use stack, we decrement this pointer by the amount of space we need and then fill it with info
- So, how do we compile this?

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```



# Using the Stack (2/2)

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y; }
```

sumSquare:

“push”

```
addi sp,sp,-8    # reserve space on stack  
sw ra, 4(sp)     # save ret addr  
sw a1, 0(sp)     # save y  
mv a1,a0         # mult(x,x)  
jal mult         # call mult  
lw a1, 0(sp)     # restore y  
add a0,a0,a1     # mult()+y  
lw ra, 4(sp)     # get ret addr  
addi sp,sp,8     # restore stack  
jr ra
```

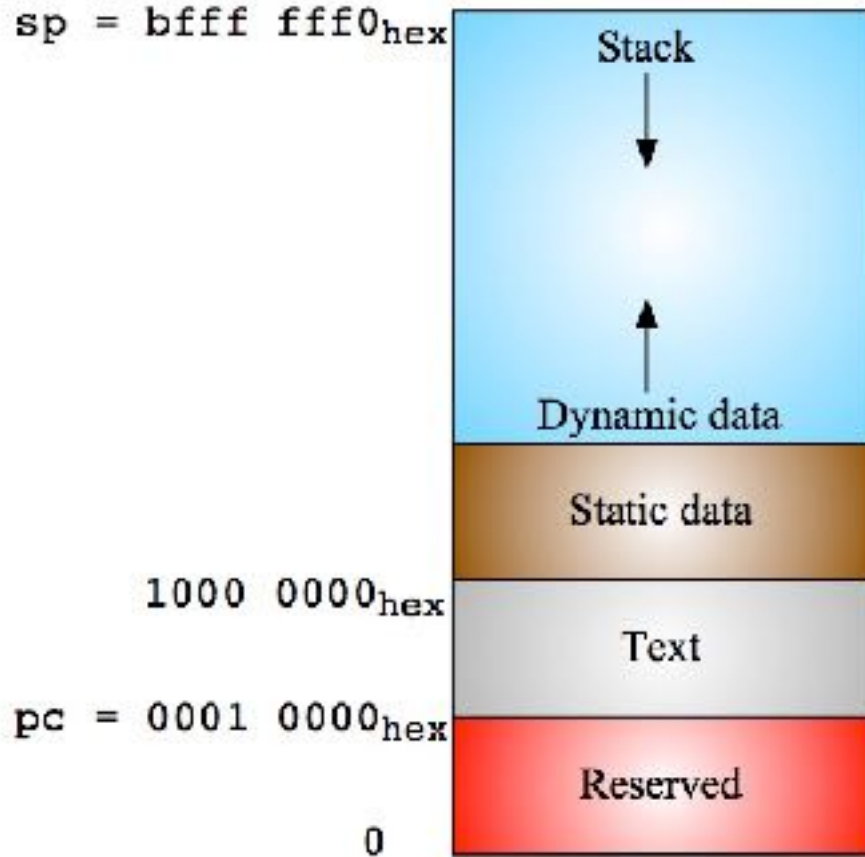
“pop”

mult: ...

# Where is the Stack in Memory?

- RV32 convention (RV64 and RV128 have different memory layouts)
- Stack starts in high memory and grows down
  - Hexadecimal (base 16) : **ffff\_fff0**<sub>hex</sub>
- RV32 programs (*text segment*) in low end
  - **0001\_0000**<sub>hex</sub>
- *static data segment* (constants and other static variables) above text for static variables
  - RISC-V convention *global pointer* (**gp**) points to static
  - RV32 **gp** = **1000\_0000**<sub>hex</sub>
- *Heap* above static for data structures that grow and shrink ; grows up to high addresses

# RV32 Memory Allocation



# Outline

- RISC-V ISA and C-to-RISC-V Review
- Program Execution Overview
- Function Call
- Function Call Example
- **And in Conclusion ...**

# And in Conclusion ...

- Functions called with **jal**, return with **jr ra**.
- The stack is your friend: Use it to save anything you need. Just leave it the way you found it!
- Instructions we know so far...
  - Arithmetic: **add, addi, sub**
  - Memory: **lw, sw, lb, lbu, sb**
  - Decision: **beq, bne, blt, bge**
  - Unconditional Branches (Jumps): **j, jal, jr**
- Registers we know so far
  - All of them!
  - **a0–a7** for function arguments, **a0–a1** for return values
  - **sp**, stack pointer, **ra** return address
  - **s0–s11** saved registers
  - **t0–t6** temporaries
  - **zero**