

CS 61C:  
Great Ideas in Computer Architecture  
Lecture 11: *RISC-V Processor Datapath*

John Wawrzynek & Nick Weaver  
<http://inst.eecs.berkeley.edu/~cs61c/sp18>

# Great Idea #1: Abstraction

## (Levels of Representation/Interpretation)

Computer Science 61C Spring 2018

Wawrzynek and Weaver

```
lw  t0, t2, 0
lw  t1, t2, 4
sw  t1, t2, 0
sw  t0, t2, 4
```

High Level Language  
Program (e.g., C)

*Compiler*

Assembly Language Program  
(e.g., RISC-V)

*Assembler*

Machine Language Program  
(RISC-V)

```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
```

Anything can be represented  
as a *number*,  
i.e., data or instructions

```
0000 1001 1100 0110 1010 1111 0101 1000
1010 1111 0101 1000 0000 1001 1100 0110
1100 0110 1010 1111 0101 1000 0000 1001
0101 1000 0000 1001 1100 0110 1010 1111
```



**We are here!**

*Machine  
Interpretation*

Hardware Architecture Description  
(e.g., block diagrams)

*Architecture  
Implementation*

Logic Circuit Description  
(Circuit Schematic Diagrams)



# Recap: Complete RV32I ISA

imm[31:12]				rd	0110111	LUI
imm[31:12]				rd	0010111	AUIPC
imm[25:10:1][11]10:12				rd	1101111	JAL
imm[11:2]				rd	1101111	JALR
imm[12]10:5	rs2	rs1	000	imm[4:1]11	1100011	BEQ
imm[12]10:5	rs2	rs1	001	imm[4:1]11	1100011	BNE
imm[12]10:5	rs2	rs1	100	imm[4:1]11	1100011	BLT
imm[12]10:5	rs2	rs1	101	imm[4:1]11	1100011	BCE
imm[12]10:5	rs2	rs1	110	imm[4:1]11	1100011	BLTU
imm[12]10:5	rs2	rs1	111	imm[4:1]11	1100011	BCEU
imm[11:2]		rs1	000	rd	0001001	LB
imm[11:2]		rs1	001	rd	0001001	LH
imm[11:2]		rs1	010	rd	0001001	LSW
imm[11:2]		rs1	100	rd	0001001	LBU
imm[11:2]		rs1	101	rd	0001001	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	SW
imm[11:2]		rs1	000	rd	0010011	ADDI
imm[11:2]		rs1	010	rd	0010011	SLLI
imm[11:2]		rs1	011	rd	0010011	SLLI
imm[11:2]		rs1	100	rd	0010011	XORI
imm[11:2]		rs1	110	rd	0010011	ORI
imm[11:2]		rs1	111	rd	0010011	ANDI

0000000	shamt	rs1	001	rd	0010011	SLLI	
0000000	shamt	rs1	101	rd	0010011	SRLI	
0100000	shamt	rs1	101	rd	0010011	SRAI	
0000000	rs2	rs1	000	rd	0110011	ADD	
0100000	rs2	rs1	000	rd	0110011	SUB	
0000000	rs2	rs1	001	rd	0110011	SLL	
0000000	rs2	rs1	010	rd	0110011	SLLI	
0000000	rs2	rs1	011	rd	0110011	SLLIU	
0000000	rs2	rs1	100	rd	0110011	XOR	
0000000	rs2	rs1	101	rd	0110011	SRL	
0100000	rs2	rs1	101	rd	0110011	SRA	
0000000	rs2	rs1	110	rd	0110011	OR	
0000000	rs2	rs1	111	rd	0110011	AND	
0000	pred	rs2c	0000	000	00000	0001111	FENCE
0000	0000	0000	0000	001	00000	0001111	FENCEI
0000000000000000			00000	000	00000	1110011	ECALL
0000000000000000			00000	000	00000	1110011	EBREAK
csr			rs1	001	rd	1110011	CSRRW
csr			rs1	000	rd	1110011	CSRRS
csr			rs1	011	rd	1110011	CSRRC
csr			shamt	101	rd	1110011	CSRRWI
csr			shamt	110	rd	1110011	CSRRSI
csr			shamt	111	rd	1110011	CSRRCI

Not in CS61C

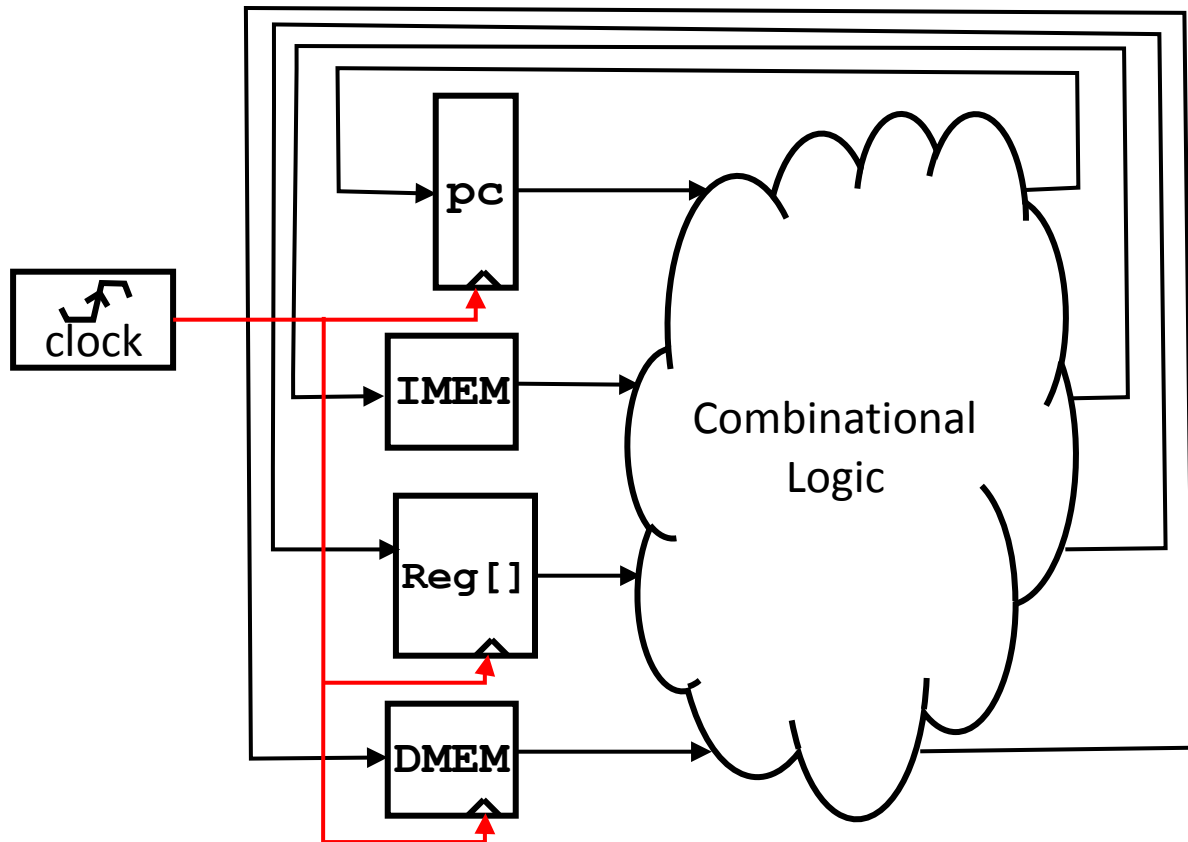
# “State” Required by RV32I ISA

Each instruction reads and updates this state during execution:

- Registers (**x0** . . **x31**)
  - Register file (or *regfile*) **Reg** holds 32 registers x 32 bits/register: **Reg**[0] . . **Reg**[31]
  - First register read specified by *rs1* field in instruction
  - Second register read specified by *rs2* field in instruction
  - Write register (destination) specified by *rd* field in instruction
  - **x0** is always 0 (writes to **Reg**[0] are ignored)
- Program Counter (**PC**)
  - Holds address of current instruction
- Memory (**MEM**)
  - Holds both instructions & data, in one 32-bit byte-addressed memory space
  - We’ll use separate memories for instructions (**IMEM**) and data (**DMEM**)
    - *Later we’ll replace these with instruction and data caches*
  - Instructions are read (*fetched*) from instruction memory (assume **IMEM** read-only)
  - Load/store instructions access data memory



# One-Instruction-Per-Cycle RISC-V Machine

On every tick of the clock, the computer executes one instruction



1. Current state outputs drive the inputs to the combinational logic, whose outputs settle at the values of the state before the next clock edge
2. At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle

# Basic Phases of Instruction Execution



Clock

time

# Implementing the **add** instruction

0000000	rs2	rs1	000	rd	0110011	ADD
---------	-----	-----	-----	----	---------	-----

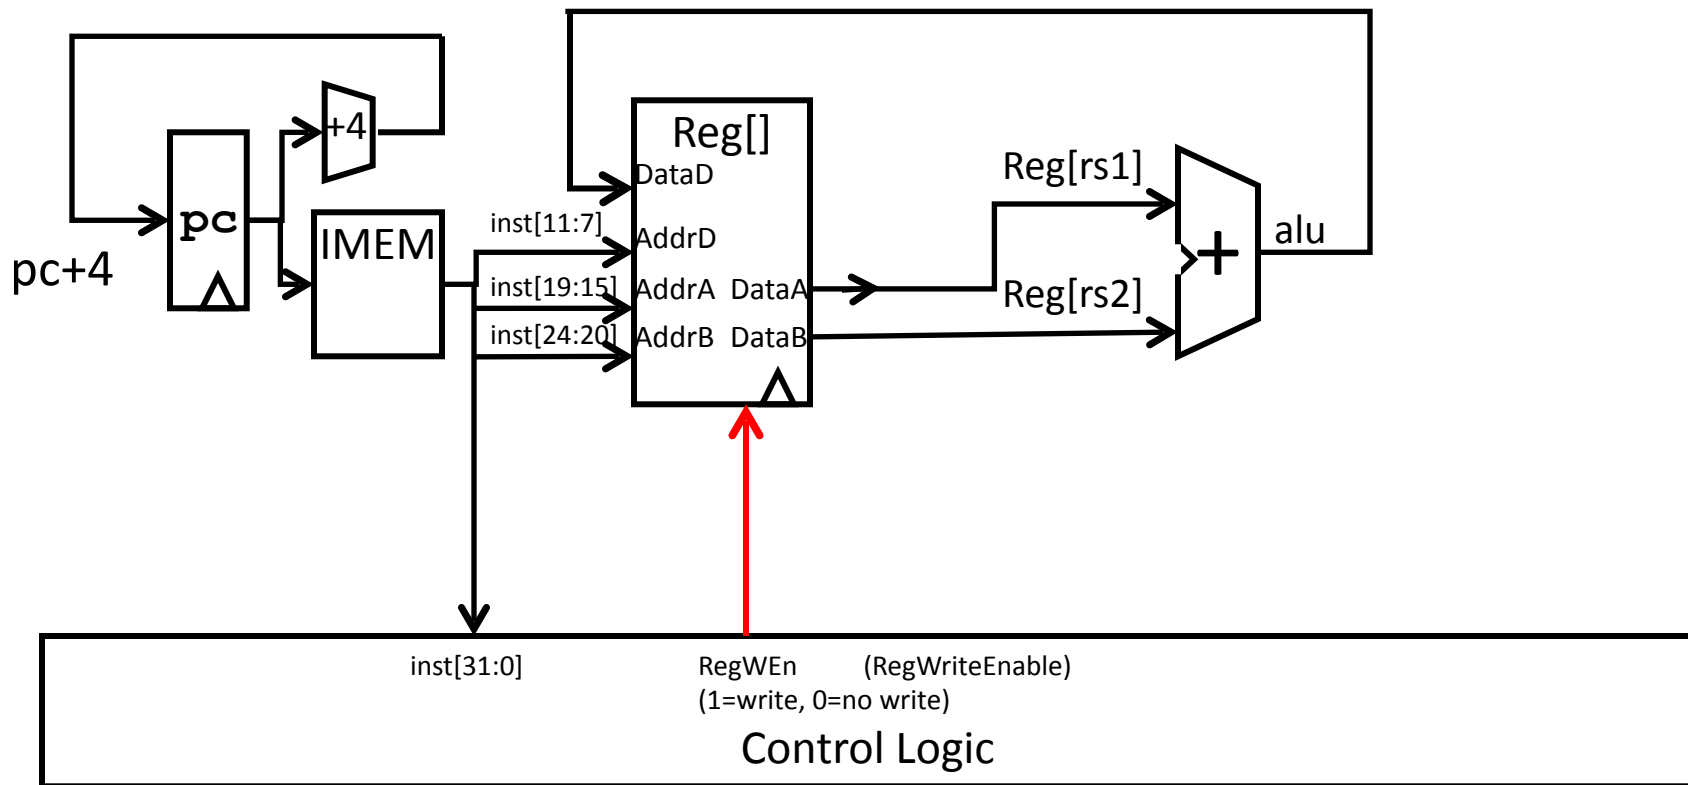
**add rd, rs1, rs2**

- Instruction makes two changes to machine's state:

**Reg[rd] = Reg[rs1] + Reg[rs2]**

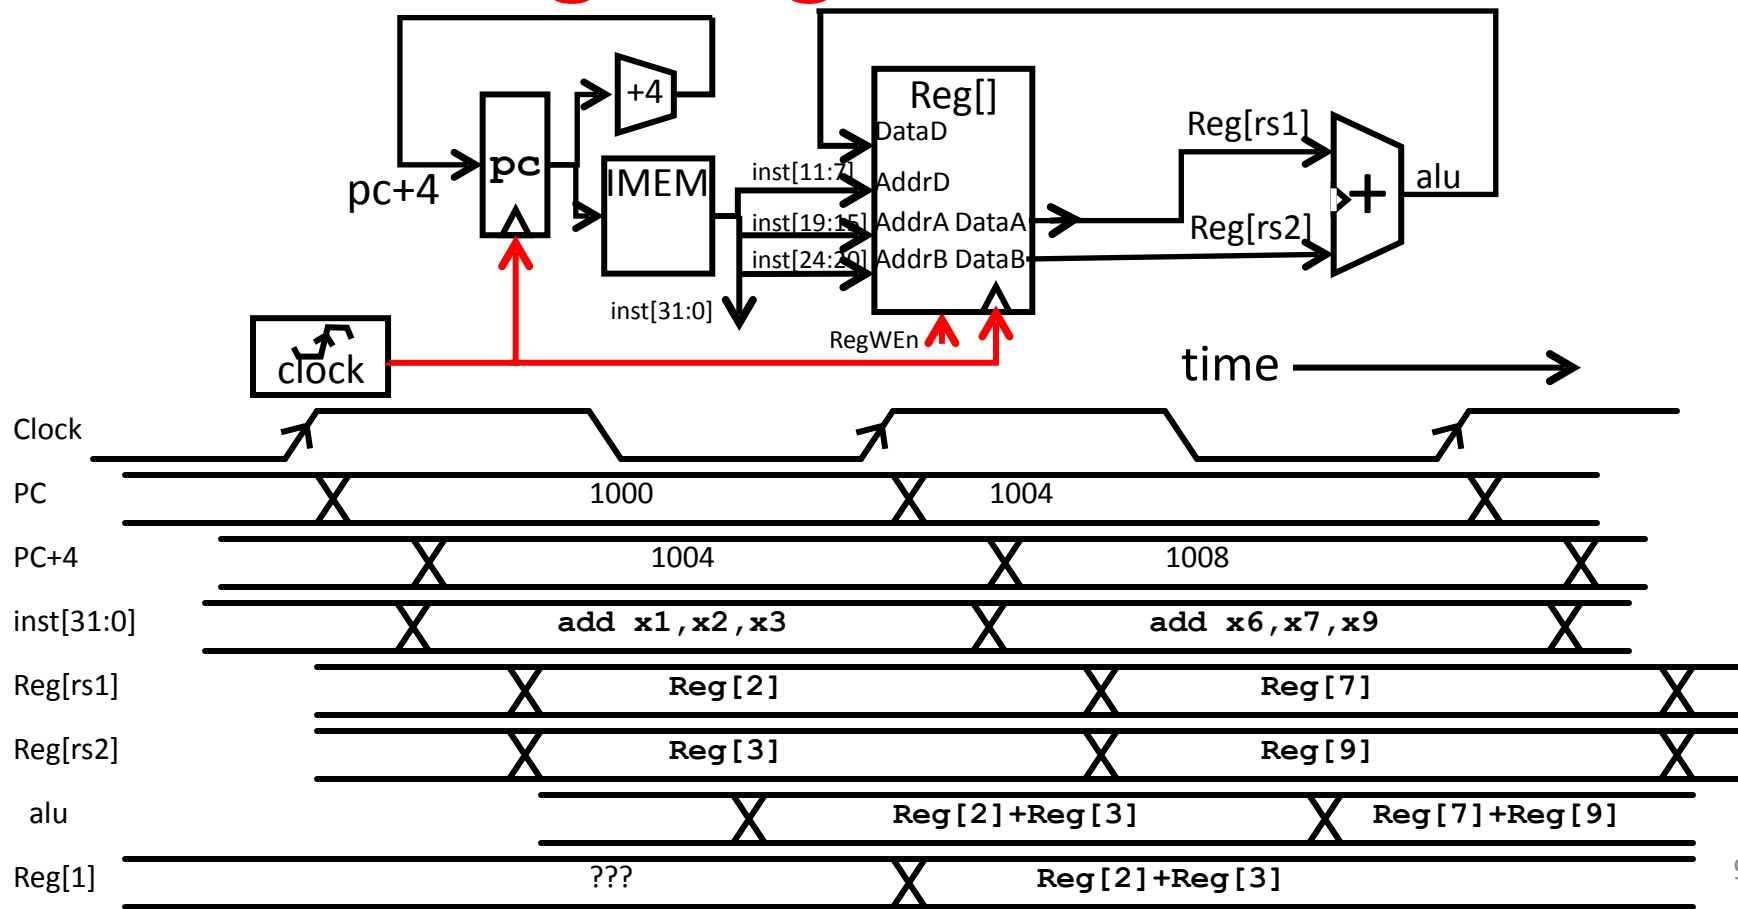
**PC = PC + 4**

# Datapath for add





# Timing Diagram for add



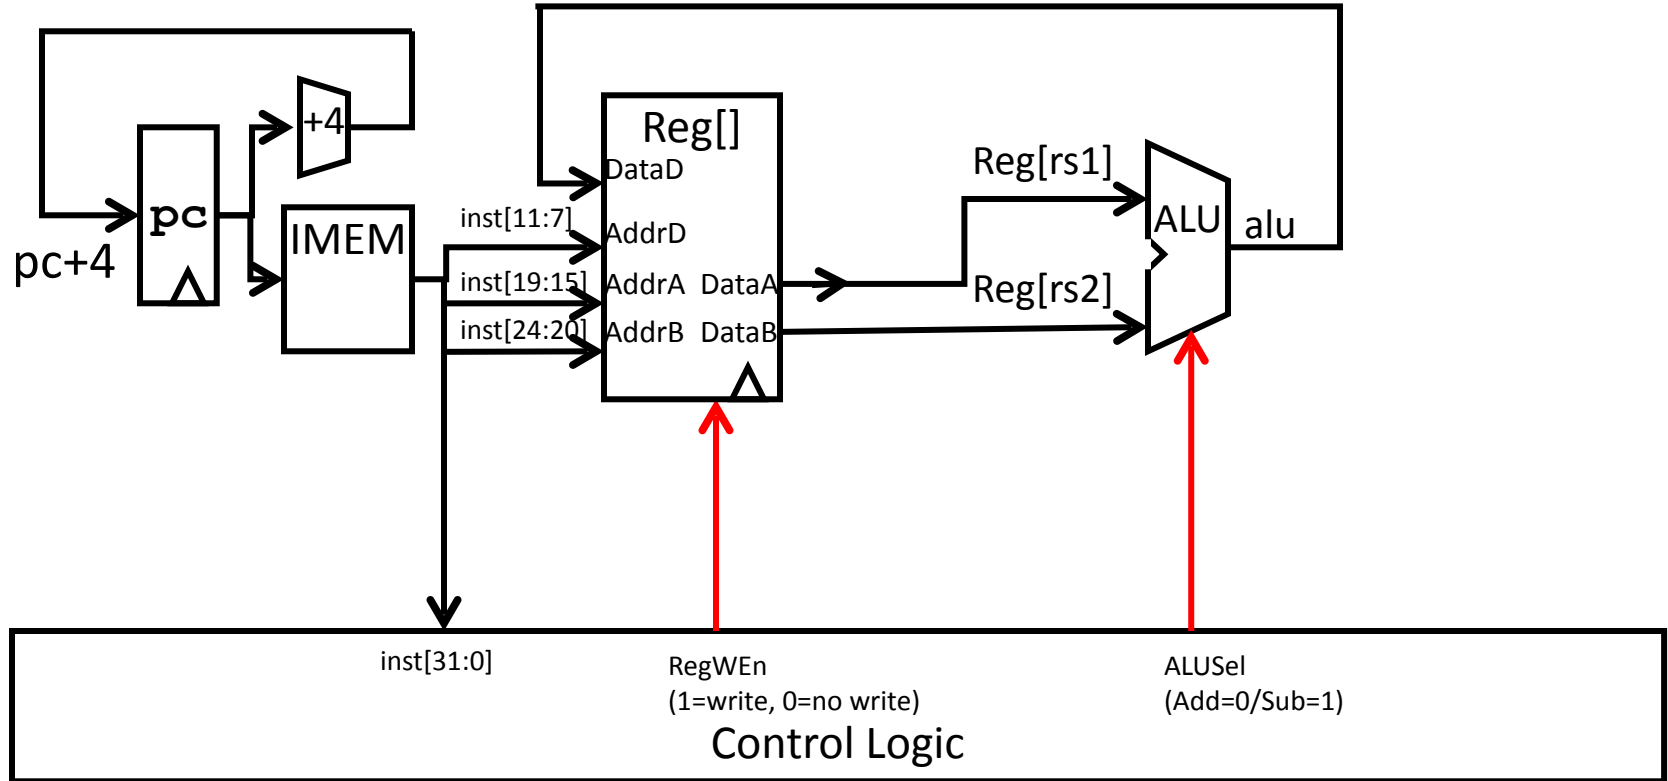
# Implementing the **sub** instruction

0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB

**sub rd, rs1, rs2**

- Almost the same as add, except now have to subtract operands instead of adding them
- **inst[30]** selects between add and subtract

# Datapath for add/sub



# Implementing other R-Format instructions


0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND

- All implemented by decoding funct3 and funct7 fields and selecting appropriate ALU function

# Implementing the **addi** instruction

- RISC-V Assembly Instruction:

**addi    x15, x1, -50**



imm=-50

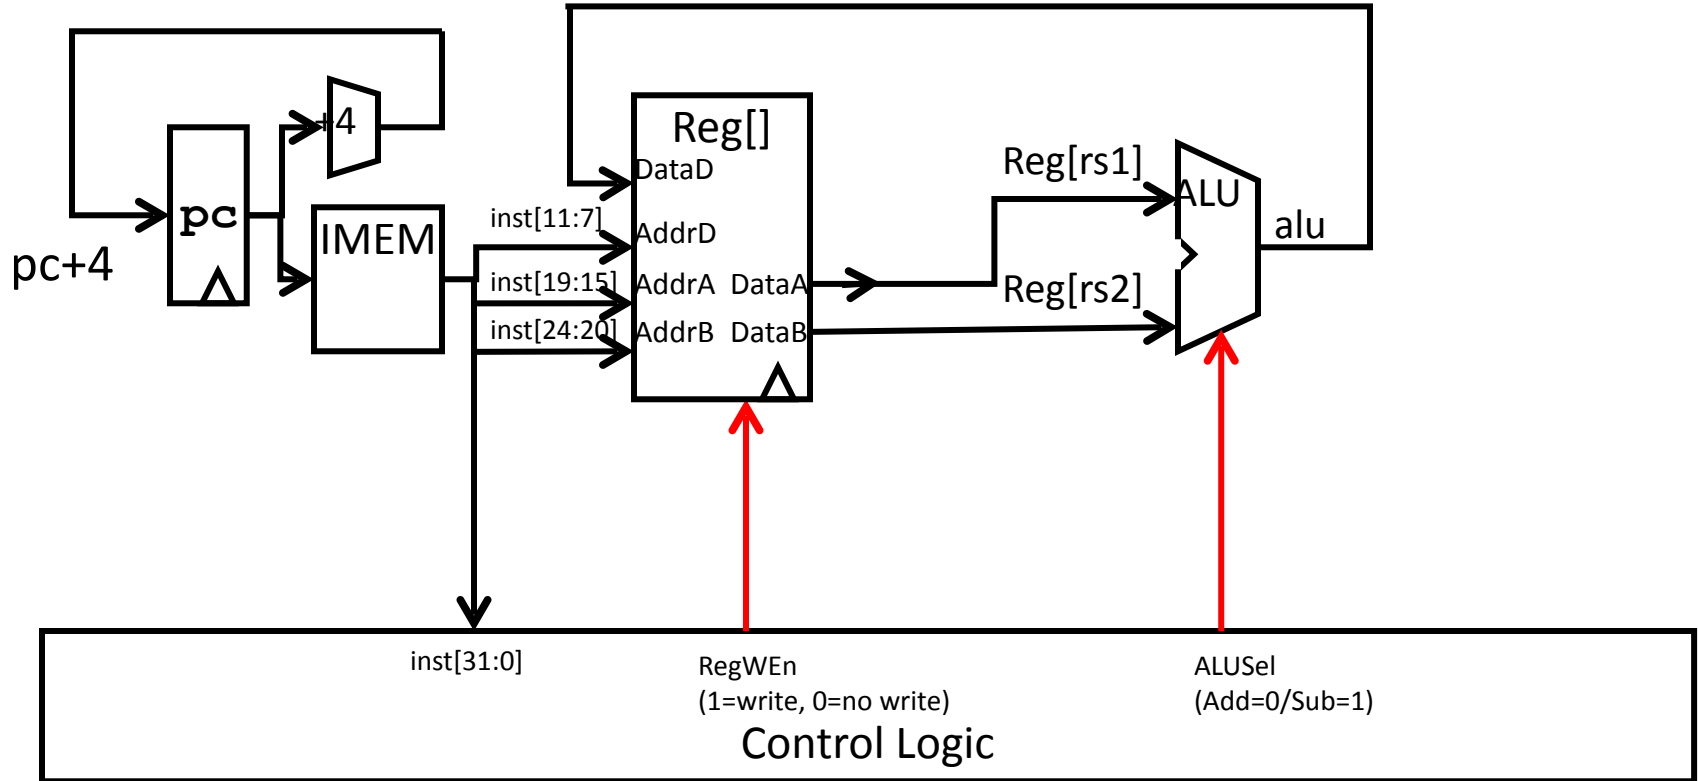
rs1=1

ADD

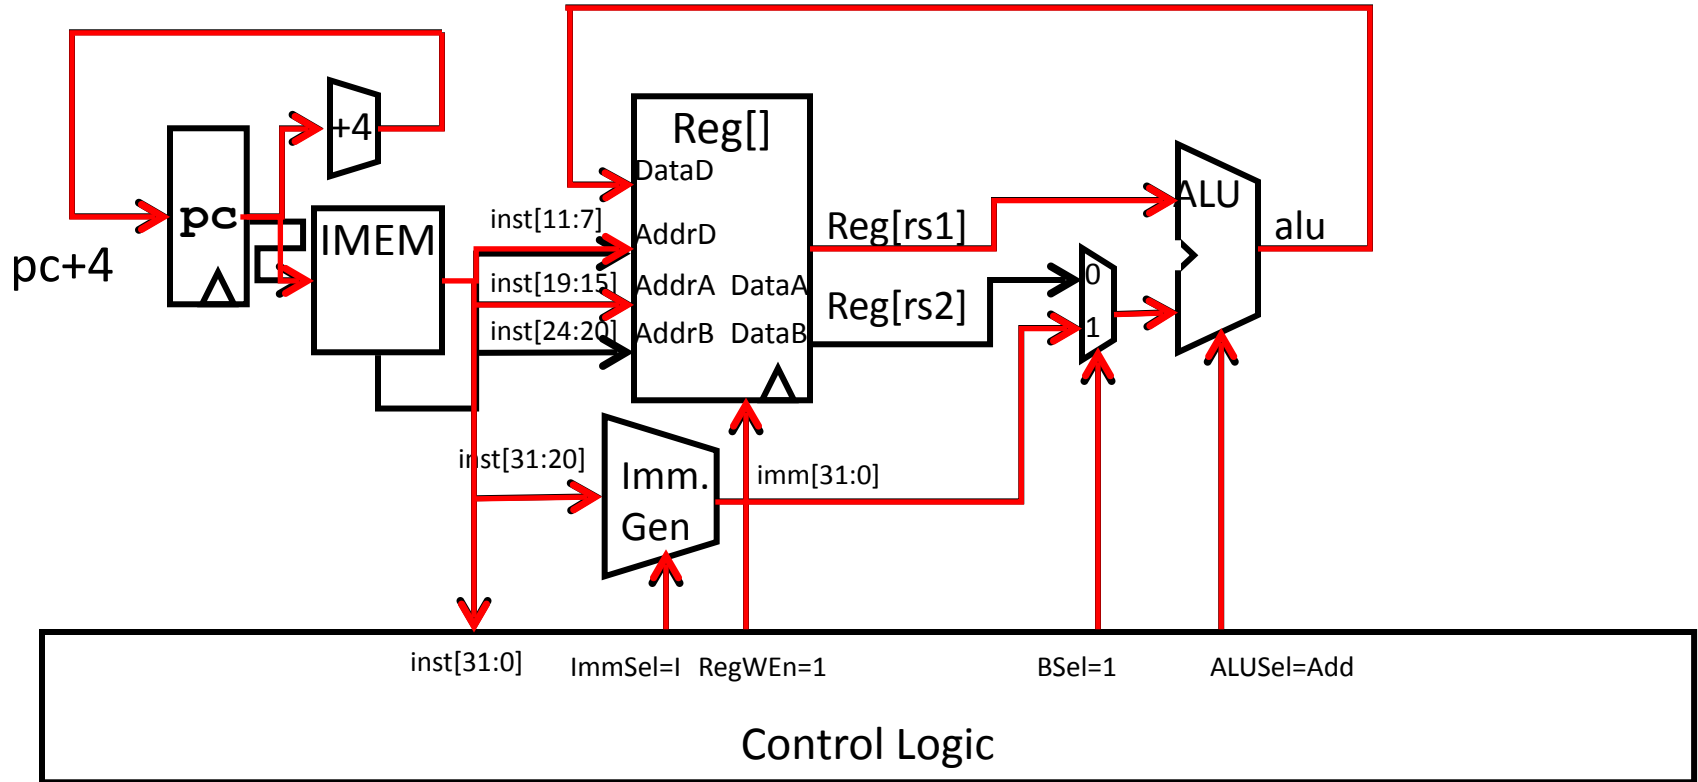
rd=15

OP-Imm



# Datapath for add/sub



# Adding **addi** to datapath



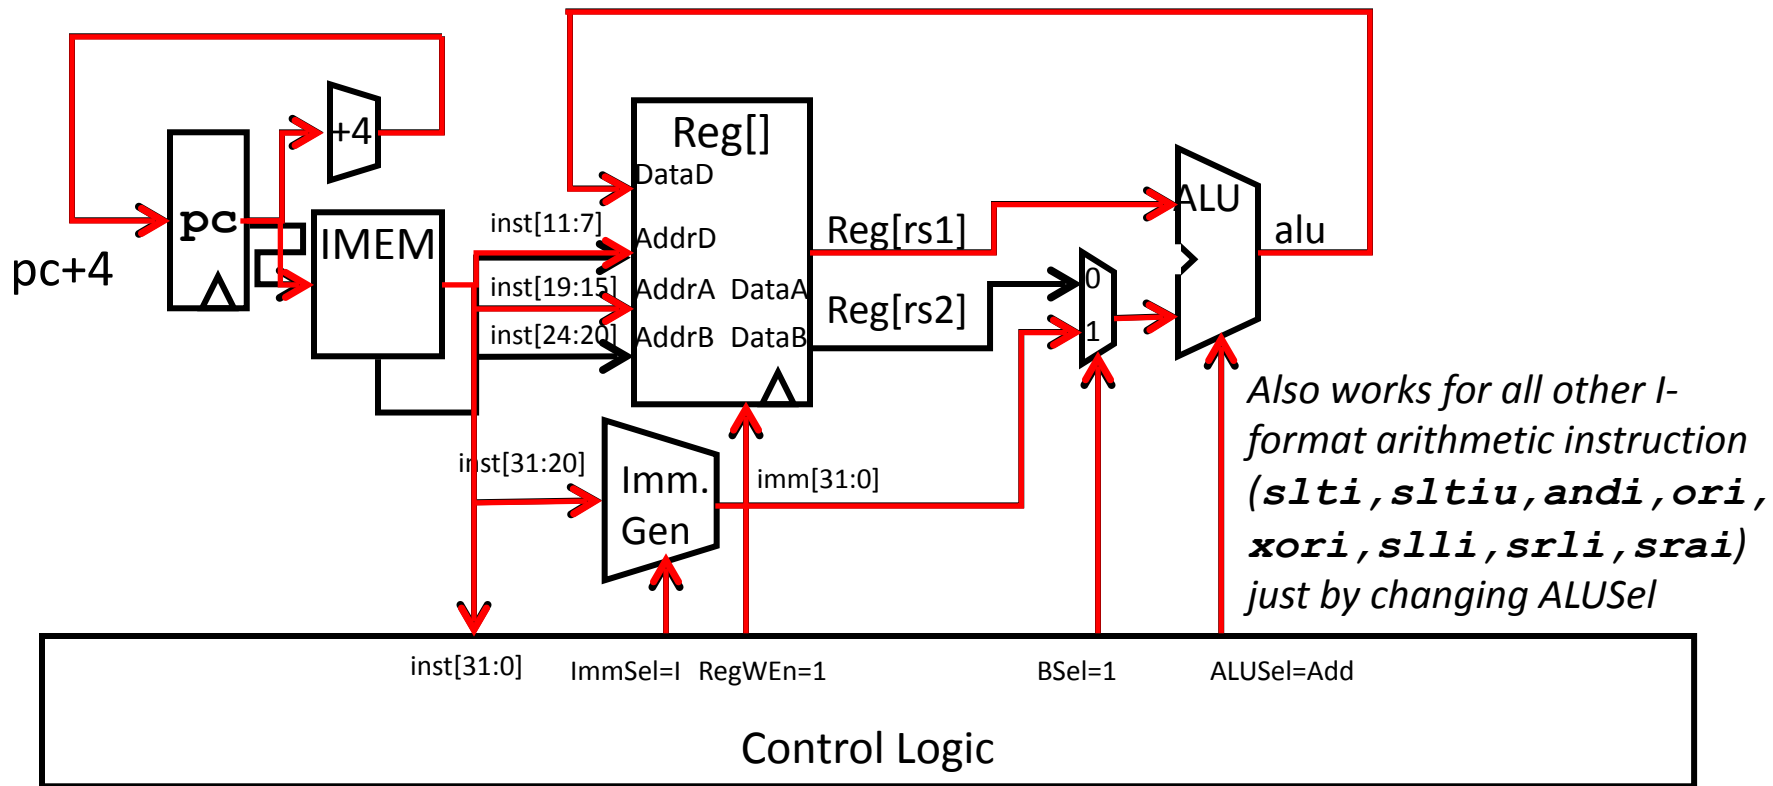
# I-Format immediates



- High 12 bits of instruction (`inst[31:20]`) copied to low 12 bits of immediate (`imm[11:0]`)
- Immediate is sign-extended by copying value of `inst[31]` to fill the upper 20 bits of the immediate value (`imm[31:12]`)



# Adding **addi** to datapath





# Break!



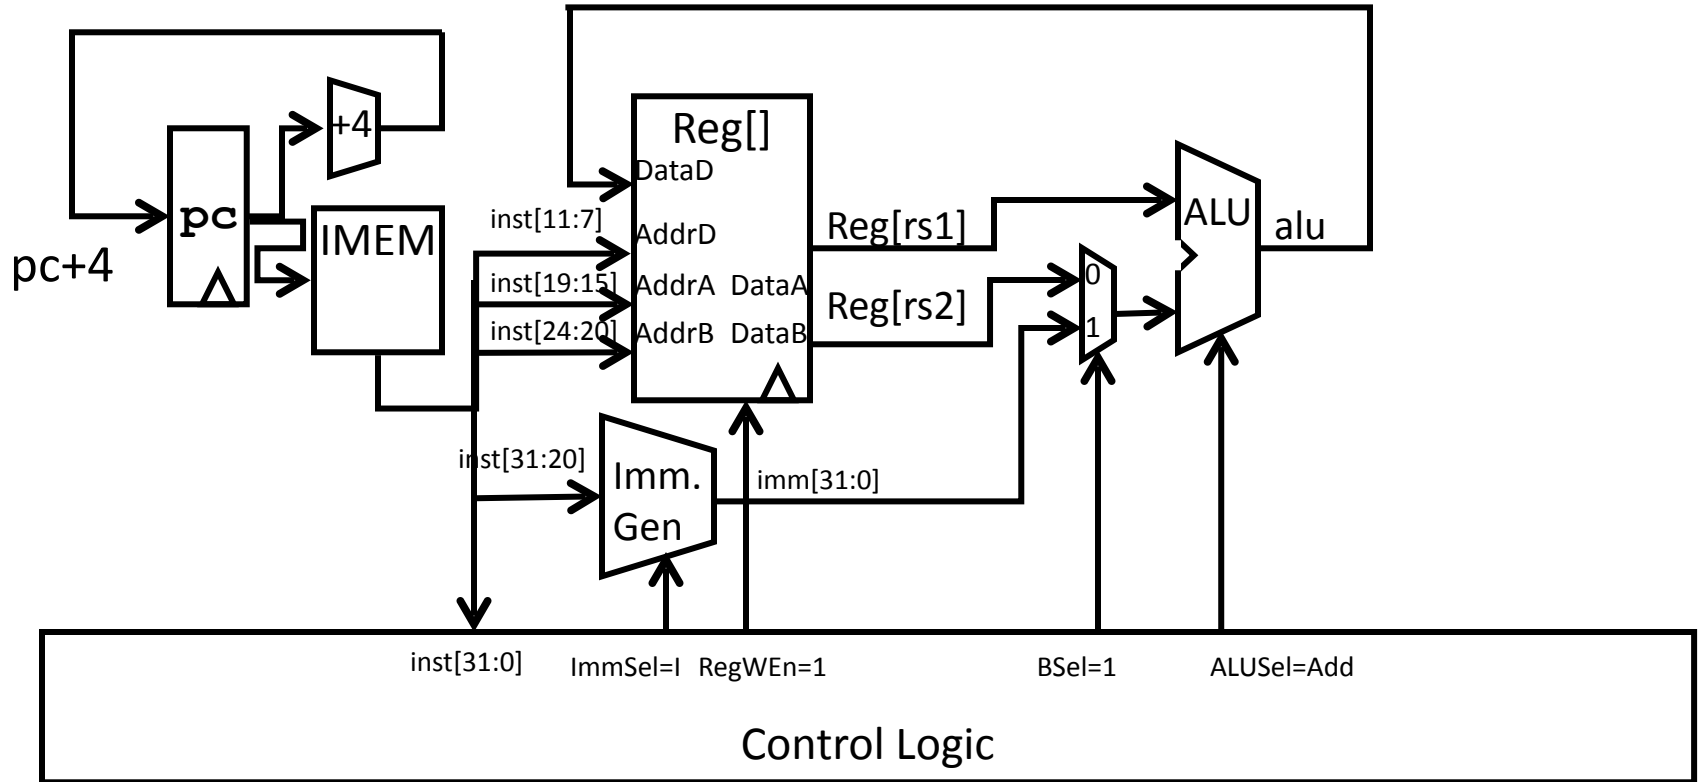
# Implementing Load Word instruction

- RISC-V Assembly Instruction:


**lw x14, 8(x2)**



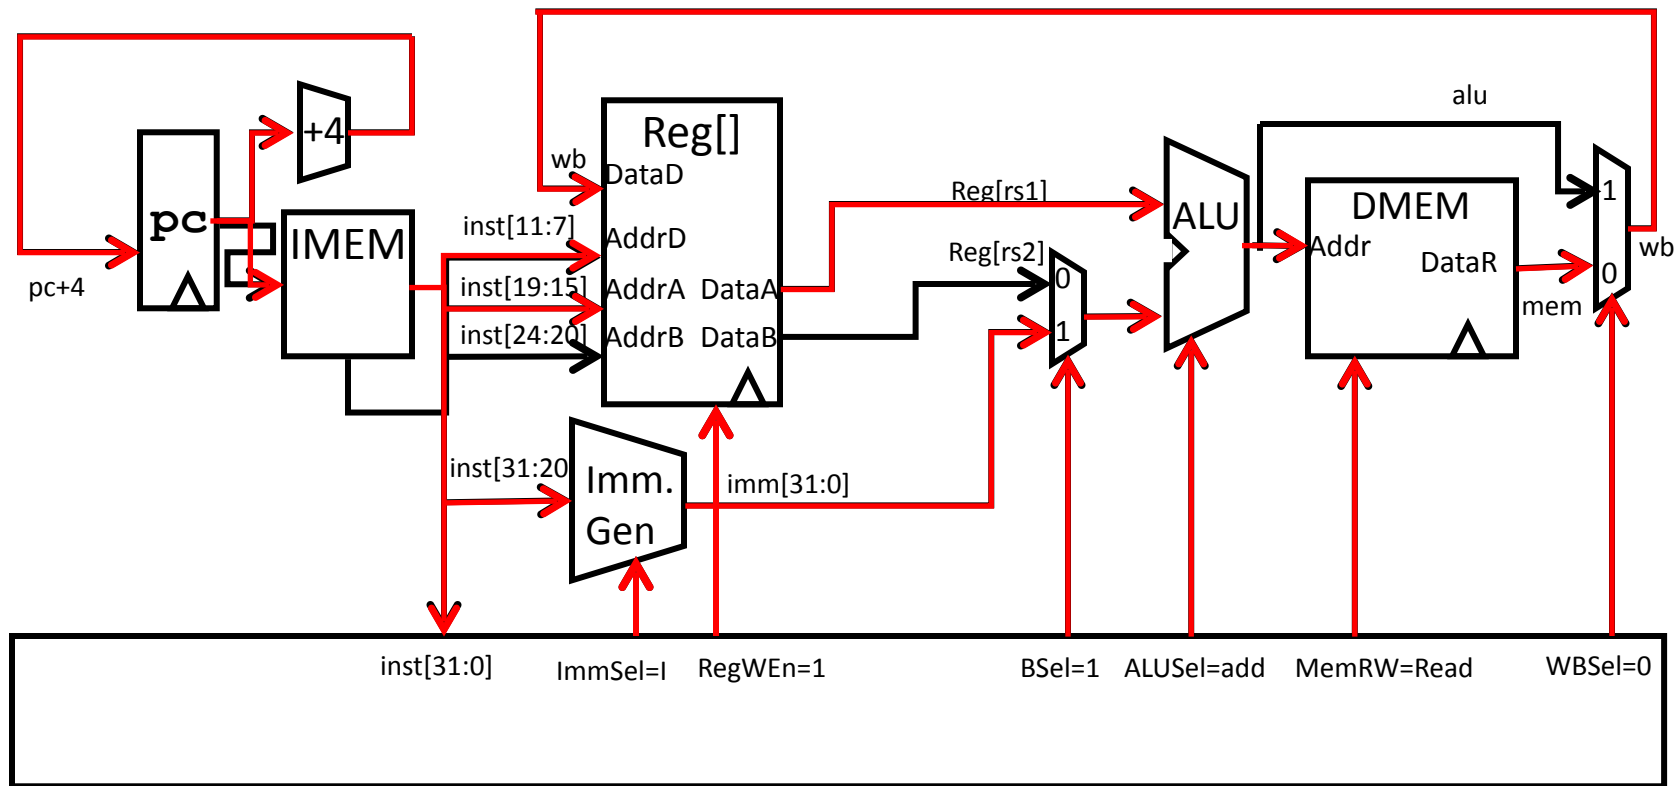
# Adding addi to datapath



# Adding lw to datapath



# Adding `lw` to datapath



# All RV32 Load Instructions

imm[11:0]	rs1	000	rd	0000011	LB
imm[11:0]	rs1	001	rd	0000011	LH
imm[11:0]	rs1	010	rd	0000011	LW
imm[11:0]	rs1	100	rd	0000011	LBU
imm[11:0]	rs1	101	rd	0000011	LHU


↑  
funct3 field encodes size and  
signedness of load data

- Supporting the narrower loads requires additional circuits to extract the correct byte/halfword from the value loaded from memory, and sign- or zero-extend the result to 32 bits before writing back to register file.


# Implementing Store Word instruction

- RISC-V Assembly Instruction:

**sw x14, 8(x2)**



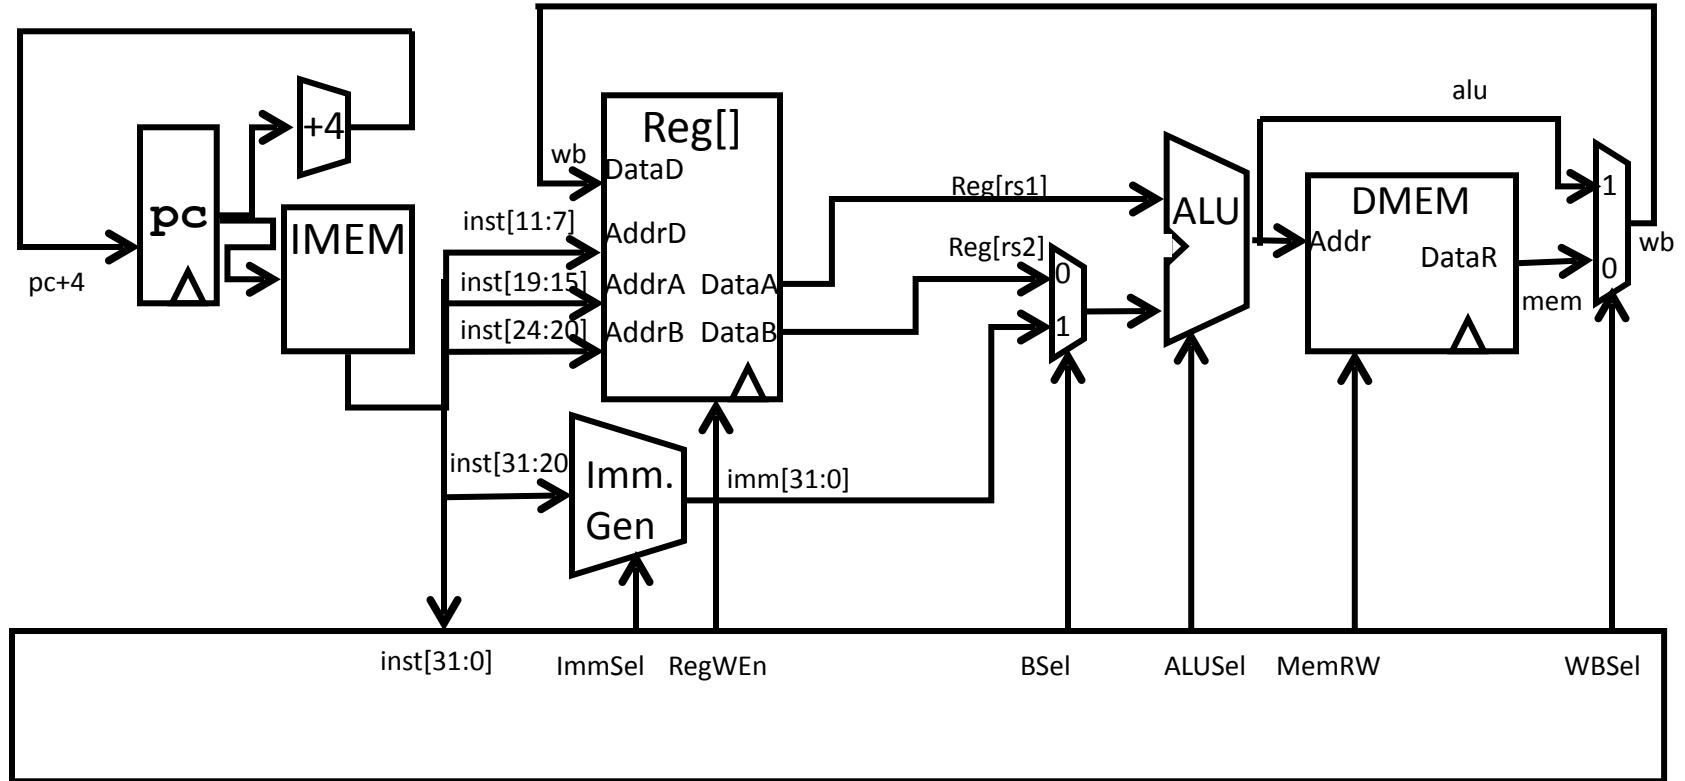
offset[11:5] = 0      rs2=14      rs1=2      SW      offset[4:0] = 8      STORE



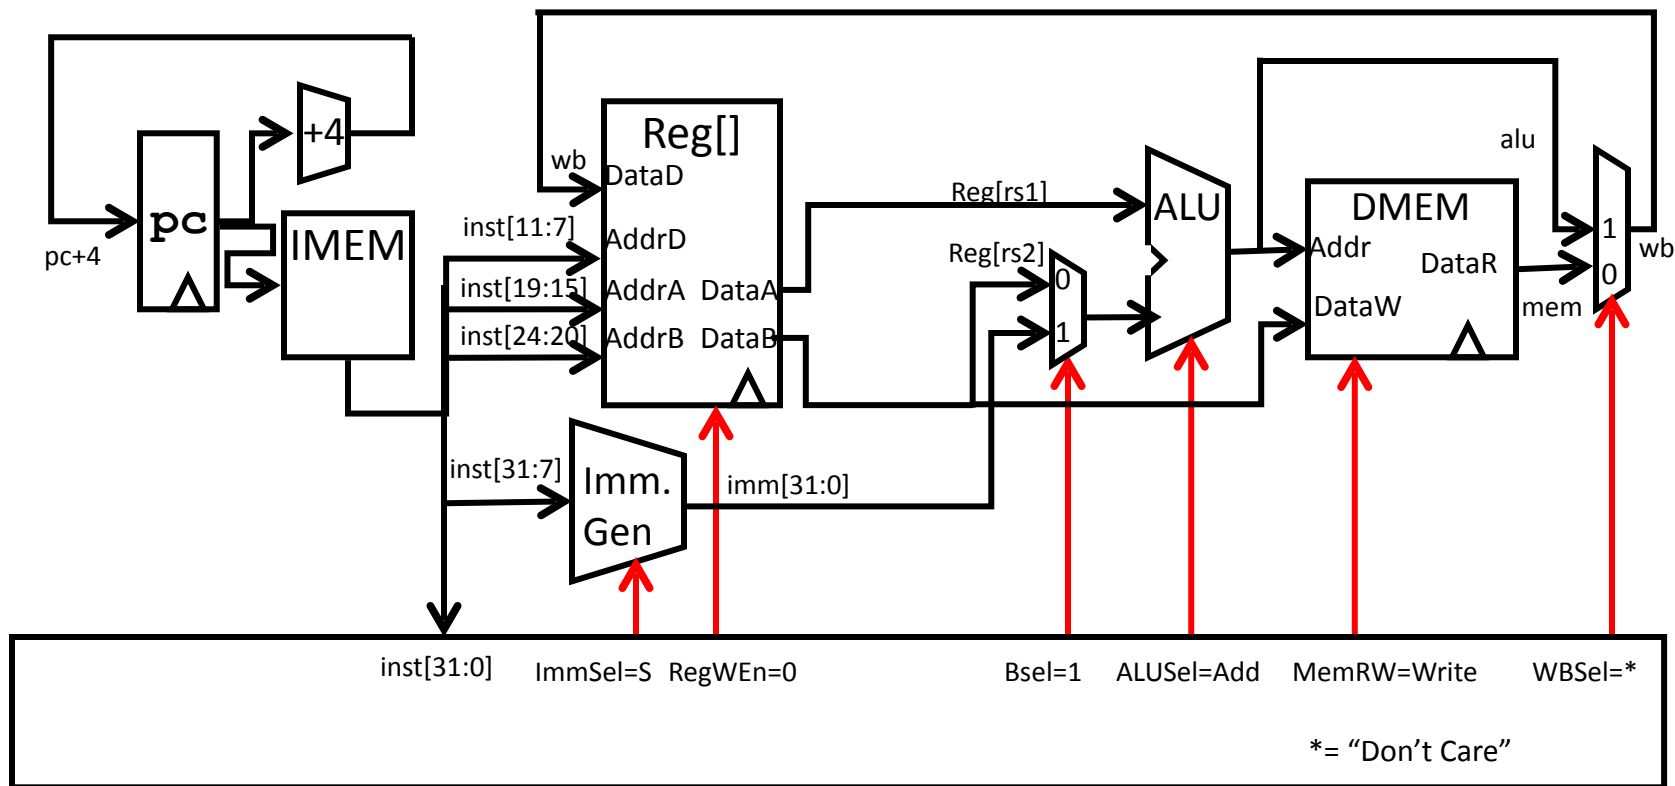
combined 12-bit offset = 8



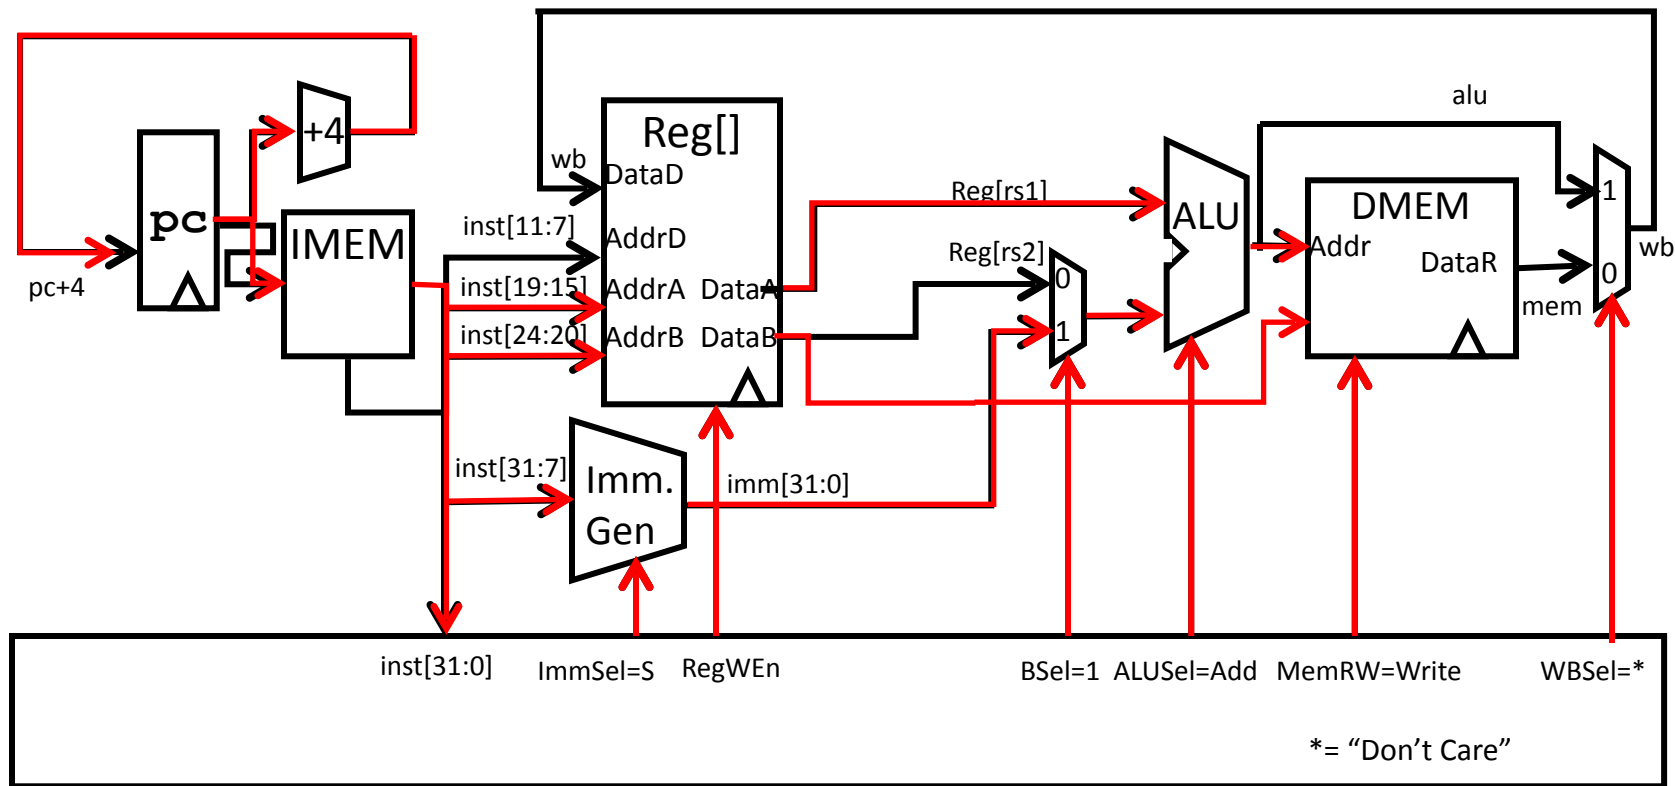
# Adding lw to datapath





# Adding **sw** to datapath



# Adding **sw** to datapath




# I-Format immediates




- High 12 bits of instruction (**inst[31:20]**) copied to low 12 bits of immediate (**imm[11:0]**)
- Immediate is sign-extended by copying value of **inst[31]** to fill the upper 20 bits of the immediate value (**imm[31:12]**)

# I & S Immediate Generator



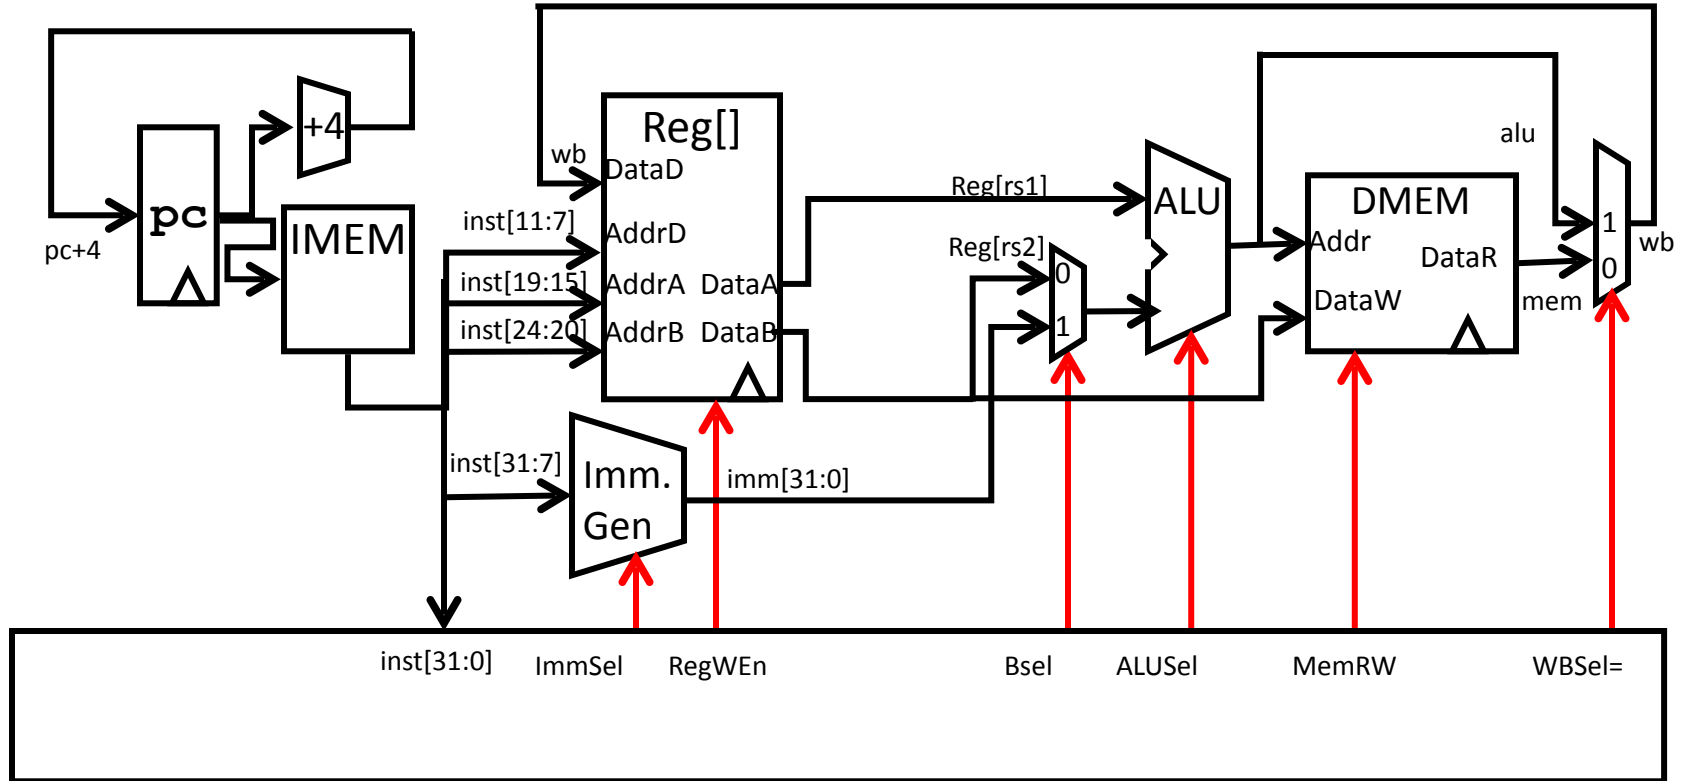
- Just need a 5-bit mux to select between two positions where low five bits of immediate can reside in instruction
- Other bits in immediate are wired to fixed positions in instruction

# Implementing Branches

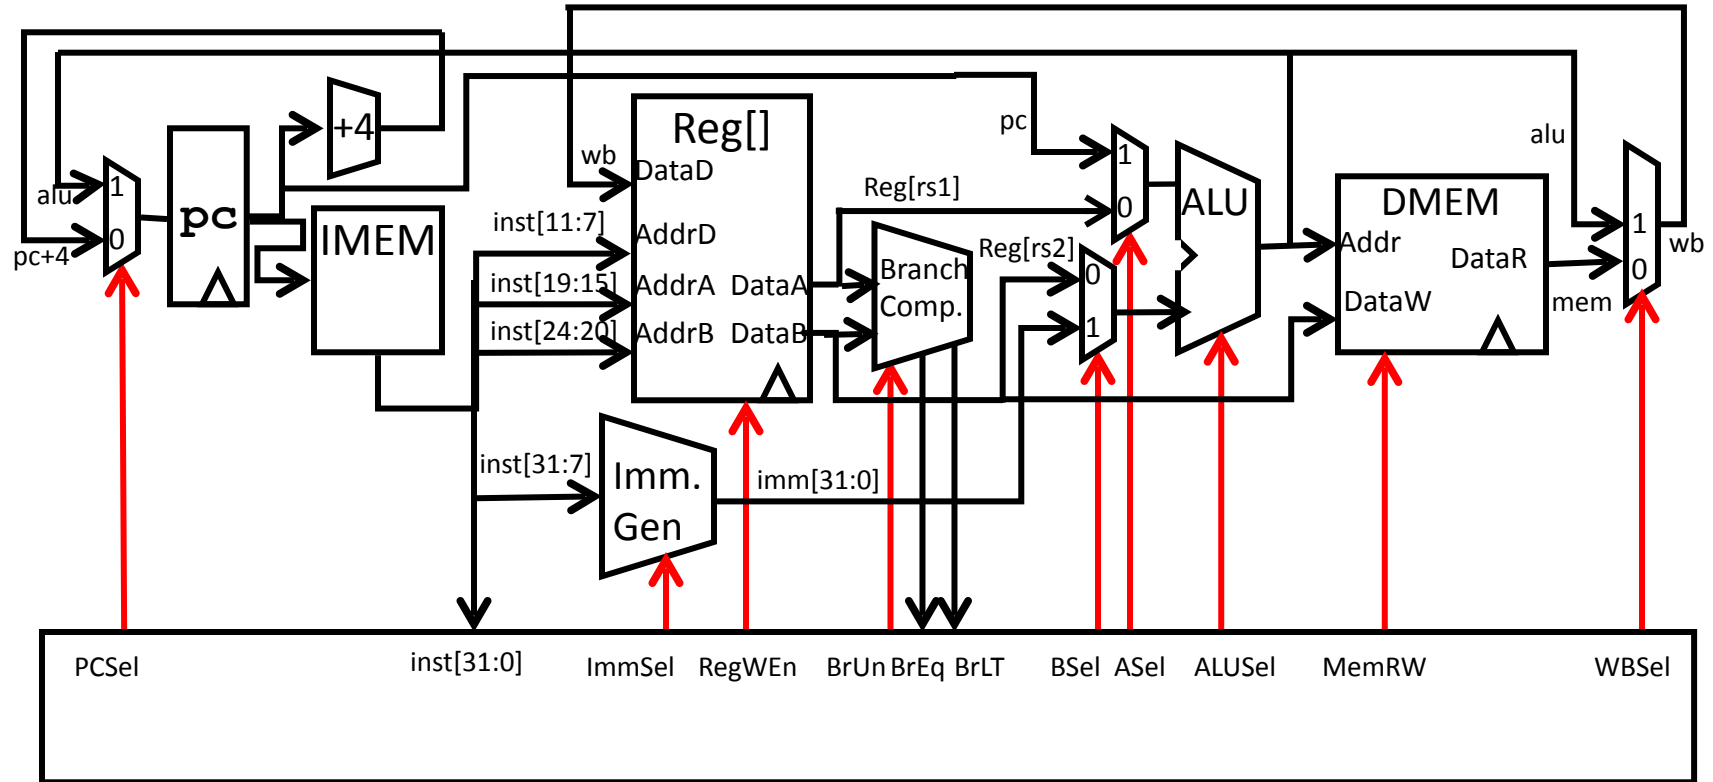


- B-format is mostly same as S-Format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode *even* 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

# Adding **sw** to datapath

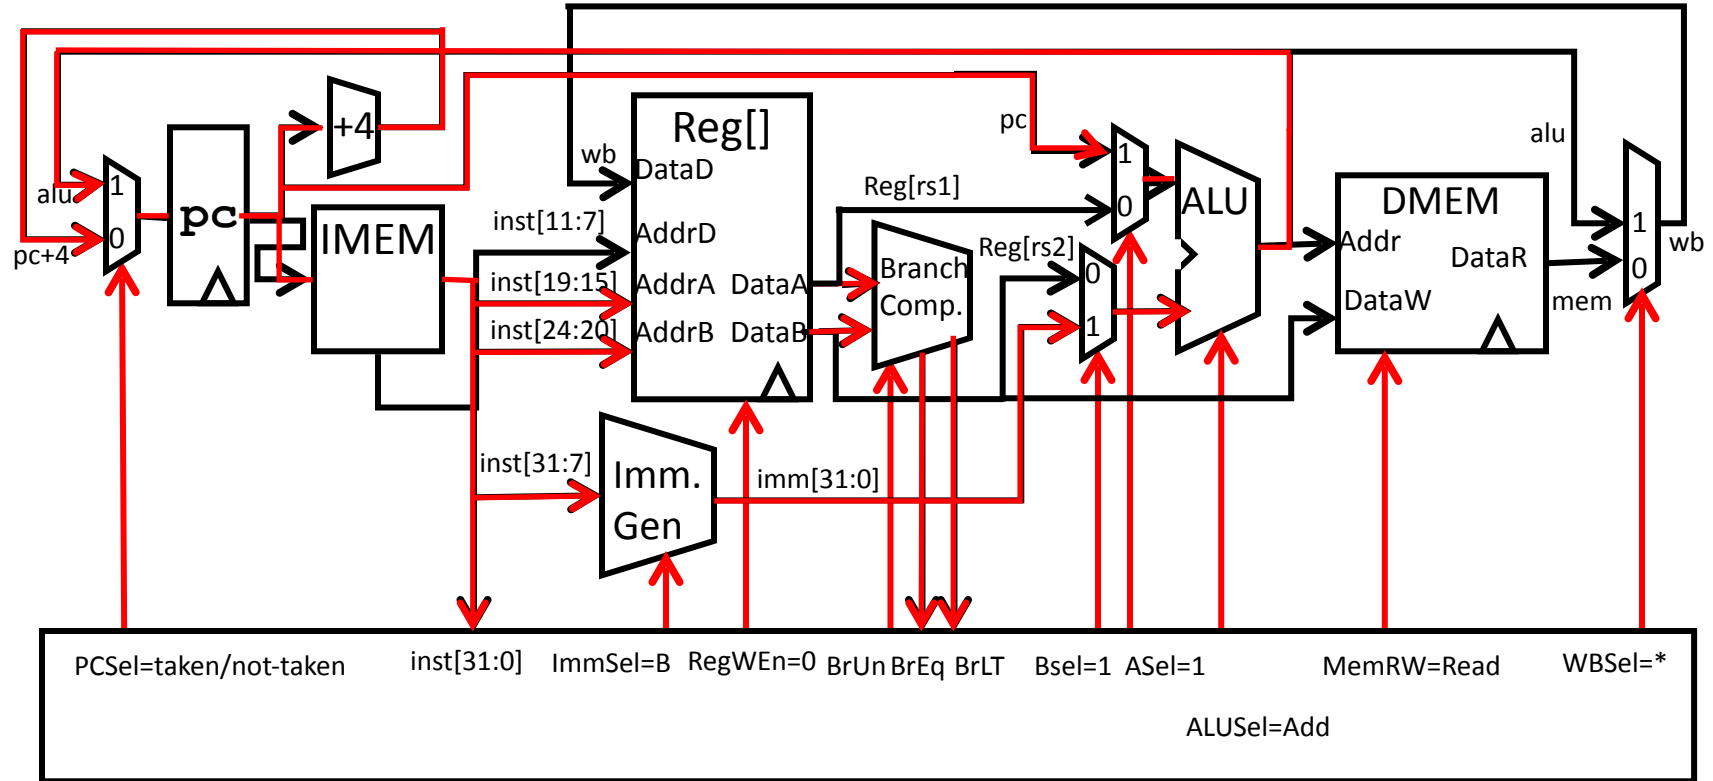


# Adding branches to datapath






# Adding branches to datapath




# Branch Comparator



- $\text{BrEq} = 1$ , if  $A=B$
- $\text{BrLT} = 1$ , if  $A < B$
- $\text{BrUn} = 1$  selects unsigned comparison for  $\text{BrLT}$ , 0=signed
- BGE branch:  $A \geq B$ , if  $\neg(A < B)$

# Administrivia (1/2)

- Midterm 1 has been regraded
- Grades Available on glookup




# Administrivia (2/2)

- Project 2.2 due Friday at 11:59pm
- Project 3.1 to be released next Wednesday (2/28)
  - RISC-V implementation in “logisim”
  - Due Tuesday
- Homework 2 is released and due next Friday at 11:59pm
- No Guerrilla Session this week—will start up again next Thursday 3/1

# Break!

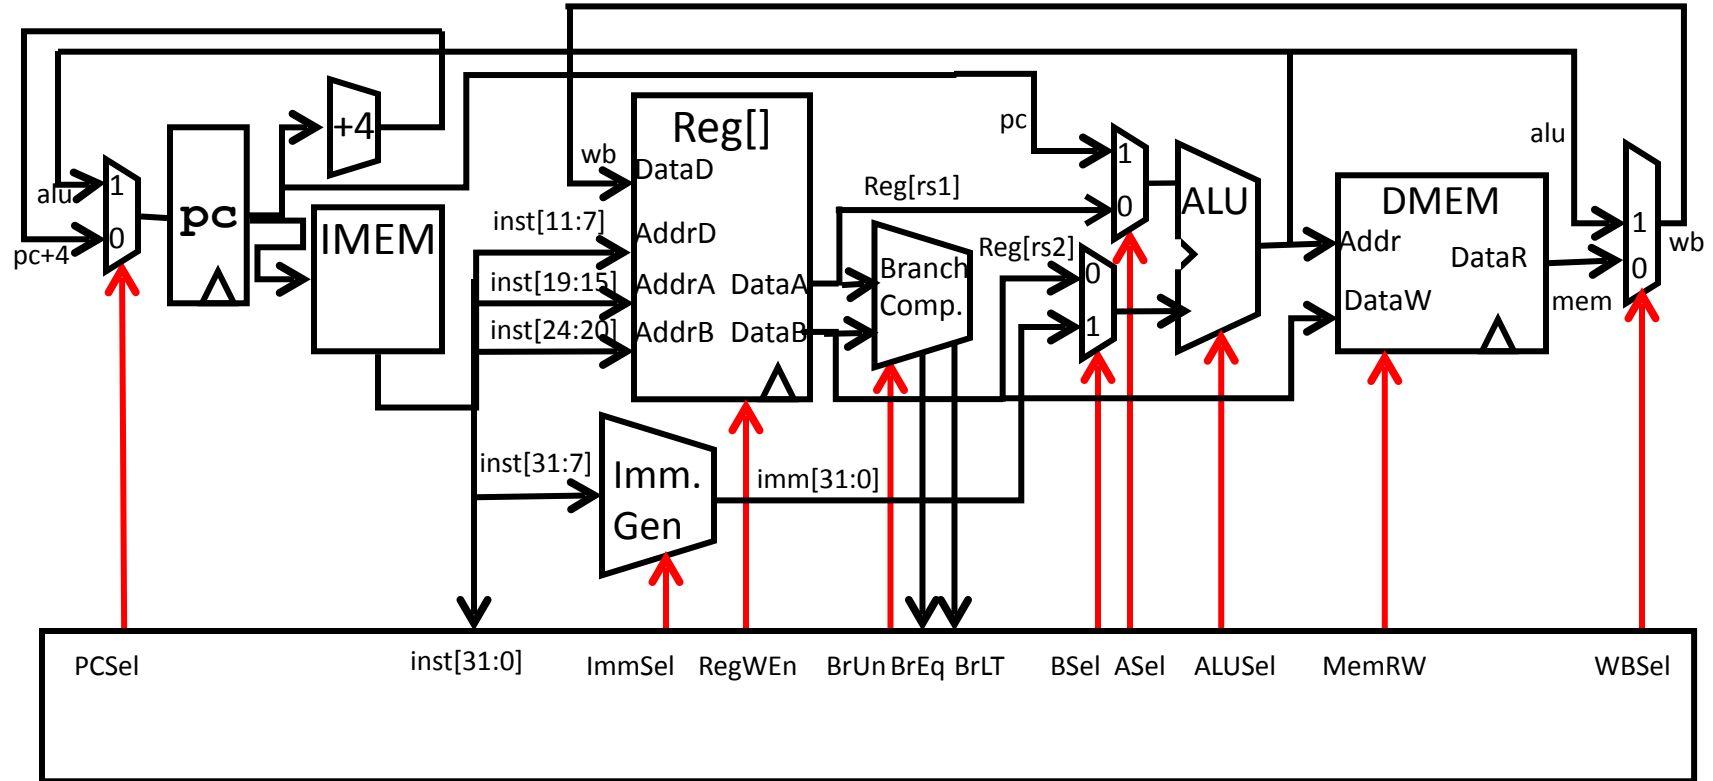


# Implementing **JALR** Instruction (I-Format)




- JALR rd, rs, immediate
  - Writes PC+4 to Reg[rd] (return address)
  - Sets PC = Reg[rs1] + immediate
  - Uses same immediates as arithmetic and loads
    - **no** multiplication by 2 bytes

# Adding branches to datapath

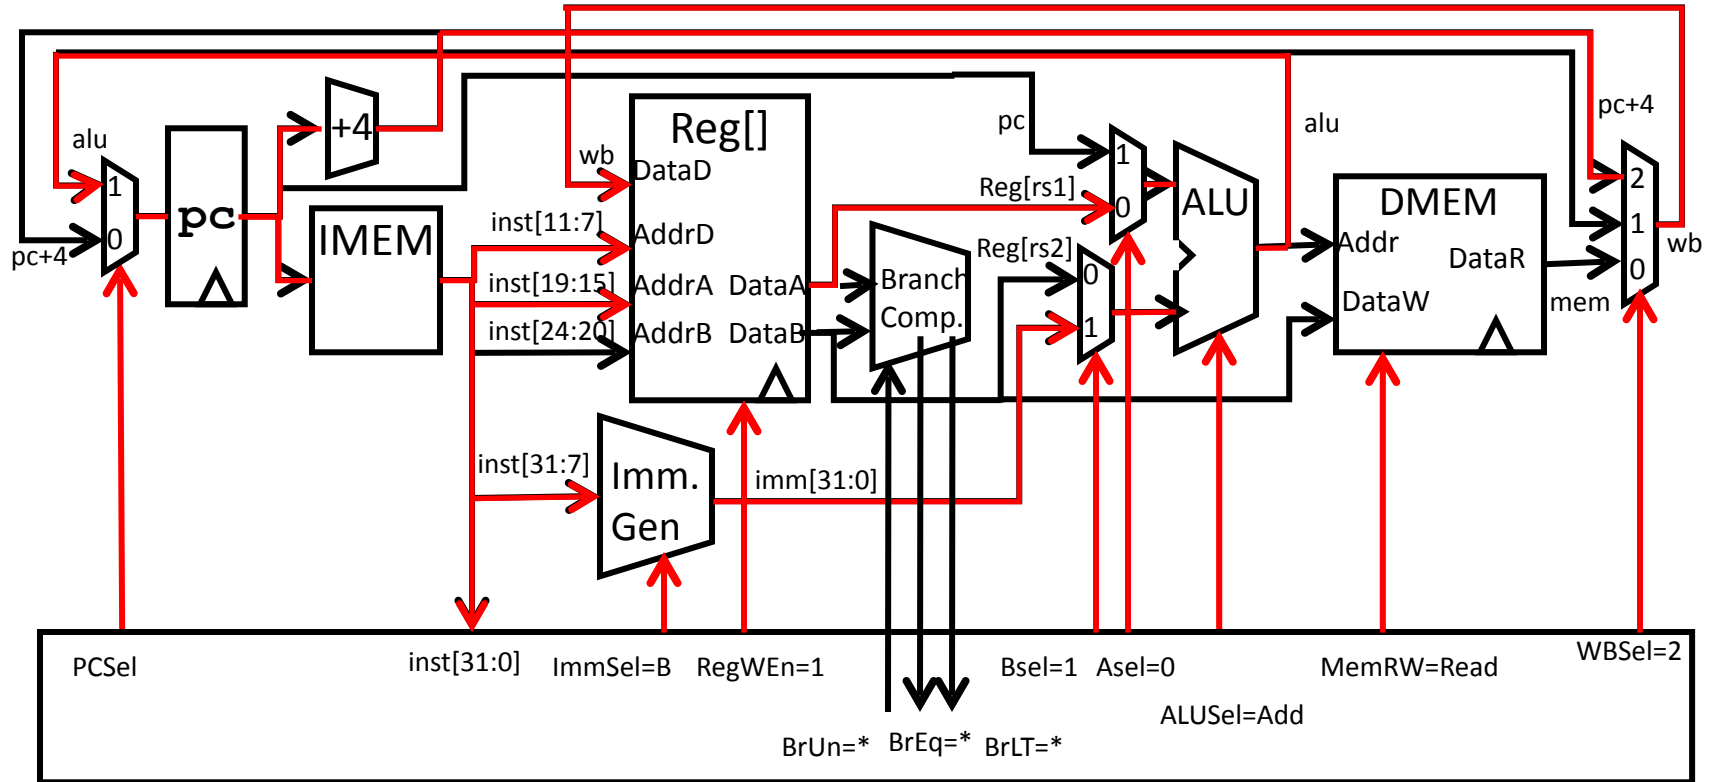


# Adding jalr to datapath






# Adding jalr to datapath

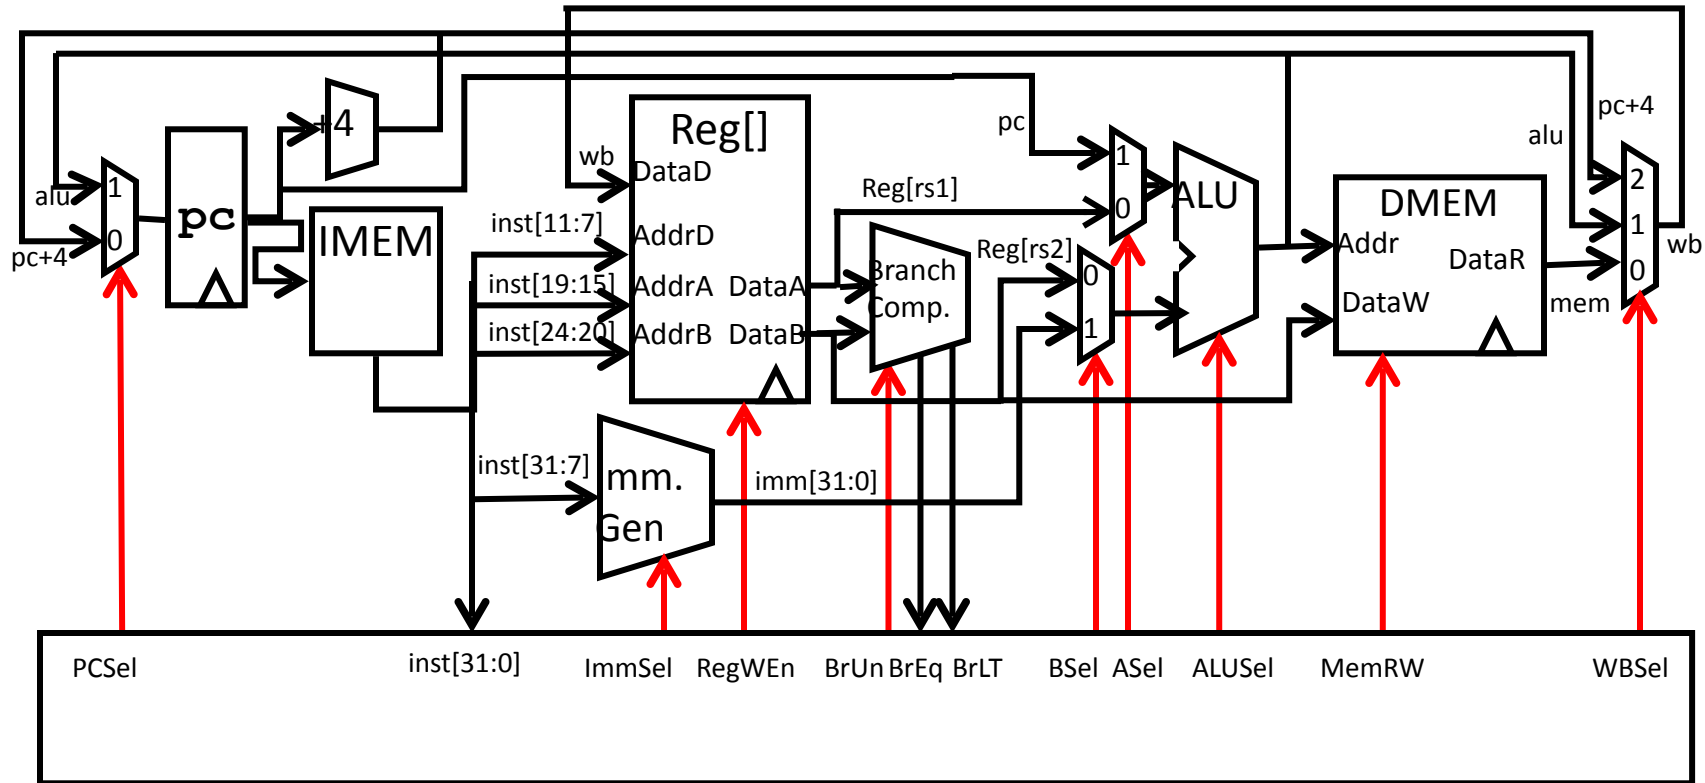


# Implementing jal Instruction

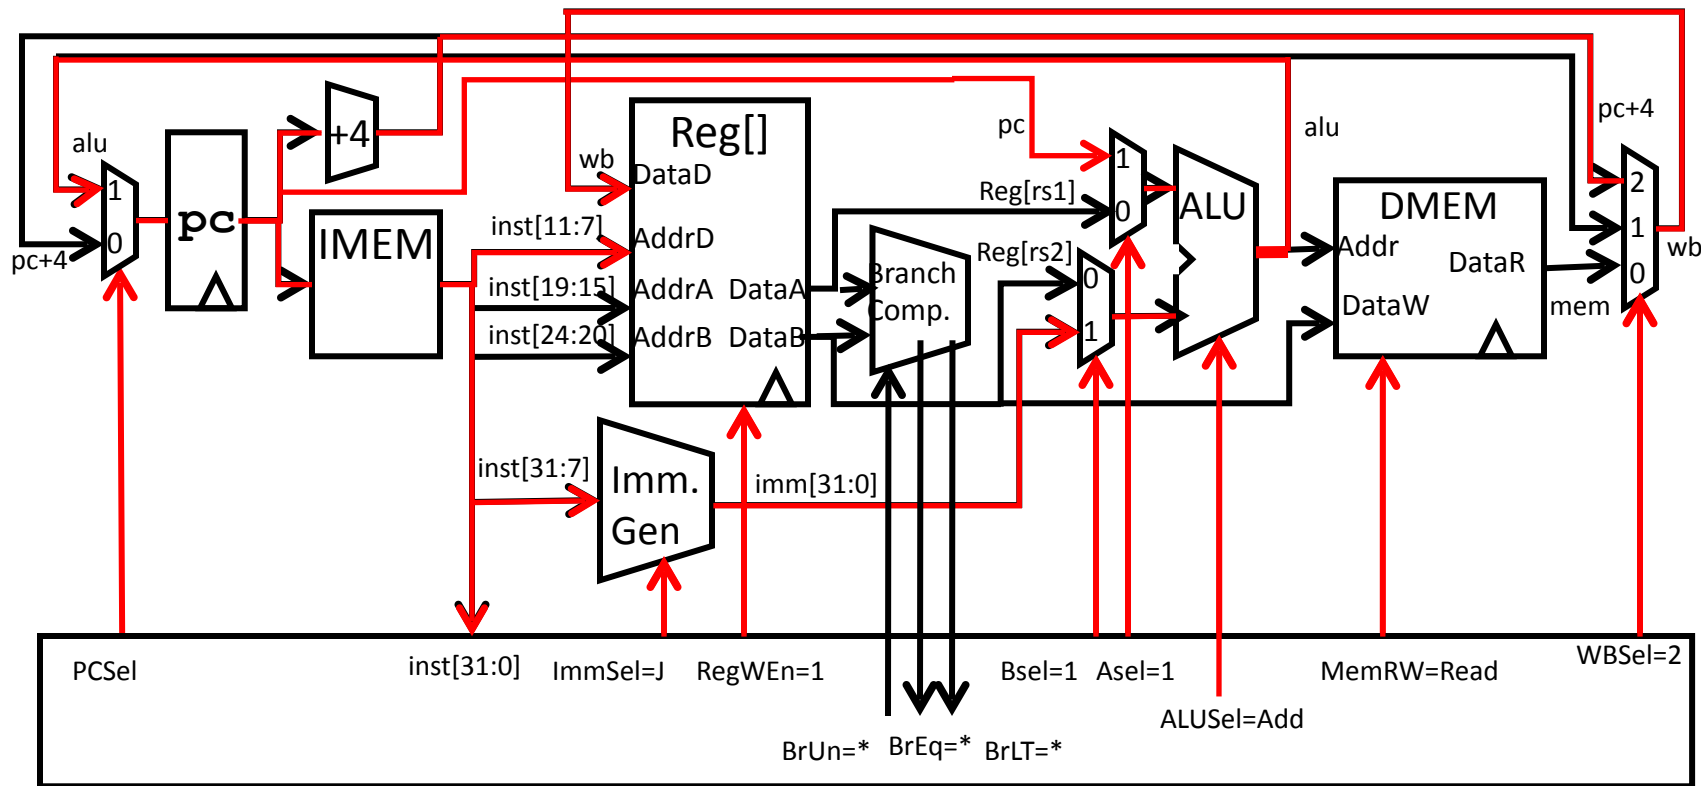


- JAL saves PC+4 in Reg[rd] (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within  $\pm 2^{19}$  locations, 2 bytes apart
  - $\pm 2^{18}$  32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

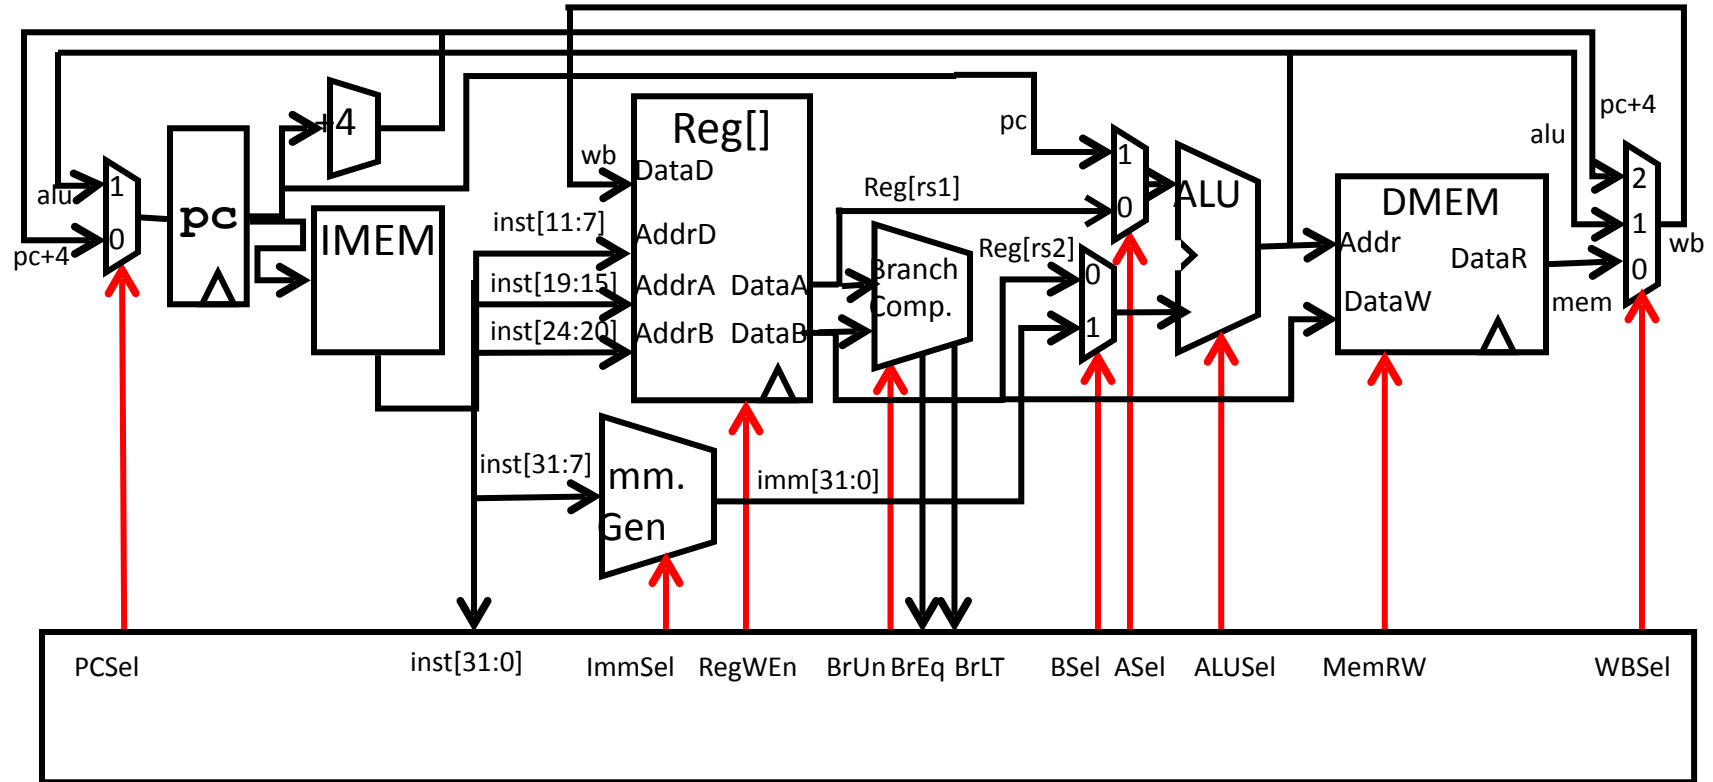
# Adding jal to datapath



# Adding jal to datapath



# Single-Cycle RISC-V RV32I Datapath



# And in Conclusion, ...

- Universal datapath
  - Capable of executing all RISC-V instructions in one cycle each
  - Not all units (hardware) used by all instructions
- 5 Phases of execution
  - IF, ID, EX, MEM, WB
  - Not all instructions are active in all phases
- Controller specifies how to execute instructions
  - what new instructions can be added with just most control?