

CS 61C:
Great Ideas in Computer Architecture

Lecture 13: *Pipelining*

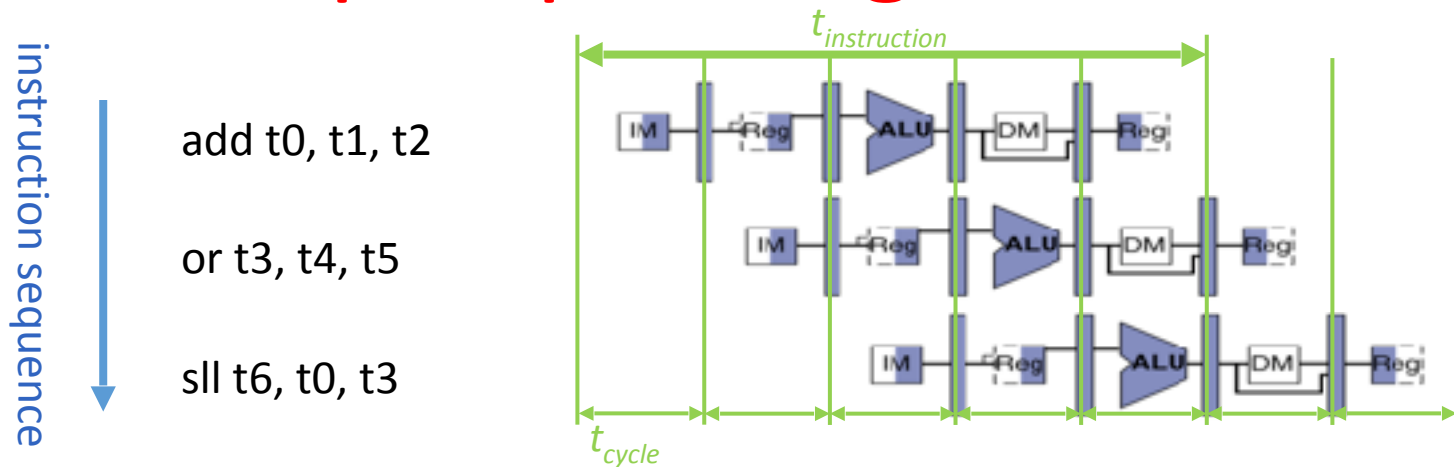
John Wawrzynek & Nick Weaver

<http://inst.eecs.berkeley.edu/~cs61c/sp18>

Agenda

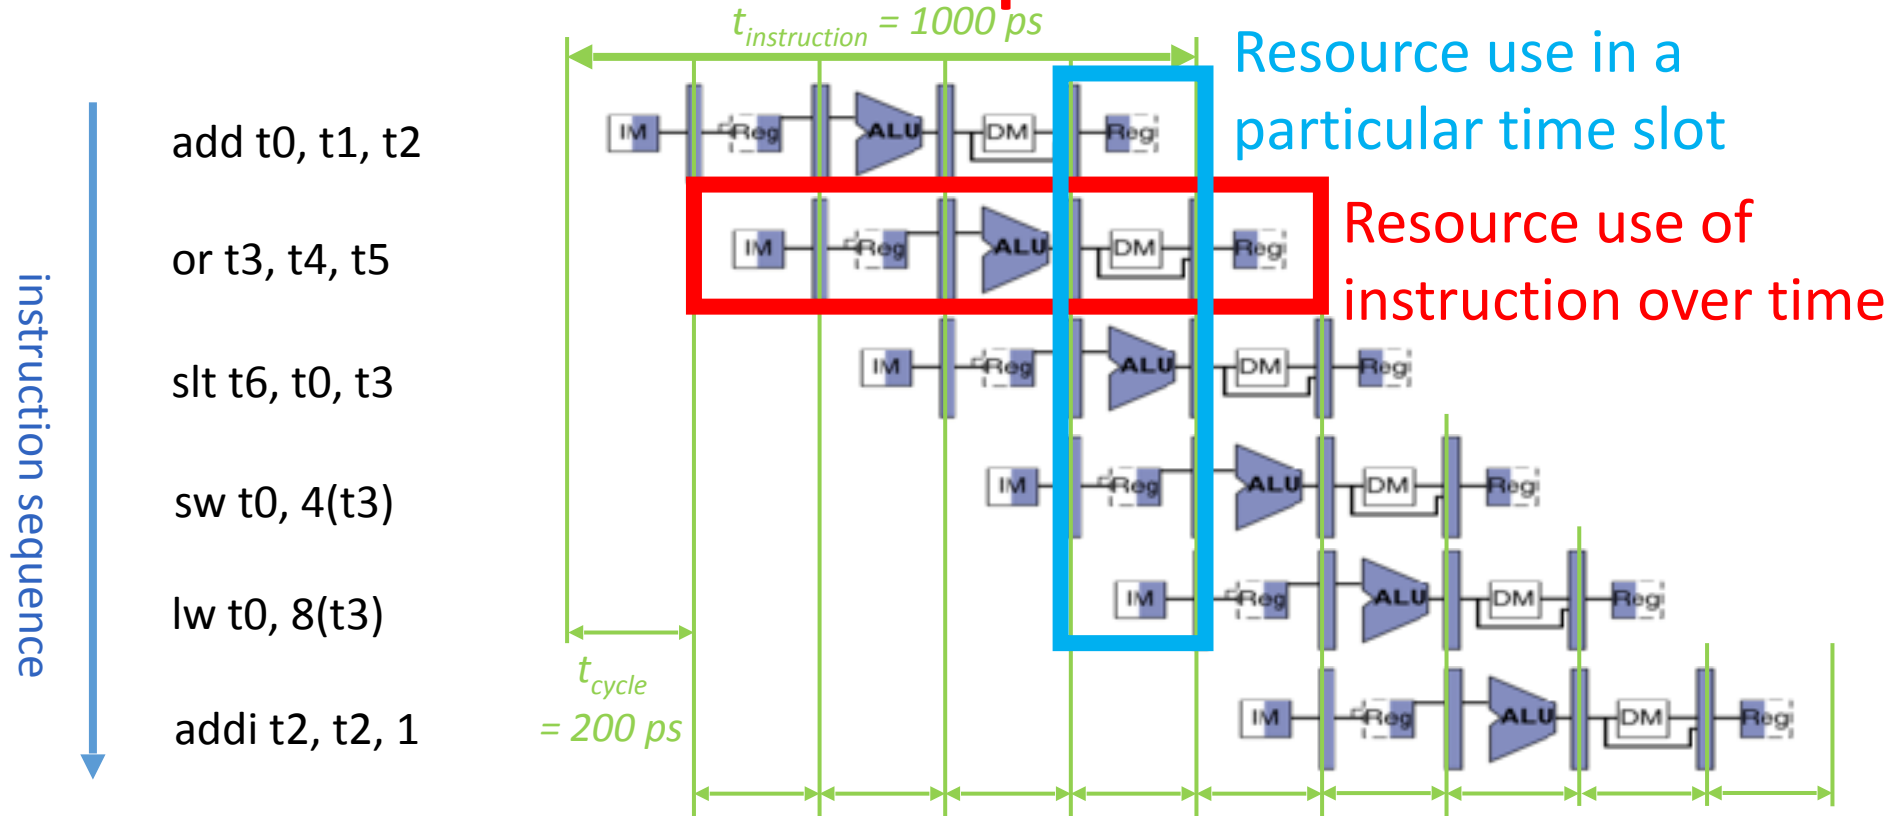
- **RISC-V Pipeline**
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Recap: Pipelining with RISC-V

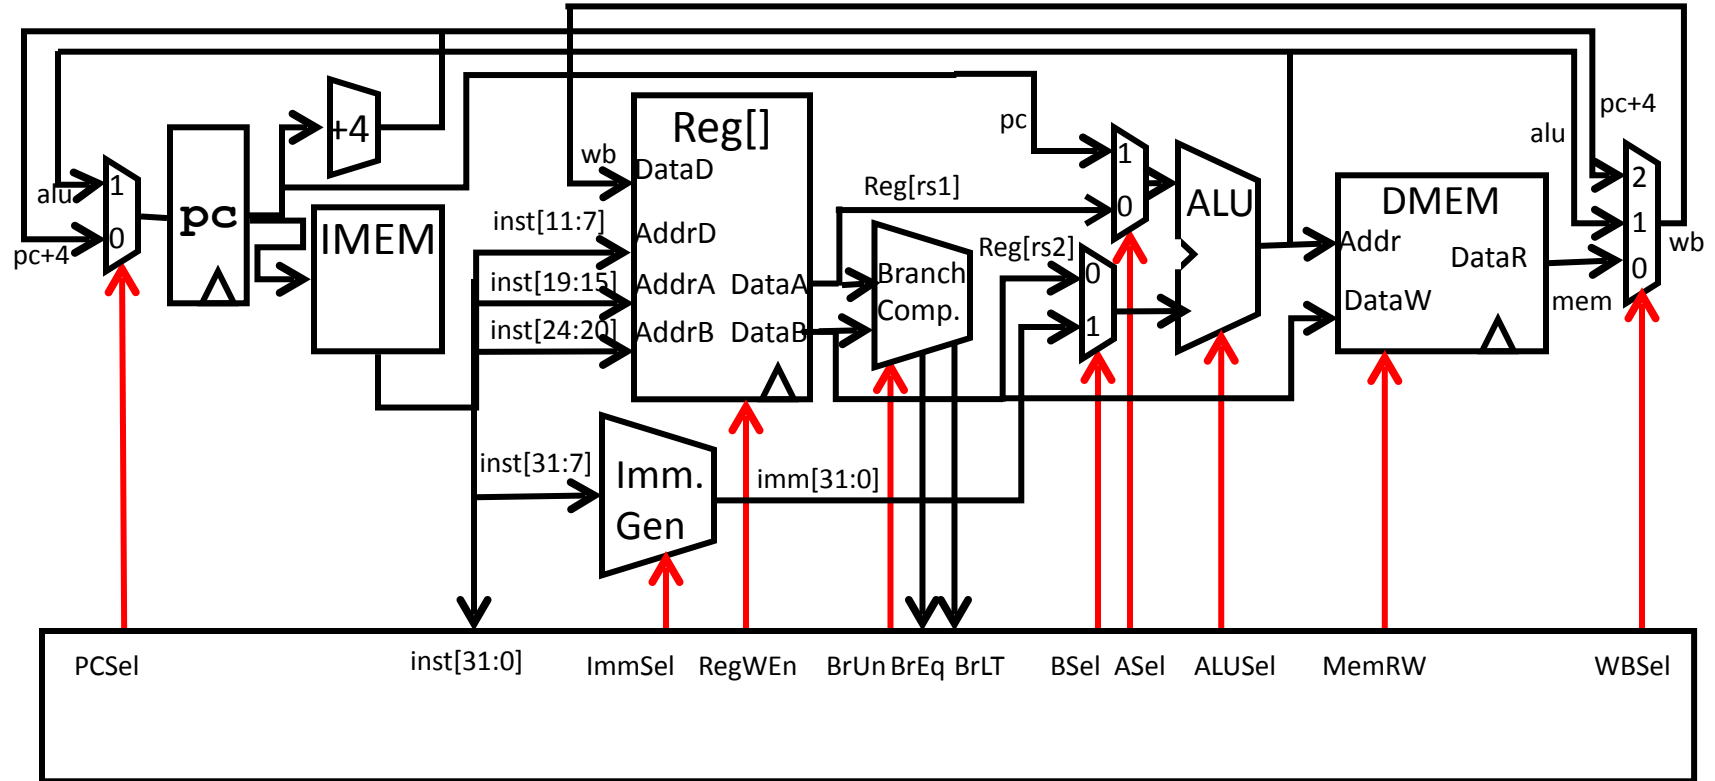


	Single Cycle	Pipelining
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, f_s	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Relative speed	1 x	4 x

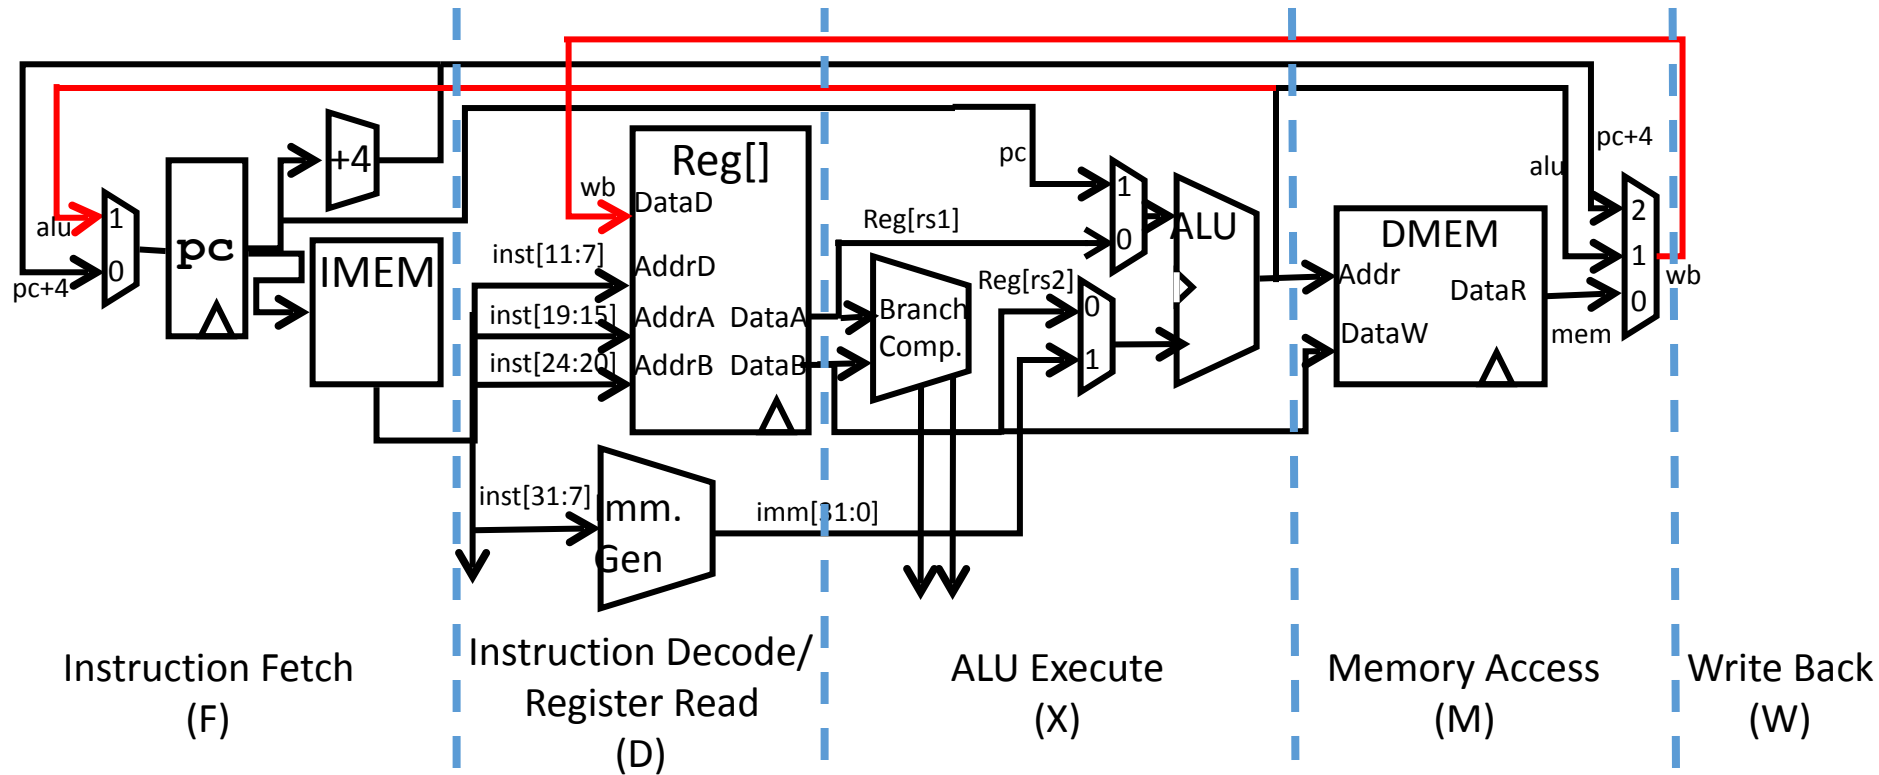
RISC-V Pipeline



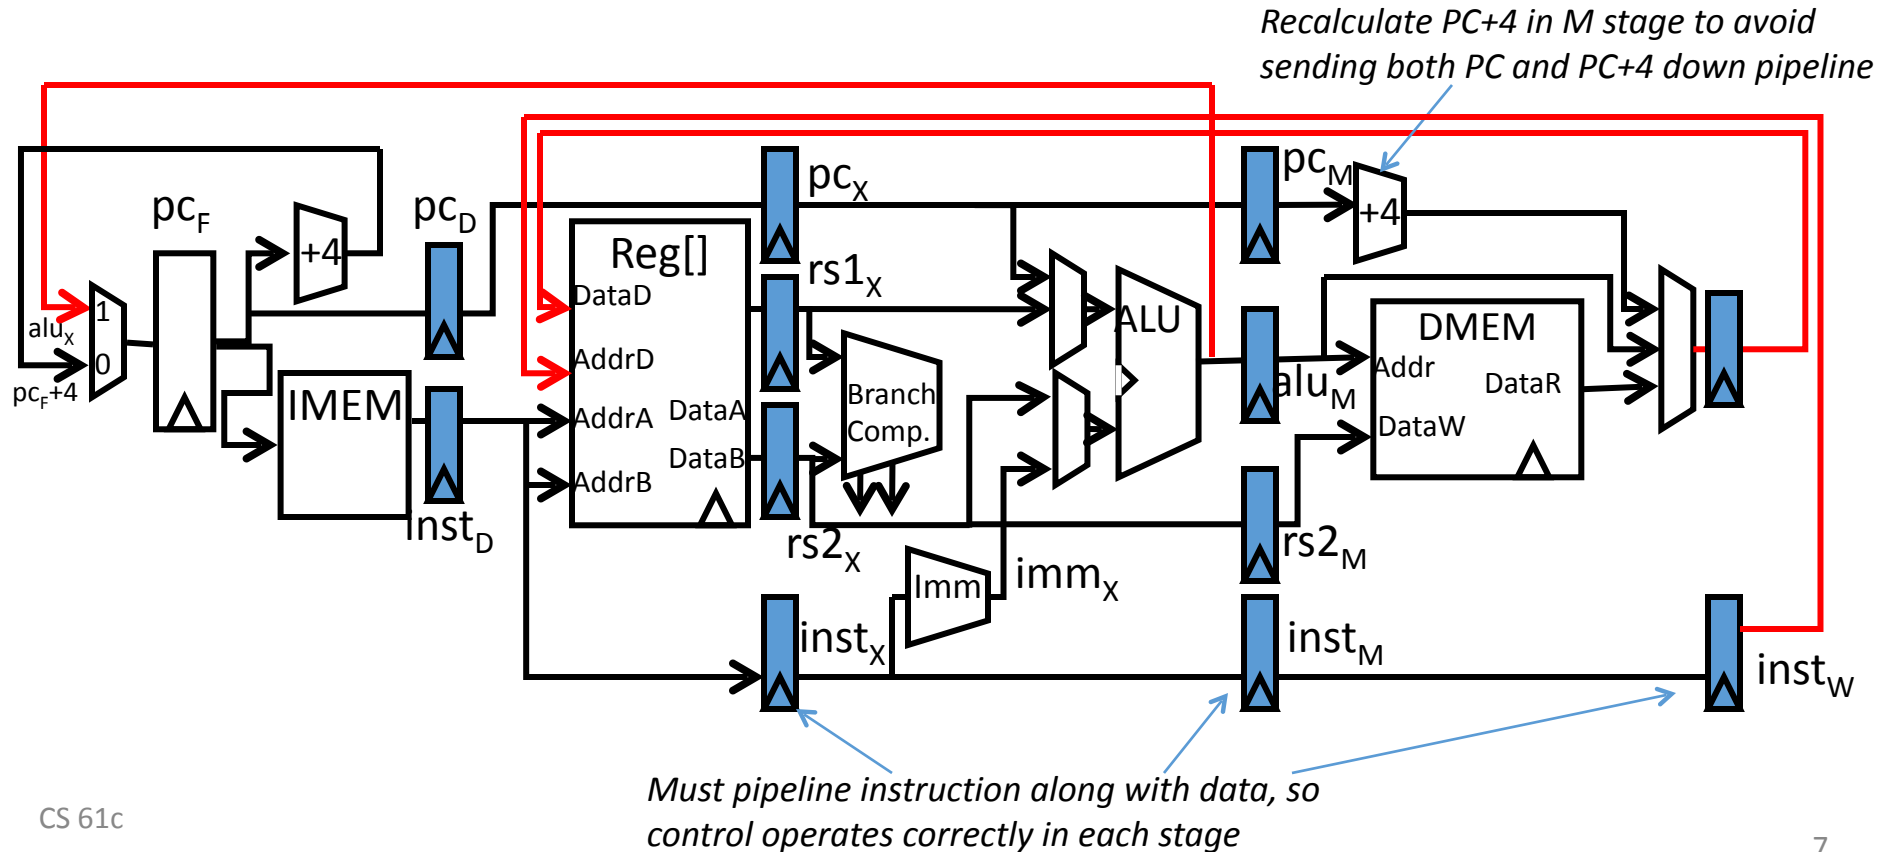
Single-Cycle RISC-V RV32I Datapath



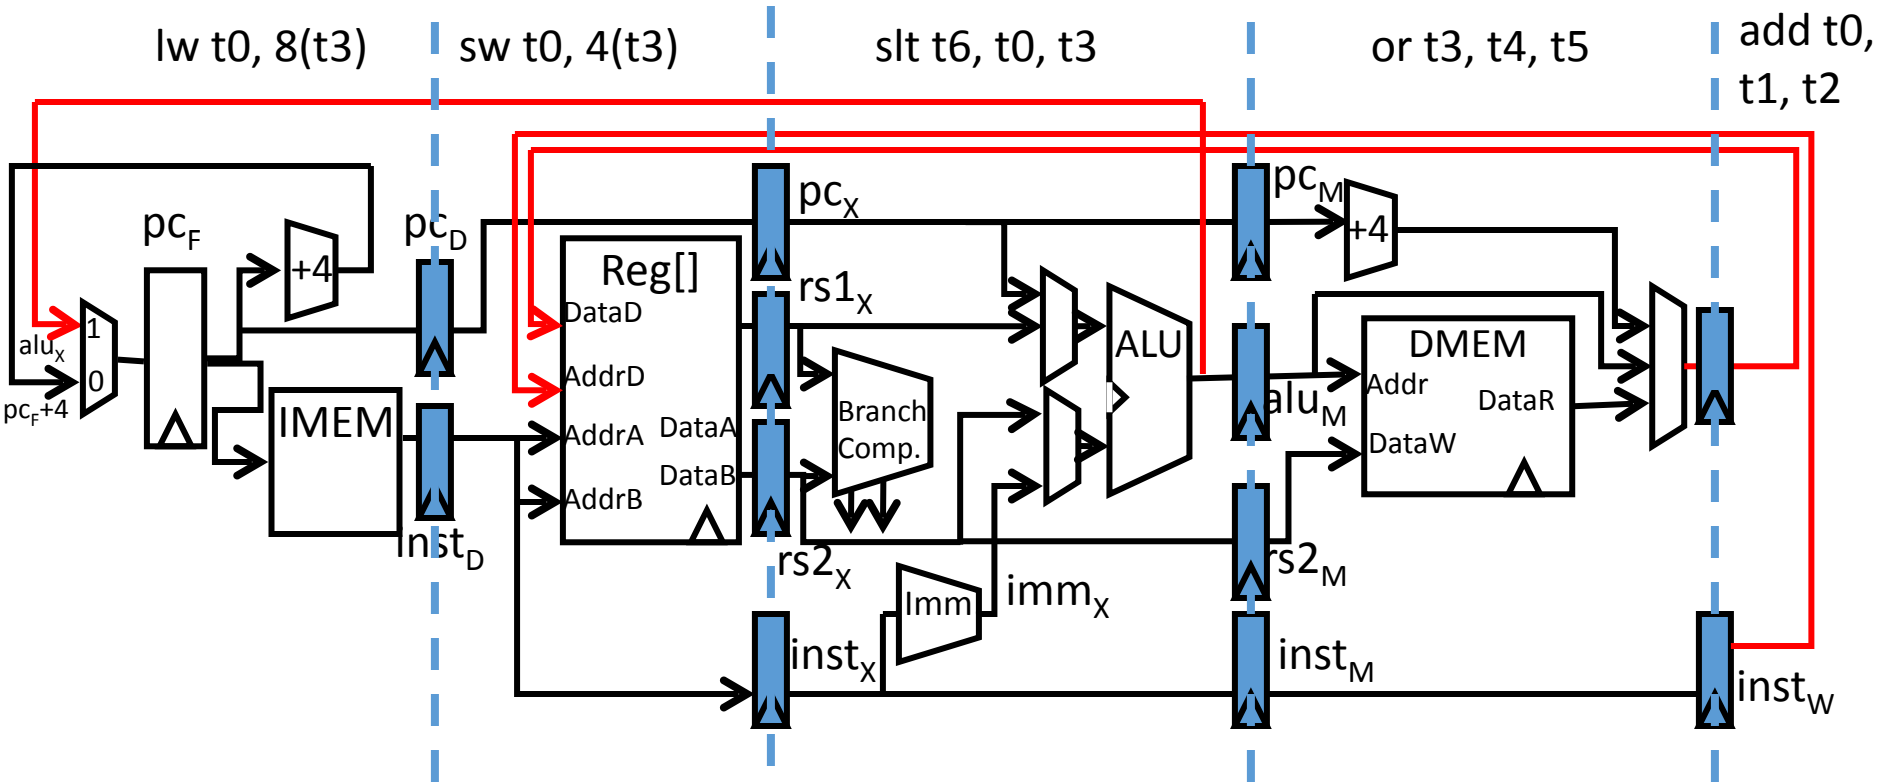
Pipelining RISC-V RV32I Datapath



Pipelined RISC-V RV32I Datapath



Each stage operates on different instruction



Clicker Question

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

Pipelining the single-cycle processor can increase processor performance by:


	Instructions /program	Cycles/ instruction	Time/cycle
A	decrease	decrease	same
B	same	increase	decrease
C	same	same	decrease
D	increase	decrease	increase

Agenda

- RISC-V Pipeline
- **Pipeline Control**
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages



Hazards Ahead



Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - **Structural**
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Structural Hazard

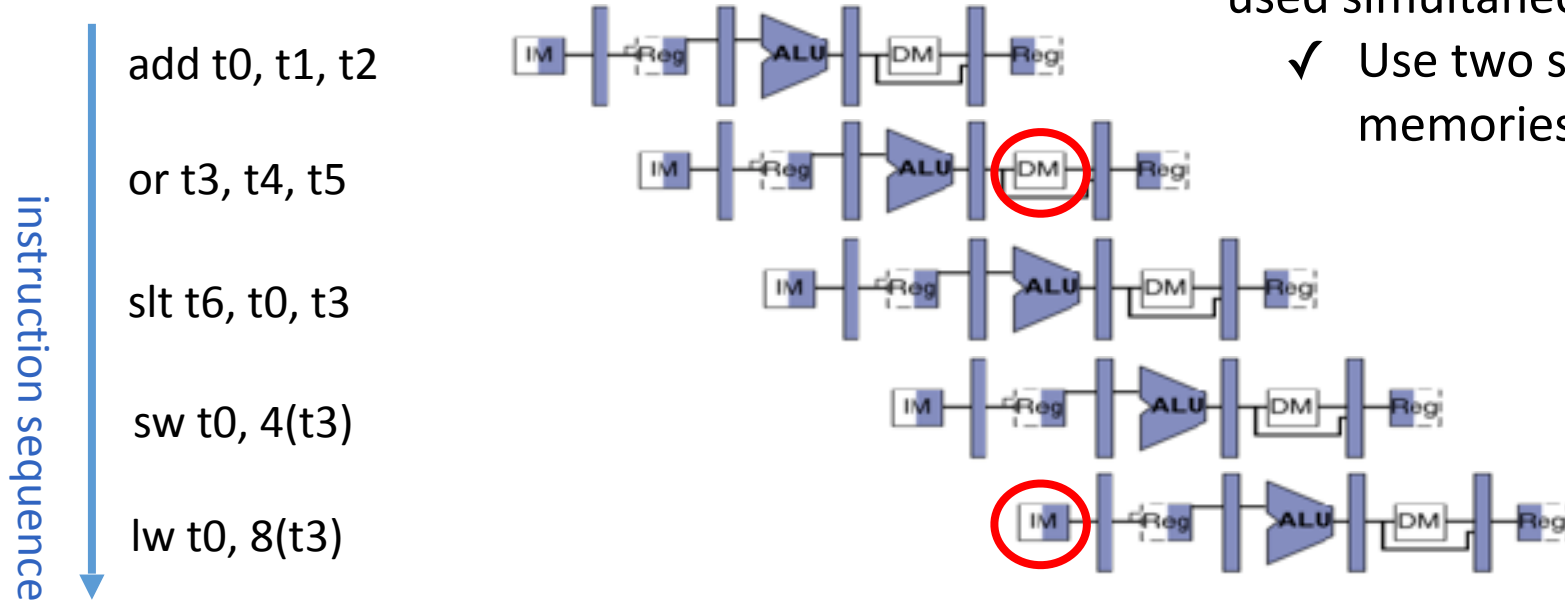
- **Problem:** Two or more instructions in the pipeline compete for access to a single physical resource
- **Solution 1:** Instructions take turns to use resource, some instructions have to stall
- **Solution 2:** Add more hardware to machine
- Can always solve a structural hazard by adding more hardware

Regfile Structural Hazards


- Each instruction:
 - can read up to two operands in decode stage
 - can write one value in writeback stage
- Avoid structural hazard by having separate “ports”
 - two independent read ports and one independent write port
- Three accesses per cycle can happen simultaneously

Structural Hazard: Memory Access

- Instruction and data memory used simultaneously
✓ Use two separate memories



Instruction and Data Caches



Caches: small and fast “buffer” memories

Structural Hazards – Summary

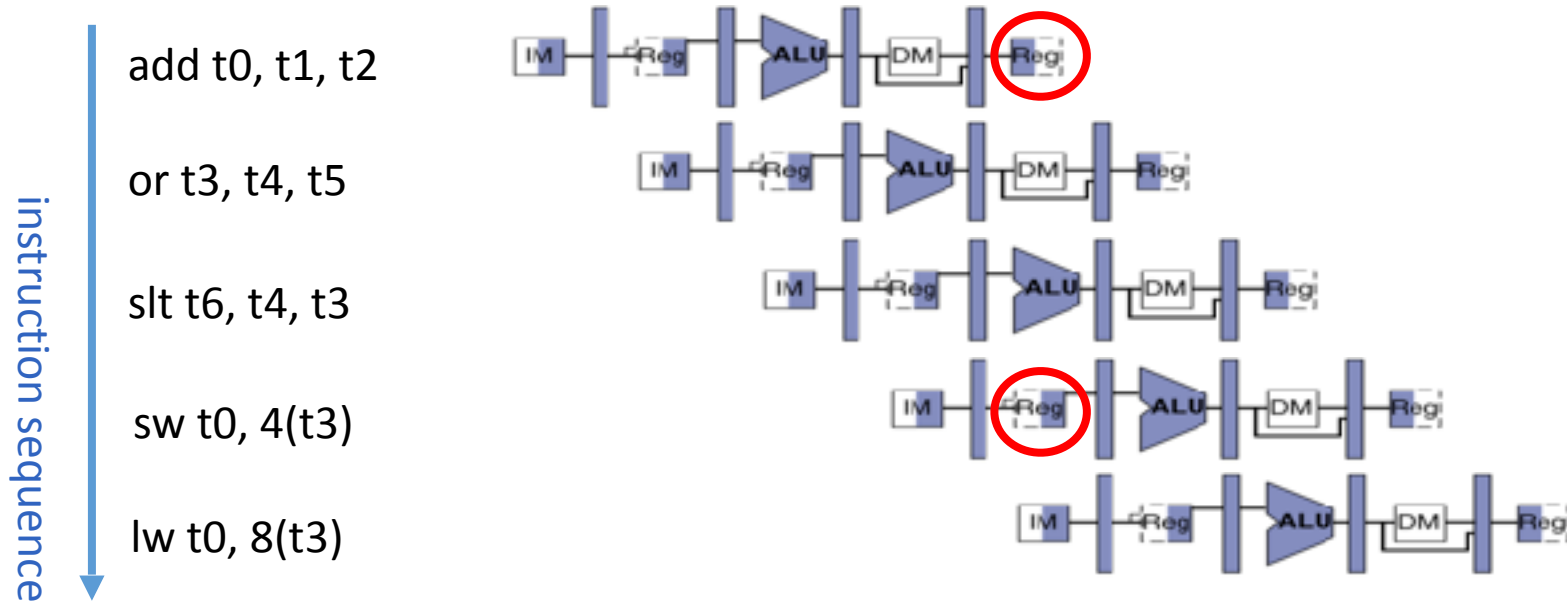
- Conflict for use of a resource
- In RISC-V pipeline with a single memory
 - Load/store requires data access
 - Without separate memories, instruction fetch would have to *stall* for that cycle
 - All other operations in pipeline would have to wait
- Pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches
- RISC ISAs (including RISC-V) designed to avoid structural hazards
 - e.g. at most one memory access/instruction

Agenda

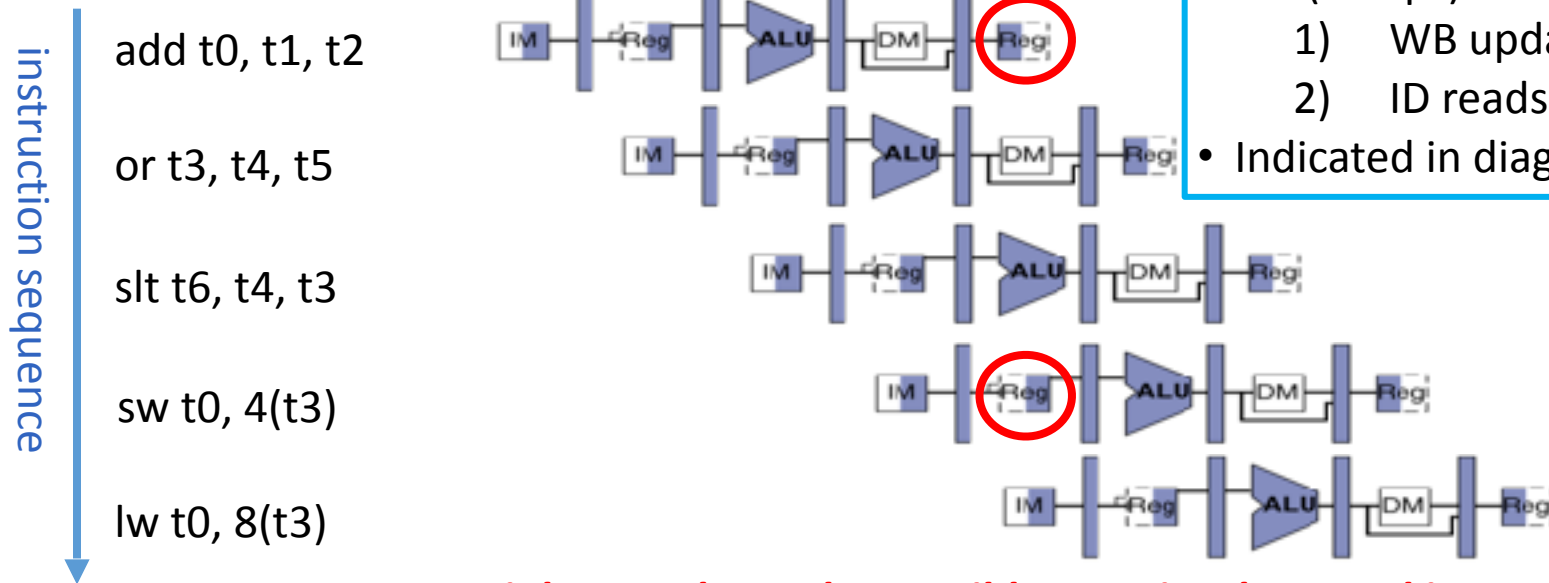
- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - **Data**
 - **R-type instructions**
 - Load
 - Control
- Superscalar processors

Data Hazard: Register Access

- Separate ports, but what if write to same value as read?
- Does **sw** in the example fetch the old or new value?



Register Access Policy




- Exploit high speed of register file (100 ps)
 - 1) WB updates value
 - 2) ID reads new value
- Indicated in diagram by shading

Might not always be possible to write then read in same cycle, especially in high-frequency designs.

Data Hazard: ALU Result

s0 holds "5" then add instr changes s0 to "9"

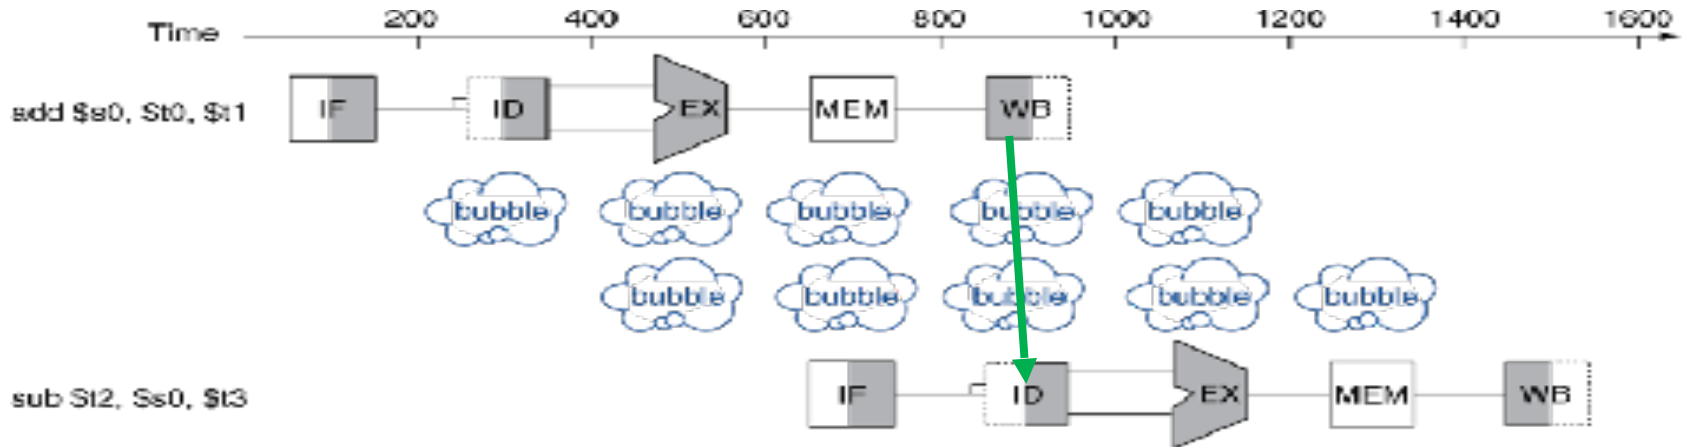


Without some fix, **sub** and **or** will calculate wrong result!

Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction

– add s0, t0, t1
 sub t2, s0, t3




- Bubble:
 - stall dependent instruction
 - effectively NOP: affected pipeline stages do “nothing”

Stalls and Performance

- Stalls reduce performance
 - But stalls are required to get correct results
- Compiler could try to arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

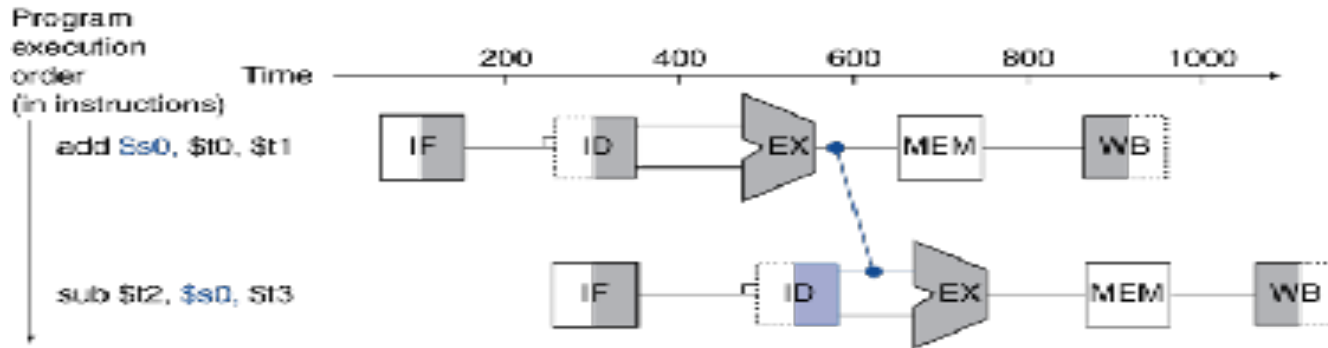
Solution 2: Forwarding



**Forwarding: grab operand from pipeline stage,
rather than register file**

Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath

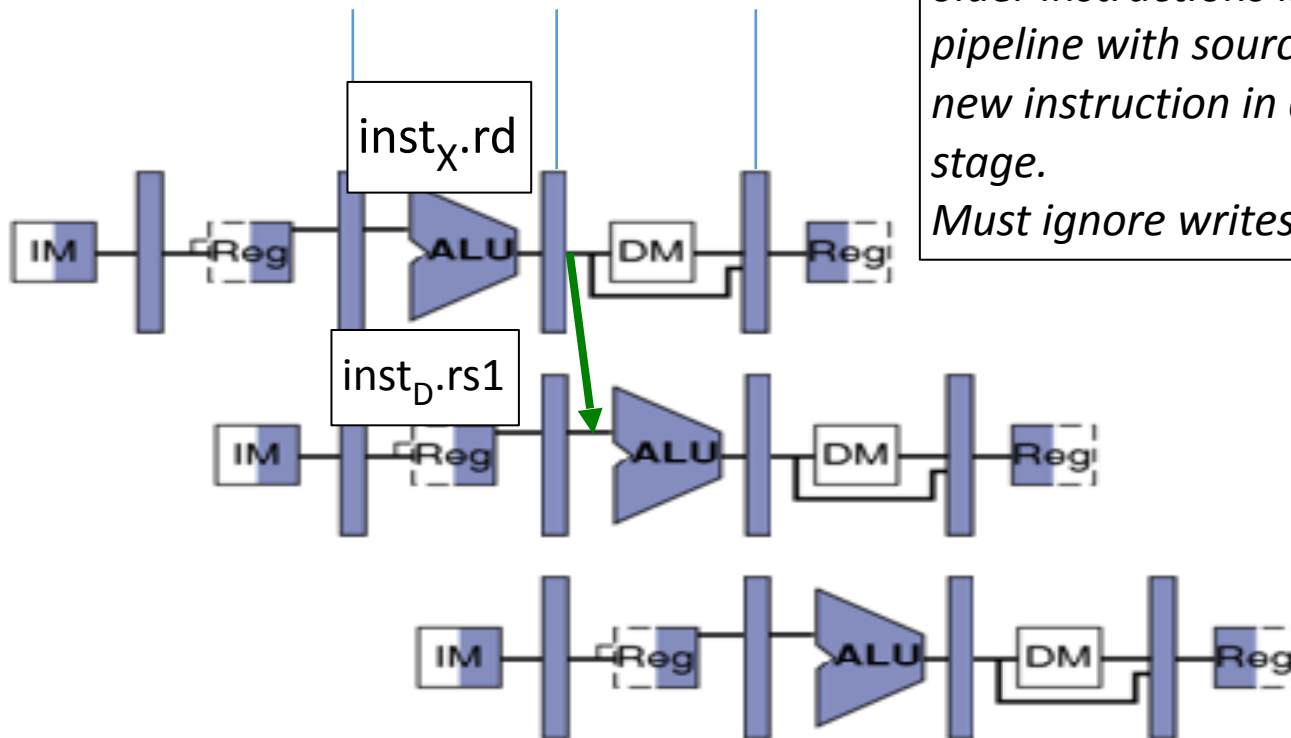


1) Detect Need for Forwarding (example)

add t0, t1, t2

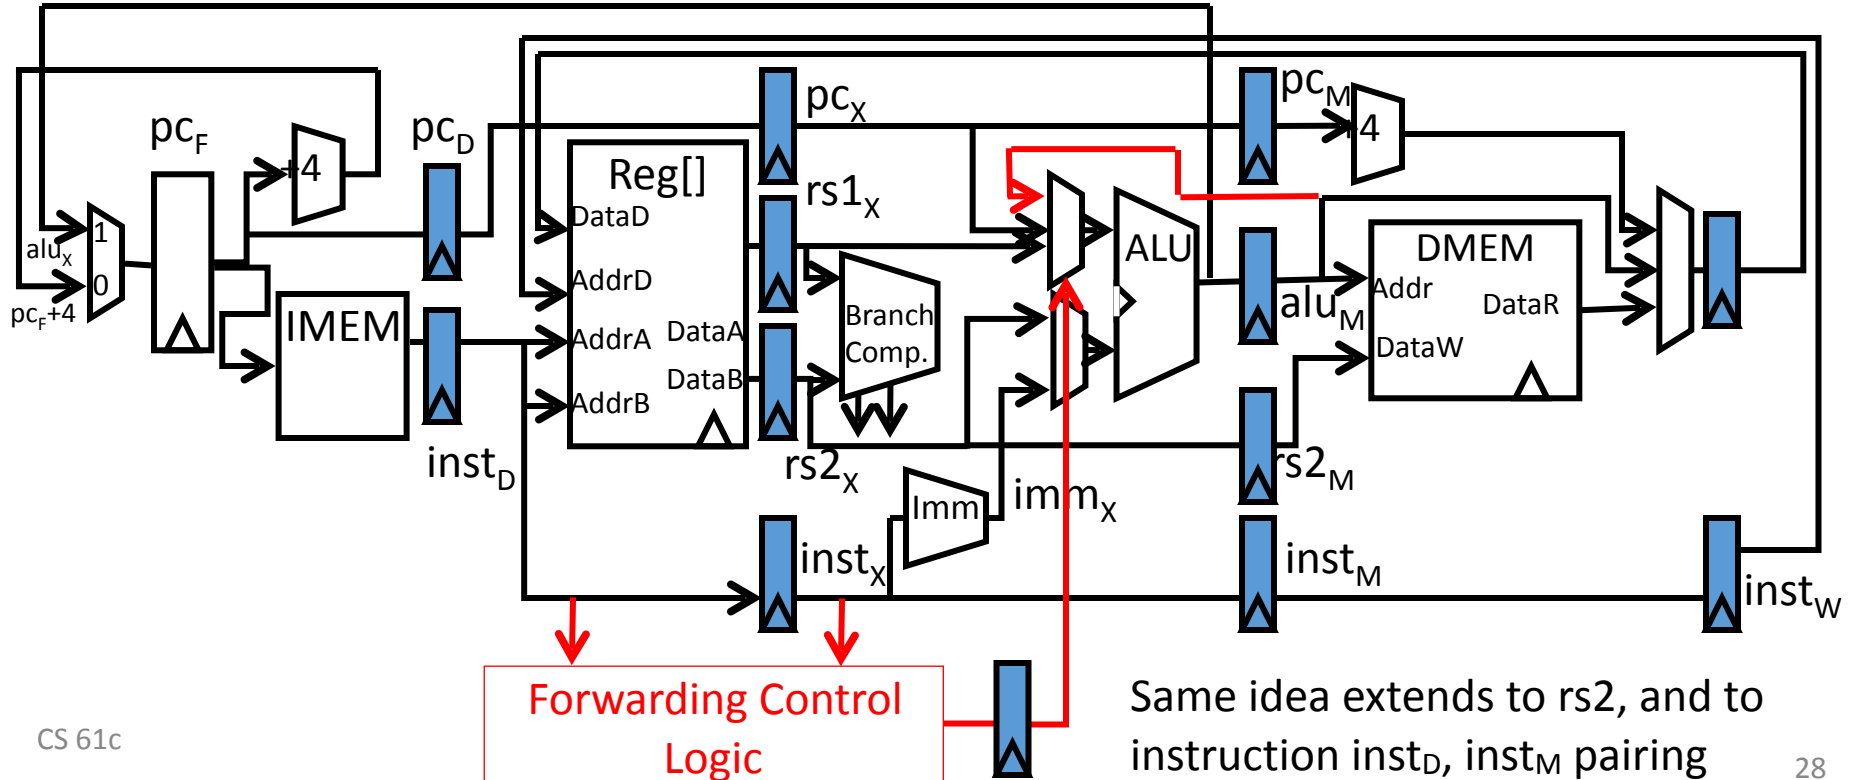
or t3, t0, t5

sub t6, t0, t3



Compare destination of older instructions in pipeline with sources of new instruction in decode stage.
Must ignore writes to x0!

Example Forwarding Path



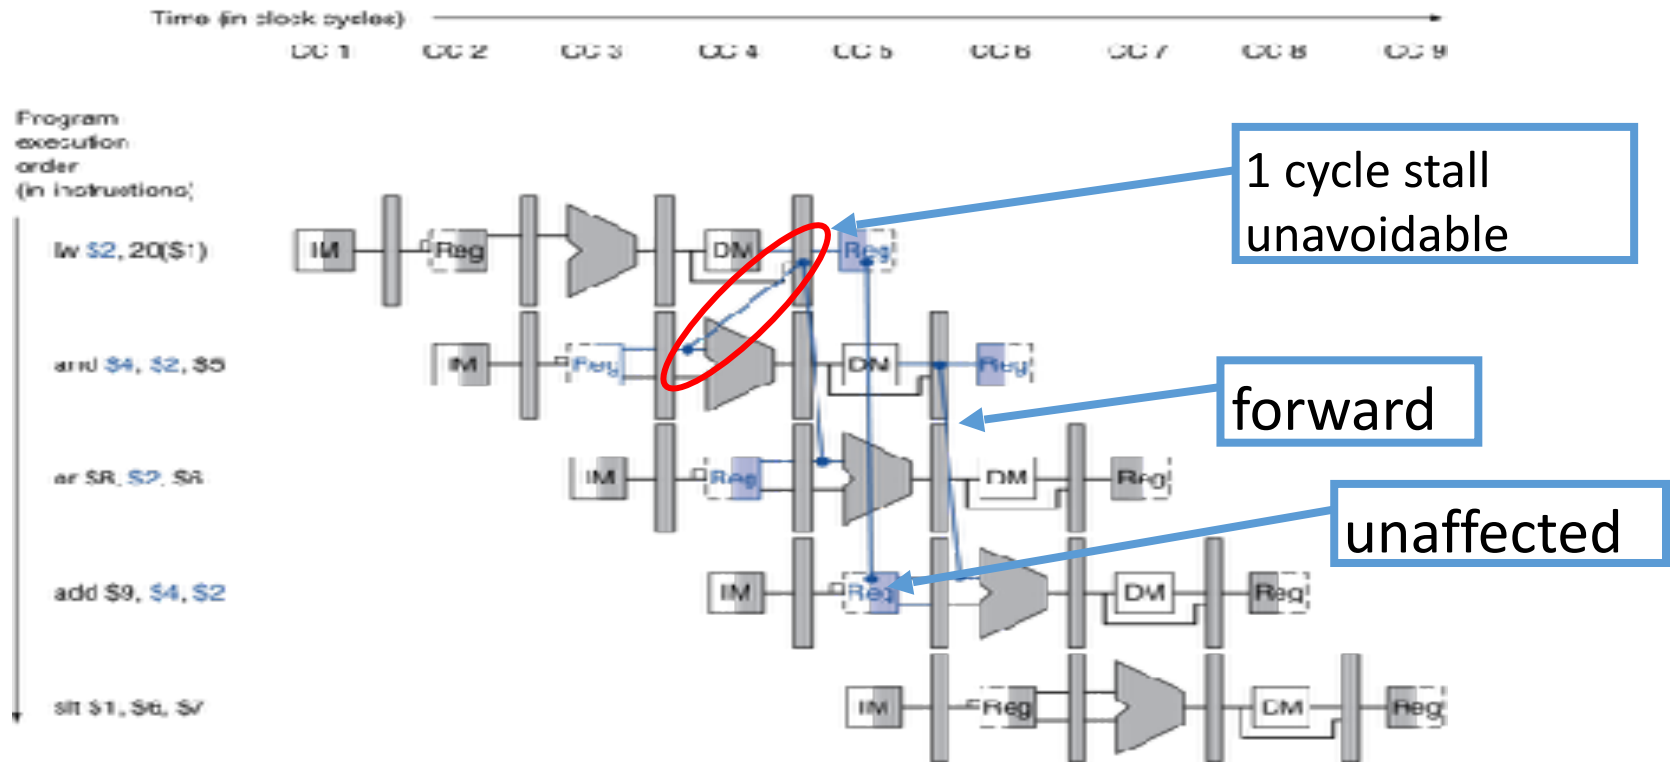
Administrivia

- *Project 3.1 still due next Wednesday (3/7)*
- Homework 2 due Friday (11:59PM)
- Project party on both *Monday (8-10, Cory 293)* and *Wednesday (7-10, Cory 293)*
- Guerrilla session Tonight 7-9pm, Barrows 20!
- Midterm 2, March 20, is moved to 8-10PM (was 7-9 on the website)
 - Alternative exam earlier, 6-8PM (so people don't need to be in exams until midnight :)
 - submit exam conflict form if they haven't

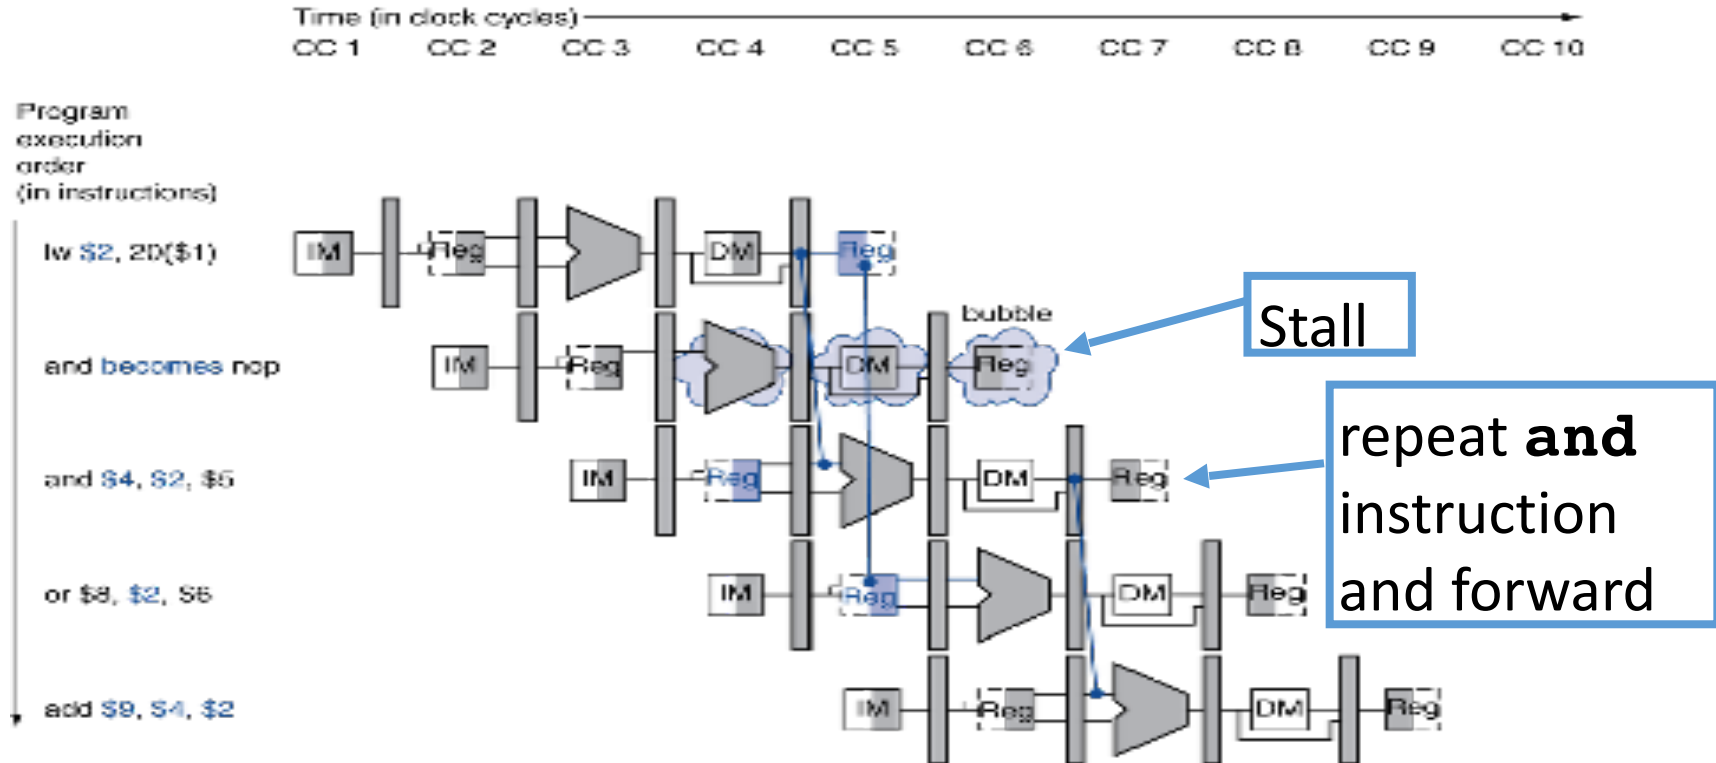
Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - **Load**
 - Control
- Superscalar processors

Load Data Hazard



Stall Pipeline



1w Data Hazard

- Slot after a load is called a *load delay slot*
 - If that instruction uses the result of the load, then the hardware will stall for one cycle
 - Equivalent to inserting an explicit **nop** in the slot
 - except the latter uses more code space
 - Performance loss!
- Idea:
 - Put unrelated instruction into load delay slot
 - No performance loss!

Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction!
- RISC-V code for $D=A+B$; $E=A+C$;

Original Order:

```
lw    t1, 0(t0)
lw    t2, 4(t0)
Stall! → add t3, t1, t2
sw    t3, 12(t0)
lw    t4, 8(t0)
Stall! → add t5, t1, t4
sw    t5, 16(t0)
```

13 cycles

Alternative:

```
lw    t1, 0(t0)
lw    t2, 4(t0)
lw    t4, 8(t0)
add    t3, t1, t2
sw    t3, 12(t0)
add    t5, t1, t4
sw    t5, 16(t0)
```

11 cycles

Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - **Control**
- Superscalar processors

Control Hazards


beq t0, t1, label

sub t2, s0, t5

or t6, s0, t3

xor t5, t1, s0

sw s0, 8(t3)



executed regardless of
branch outcome!

executed regardless of
branch outcome!!!

PC updated reflecting
branch outcome

Observation

- If branch not taken, then instructions fetched sequentially after branch are correct
- If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs


Kill Instructions after Branch if Taken

beq t0, t1, label

sub t2, s0, t5

or t6, s0, t3

label: xxxxxx



Taken branch

Convert to NOP

Convert to NOP

PC updated reflecting
branch outcome

Reducing Branch Penalties


- Every taken branch in simple pipeline costs 2 dead cycles
- To improve performance, use “branch prediction” to guess which way branch will go earlier in pipeline
- Only flush pipeline if branch prediction was incorrect

Branch Prediction

beq t0, t1, label

label:

.....



Taken branch

Guess next PC!

Check guess correct

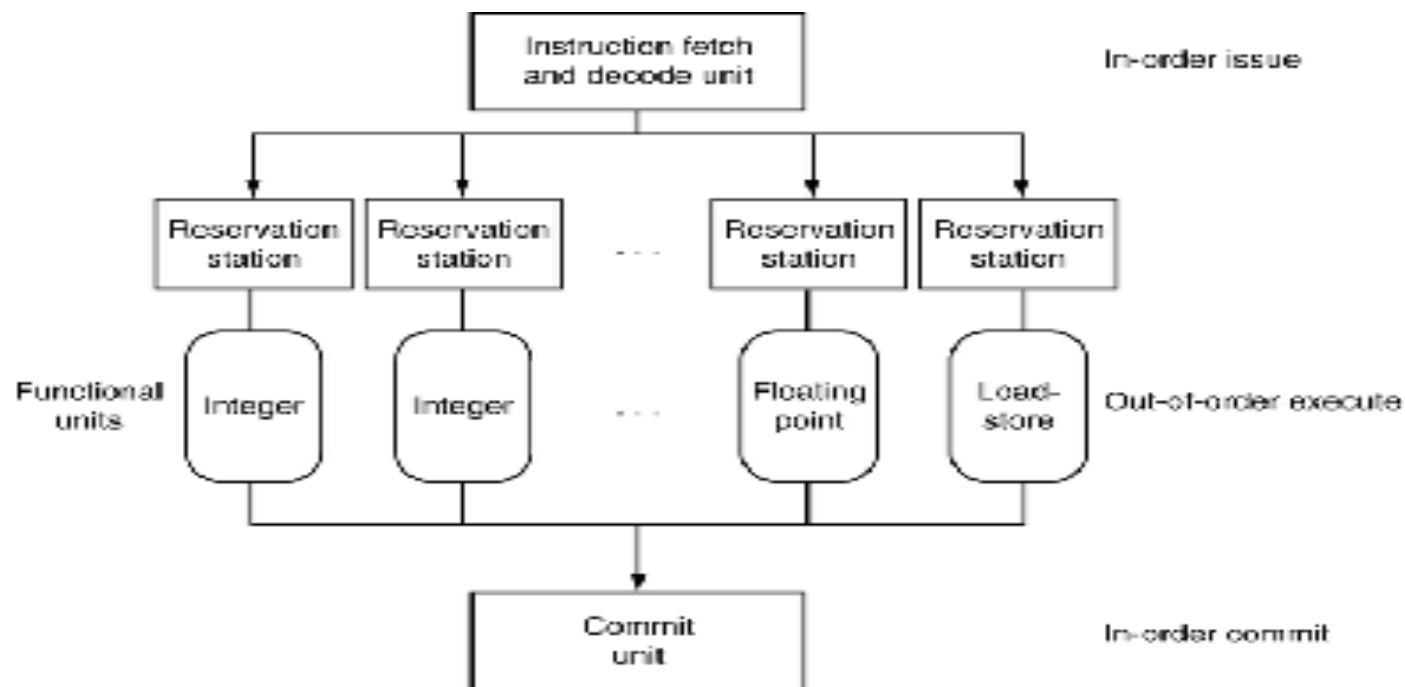
Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- **Superscalar processors**

Increasing Processor Performance


1. Clock rate
 - Limited by technology and power dissipation
2. Pipelining
 - “Overlap” instruction execution
 - Deeper pipeline: 5 \Rightarrow 10 \Rightarrow 15 stages
 - Less work per stage \rightarrow shorter clock cycle
 - But more potential for hazards ($CPI > 1$)
3. Multi-issue “super-scalar” processor
 - Multiple execution units (ALUs)
 - Several instructions executed simultaneously
 - $CPI < 1$ (ideally)

Superscalar Processor



P&H p. 340

Benchmark: CPI of Intel Core i7



P&H p. 350

In Conclusion

- Pipelining increases throughput by overlapping execution of multiple instructions
- All pipeline stages have same duration
 - Choose partition that accommodates this constraint
- Hazards potentially limit performance
 - Maximizing performance requires programmer/compiler assistance
 - E.g. Load and Branch delay slots
- Superscalar processors use multiple execution units for additional instruction level parallelism
 - Performance benefit highly code dependent

Extra Slides

Pipelining and ISA Design

- RISC-V ISA designed for pipelining
 - All instructions are 32-bits
 - Easy to fetch and decode in one cycle
 - Versus x86: 1- to 15-byte instructions
 - Few and regular instruction formats
 - Decode and read registers in one step
 - Load/store addressing
 - Calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle

Superscalar Processor

- Multiple issue “superscalar”
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - $CPI < 1$, so use Instructions Per Cycle (IPC)
 - E.g., 4GHz 4-way multiple-issue
 - 16 BIPS, peak $CPI = 0.25$, peak $IPC = 4$
 - Dependencies reduce this in practice
- “Out-of-Order” execution
 - Reorder instructions dynamically in hardware to reduce impact of hazards
- *CS152 discusses these techniques!*