

Even More Caches: Successes & Failures

Outline...

- A Simple Program to Analyze Caches
- Branch Prediction
- Multiprocessors and Caches
 - And what that means
- Virtual Memory is a Cache
 - And what that means

So Lets Analyze Caches...

- A conceptually simple set of nested for loop:

```
• int array[MAXLEN]
  for(size = 1024; size <= MAXLEN; size = size << 1){
    for(stride = 1; stride <= size; stride = stride << 1){
      // Some initialization to eliminate compulsory misses
      // Repeat this loop enough to get good timing for computing AMAT
      for(i = 0; i < size; i += stride){
        array[i] = array[i] + i
      }
      // Now add some timing information for how long the
      // for loop took
    }
  }
```

- This is ***striding access***:
 - Rather than every element in the array this accesses every i^{th} element
- This is ***designed to break caches***:
 - Caches tend to work great up until the instant they don't...
 - So by seeing where and how the cache falls down, this can deduce the internal structure

Reminder:

Cache Miss Types

- **Compulsory:** A miss that occurs on first load
 - An infinitely large cache would still incur the miss
- **Capacity:** A miss that occurs because the cache isn't big enough
 - An infinitely large cache would not miss
- **Conflict:** A miss that occurs because the associativity doesn't allow the items to be stored
 - A fully associative cache of the same size would not miss

Reminder:

Other Parameters

- Block or Line size:
 - The # of bytes in each entry
- Associativity:
 - The degree of flexibility in storing a particular address
 - Direct mapped: One location
 - N-way set associative: one of N possible locations
 - Fully associative: Any location
- AMAT: Average Memory Access Time
 - $\text{hit time} + \text{miss penalty} * \text{miss rate}$

Breaking Caches: Capacity...

Computer Science 61C Spring 2018

Wawrzynek and Weaver

- Up until the test exceeds the cache capacity...
 - Everything is fine!
- But once $\text{sizeof}(\text{array}) > \text{cache size}$...
 - Then things break down and you start getting misses
- Which increases the loop time
 - $\text{AMAT} = \text{hit time} + \text{miss penalty} * \text{miss rate}$
- So where the size breaks capacity...



Breaking Caches: Spatial Locality

- Spatial locality breaks down if only a ***single*** entry in each cache line is ever accessed
 - Since the rest of the cache line provides no benefit...
- So worst-case behavior occurs when each line is only accessed in one location
 - So when $\text{stride} * \text{sizeof}(\text{int}) == \text{block size}$
- Combined with where the capacity break occurs...
 - And you now know the line-size

Un-breaking Caches: Associativity

- If your array is 2x the cache capacity...
 - But you are striding at $\geq 2 \times \text{block size}$...
- You aren't using all the cache entries
 - So by definition, all your misses are no longer capacity misses but **conflict** misses!
- Reminder: Tag/Index/Offset...
 - The index specifies the possible location for a set associative cache
 - So when do the accesses have different indexes?
 - That is when you have stopped having conflict misses

Breaking Caches: Levels...

- Each level is its own cache...
 - To test L2, you must be using references that break L1...
- Which is fine for capacity, but...
 - If L2 line size \leq L1 line size, we can't reliably tell
 - Since the L1 cache provides the spacial locality
 - But generally most multi-level caches use the same linesize, defined by the external memory interface
 - If the L2 associativity \leq L1 associativity
 - The conflict misses will be removed in L1

Actual Test:

Raspberry Pi: L1 Cache hitting...

| | | | | | |
|-----------------|-------|-----------------|-------|-------------|------|
| • Size (bytes): | 32768 | Stride (bytes): | 4 | read+write: | 9 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 8 | read+write: | 7 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 16 | read+write: | 8 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 32 | read+write: | 9 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 64 | read+write: | 9 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 128 | read+write: | 7 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 256 | read+write: | 9 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 512 | read+write: | 9 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 1024 | read+write: | 8 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 2048 | read+write: | 9 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 4096 | read+write: | 8 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 8192 | read+write: | 8 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 16384 | read+write: | 8 ns |
| • Size (bytes): | 32768 | Stride (bytes): | 32768 | read+write: | 9 ns |

Actual Test:

Raspberry Pi: L1 missing...

| | | | | | | |
|---|---------------|-------|-----------------|-------|-------------|-------|
| • | Size (bytes): | 65536 | Stride (bytes): | 4 | read+write: | 7 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 8 | read+write: | 8 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 16 | read+write: | 9 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 32 | read+write: | 8 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 64 | read+write: | 8 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 128 | read+write: | 10 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 256 | read+write: | 10 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 512 | read+write: | 14 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 1024 | read+write: | 16 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 2048 | read+write: | 18 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 4096 | read+write: | 20 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 8192 | read+write: | 18 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 16384 | read+write: | 8 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 32768 | read+write: | 8 ns |
| • | Size (bytes): | 65536 | Stride (bytes): | 65536 | read+write: | 9 ns |

Actual Testing: Watching L2 Fail

- Size (bytes): 1048576 Stride (bytes): 4 read+write: 8 ns
- Size (bytes): 1048576 Stride (bytes): 8 read+write: 9 ns
- Size (bytes): 1048576 Stride (bytes): 16 read+write: 12 ns
- Size (bytes): 1048576 Stride (bytes): 32 read+write: 29 ns
- Size (bytes): 1048576 Stride (bytes): 64 read+write: 61 ns
- Size (bytes): 1048576 Stride (bytes): 128 read+write: 61 ns
- Size (bytes): 1048576 Stride (bytes): 256 read+write: 61 ns
- Size (bytes): 1048576 Stride (bytes): 512 read+write: 117 ns
- Size (bytes): 1048576 Stride (bytes): 1024 read+write: 117 ns
- Size (bytes): 1048576 Stride (bytes): 2048 read+write: 124 ns
- Size (bytes): 1048576 Stride (bytes): 4096 read+write: 140 ns
- Size (bytes): 1048576 Stride (bytes): 8192 read+write: 61 ns
- Size (bytes): 1048576 Stride (bytes): 16384 read+write: 24 ns
- Size (bytes): 1048576 Stride (bytes): 32768 read+write: 24 ns
- Size (bytes): 1048576 Stride (bytes): 65536 read+write: 21 ns
- Size (bytes): 1048576 Stride (bytes): 131072 read+write: 18 ns
- Size (bytes): 1048576 Stride (bytes): 262144 read+write: 8 ns
- Size (bytes): 1048576 Stride (bytes): 524288 read+write: 7 ns
- Size (bytes): 1048576 Stride (bytes): 1048576 read+write: 8 ns

Can Already See How Caches Fail Catastrophically...

- For a memory-bound task...
 - Fit in L1 cache: AMAT 7ns
 - Fit in L2 cache: AMAT ~20ns
 - Exceed L2 cache: AMAT 100+ns
- Performance drops by an ***order of magnitude*** when you exceed the capabilities of the cache even by not that much!

Complications...

- Of course, why we don't inflict this particular assignment on you anymore...
 - There are a lot of additional complications on modern processors
- Memory fetching/prefetching...
 - L2 cache is starting to hit memory at a stride of 64, but...
 - Performance keeps getting worse until the stride is larger
 - Memory is probably transferring 256B at a time to L1 as well as L2...
- There may be a "victim cache"
 - A small fully associative cache that holds the last few evicted cache lines:
So although L2 is only 16 way according to the documentation I find, we still are getting good performance on the 32 way and 64 way test
- And this is on a **simple** modern processor
 - **Only** 4 cores, **only** 2-issue superscalar
- This **used** to be a good homework assignment

More on Victim Caches...

- Observation: Conflict misses are a pain, but...
 - Perhaps a little associativity can help without having to be a fully associative cache
- In addition to the main cache...
 - Have a very small (16-64 entry) **fully associative** "victim" cache
- Whenever we evict a cache entry
 - Don't just get rid of it, put it in the victim cache
- Now on cache misses...
 - Check the victim cache first, if it is in the victim cache you can just reload it from there

Another Pathological Example...

- ```
int j, k, array[256*256];
for (k = 0; k < 255; k++) {
 for (j = 0; j < 256; j++) {
 array[256*j] += array[256*j + k + 1]
 }
}
```
- This from the homework has a nasty pathology...
  - It experiences **no** spacial locality of note: both array reads and the array write are a stride of 256 entries
  - And in the example for the homework it also generates a huge number of capacity misses



# But a minor tweak...

- ```
int j, k, array[256*256];  
for (j = 0; j < 256; j++) {  
    for (k = 0; k < 255; k++) {  
        array[256*j] += array[256*j + k + 1]  
    }  
}
```
- And now it runs vastly better as it changes from stride 256 to stride 1 and stride 0...
 - Stride 0 == best case temporal locality
 - Stride 1 == best case spacial locality

Clicker Question...

- Adding an additional layer of cache hierarchy (e.g. L4 in a system with L1-L3) will:
A) Improve. B) Not affect. C) Make Worse. D) 🙄
the following properties:
 - Lower level cache hit time
 - Lower level cache miss rate
 - Lower level cache miss penalty
 - AMAT on all workloads
 - (hopefully) AMAT on most workloads

Administrivia

- Project Party: Wednesday, 7-10pm in 293 Cory
- Guerrilla session on Caches Thursday from 7-9pm, Barrows 20
- Midterm 2 is Tuesday, 3/20 from 8-10pm.
 - Everyone (including DSP and conflict) should have already heard from Dylan or me if you have alternative exam arrangements
 - Reminder, you are allowed **two** hand-written, double sided note sheets
- Midterm Review this Sunday March 18th 1-4PM GPBB 100
- HW3 the cache part due Friday 3/16

Blocking-out Data

- Very common motif (yes, I like go lang...):
 - ```
for i := range A {
 for j := range B {
 fn(A[i], B[j])
 }
}
```
  - "Do something for every pair of elements"
- If B fits in the cache...
  - Its all good...
- But if B doesn't...
  - The inner iteration is going to be dominated by **capacity** misses as B has to keep being reloaded
    - So there is no more temporal locality for B[j]
  - But the fetches of A are still going to be fine because a lot of **temporal locality** for A[i]
- And its going to be very good **up until the instance it isn't**
  - Caches don't tend to degrade performance gracefully... Instead you get step-functions

# Implications...

- If one array is a lot bigger than the other...
  - It should be the outer one in the loop: It doesn't generally matter if the outer one doesn't fit
    - Since there is tons of temporal locality for the outer array that the cache will take advantage of
- But if both don't fit, you need to be better
  - ```
for i := 0; i < len(A); i += BLOCK {  
    for j := 0; j < len(B); j += 1 {  
        for k := 0; k < BLOCK; k++ {  
            if (k + i) < len(A) {  
                fn(A[i*block+k], B[j])  
            }  
        }  
    }  
}
```
 - Now have a lot more temporal locality for the entries of both A and B, **if *BLOCK* is set correctly**

Another Cache: Branch Predictor

- In our simple pipeline, we assume branches-not-taken
 - Always start fetching the next instruction
- If a branch or jump is taken...
 - Then we have to kill the non-taken instructions so they don't cause side effects
- But both branches and jumps are PC relative...
 - So if we can quickly look at the instruction and decide 'eh, probably taken/not' we can compute the new location for the PC if we can guess right
 - Which for `ja1` we always can, but branches we need to guess
- Idea: **branches** have temporal locality!
 - Loops: `for (x = 0; x < n...)`
 - Rare conditionals: `if (err) ...`

A Simple Branch Predictor

- Have an N entry, direct-mapped memory
 - EG, a 1024x1b memory
- If fetched instruction is a branch...
 - Check if the bit for pc[12:2] is set...
 - If so, set next PC to PC + branch offset for fetching
 - Set bit in pipeline to say "branch predicted-taken"
- When actually evaluating branch in EX...
 - Set pc[12:2] to branch taken/not-taken status
 - If branch taken but predicted not-taken
 - ***Kill untaken instructions***
 - If branch not taken but predicted taken
 - ***Kill predicted instructions***

Improving it slightly...

- How about 2 bits:
 - Each entry starts at 01...
 - If taken, increment with saturating arithmetic (so max is 11)
 - If not taken, decrement by one (so min is 00)
- If the upper bit is 1, assume the branch is taken
 - Now a function with a commonly taken loop will only mispredict once rather than twice
- Miss penalty for a mispredicted branch: The # of instructions that got terminated because of the wrong prediction
 - EG, on a dual-issue, 10-deep 2x superscalar like a Raspberry Pi or smartphone: Probably ~10 instructions
 - On a modern x86? It could be 20+
- Can then try even fancier predictors...
 - But we get into ***diminishing returns***:
 - The simple-smart thing (e.g. a two bit branch predictor) is a big win...
 - But trying to get fancier no longer helps nearly as much

Approximate Cache...

- The data caches are exact:
 - They will return an answer that is exactly what is asked for
- But this branch predictor is approximate:
 - It can make mistakes due to aliasing:
Its not actually storing the full address as a tag to check
- Sometimes its OK to be wrong
 - So data structures that are this way are also particularly interesting...
 - EG "Bloom Filters": One of the **coolest** data structures on the planet.
 - Is this element in the set? Using a small amount of memory
 - Yes it is? Will always return "Yes"
 - No it isn't? Will mostly return "No", but can have an error

A related cache: return target location...

- Observation:
 - On RISC-V, you call a function with `jal` or `jalr` (object oriented) with the return set to `ra`
 - And you return with a `jalr` with the source register as `ra`
- So lets maintain a small stack in hardware...
 - Whenever we see `jal` or `jalr` writing to `ra`:
We write `PC+4` into the stack
 - Whenever we see `jalr` reading from `ra`:
We predict the top of this stack as the next PC, and pop this stack
- Result: We should always correctly predict a function return address...
 - Works as long as we don't exceed the stack depth: once we hit that we will start getting misses

And A Final Related Cache: Branch Target Buffer

- Function calls using `jal` we will never mispredict on RISC-V
 - Since they are all PC relative
- But so much today is object-oriented programming which uses `jalr`:
C++ and Java object calls are equivalent to calling pointers to functions
 - `foo.bar()` is implemented as something like this:

```
lw t0 0(a0) # Get the pointer to the "virtual function table" in the object
lw t0 8(t0) # Get the pointer to the function to actually call
jalr ra t0  # Do a JALR to call bar(),
            # with the object foo as the first implicit argument
```
- So cache this as well:
 - On a `jalr` which writes to `ra`...
Look in a small cache for the address to predict to based on current PC
 - When evaluating the jump, set the value in this cache to the address used
- It is the x86 equivalent of this cache that is part of one of the Spectre vulnerabilities

Caches and Multiprocessors...

- These days practically every computer is a multiprocessor
 - Since that is the only way we know how to increase computation by throwing more silicon at the problem
 - But we can't make single-processor performance **worse** in this process
 - So these processors **must** have significant caches
 - And because the L1 caches are integrated into the pipeline, to prevent structural hazards each processor must have its **own** caches
- What happens if multiple processors are accessing the same piece of memory?

Multiple Processor Reading Memory?

- No problem!
 - Each processor just caches the data independently
- There is ***no*** issue with multiple processors reading the same thing
 - Their own caches have a unique copy...
But the values should always be the same

Multiple Processors Writing?

- We need **coherency**: Writes from one processor **must** be reflected in memory that other processors read after some short period of time
- There have been processors made without this, but it is impossible to program these
- Goal is to guarantee the following property:
 - If processor A writes to memory location L, **within time T** , all other processors will see the updated data

Idea: Broadcasting Messages...

- We need a way for the processors to communicate
 - So we have some sort of fabric
 - It could be a shared bus
 - It could be something looking more like a packet-switched computer network
- Each processor (or more precisely its cache) can send and receive messages
 - Requests are "broadcast", a single sender can send a message to anybody...
 - Replies may or may not be broadcast: Can go to everyone or could go to a specific recipient

Broadcasting Writes...

- Processor **A** wants to write to physical location **L** for the first time...
 - First do a write-allocate...
 - It then sets the **dirty** bit on the cache
 - And **broadcasts** a message that "I am writing to **L**"
- All other processors which receive the message
 - Do I have address **L** cached?
 - If no: Do nothing
 - If yes: **invalidate** the entry in the cache
 - **Snooping** on requests:
Term comes from when all processors shared the same memory bus to communicate with the memory

Broadcasting Reads

- If there is a miss in the upper level of the cache in the processor...
 - Broadcast a read request: "Hey, does anyone have location L ?"
- If nobody has written to this location...
 - The memory controller/common cache just does a fetch and returns it: Just like any other cache miss
- If a processor has this location with the dirty bit set...
 - It goes "Hey, I have this"
 - It **flushes** the entry (writing the value to memory) and clears the dirty bit
 - It then says the new value to the requesting processor

Why Does This Work?

- If processor A wants to write to a non-dirty line...
 - If the element is in the cache...
 - It will write, and all other processors **invalidate**: If they then want to read that location they will have to broadcast that request
 - If not, it performs a read first...
 - So if reads are correct, it is going to be correct from there
- If processor B wants to read...
 - If the element is in its cache...
 - It is correct, because if someone else wrote to that location **it would be invalid already**
 - If the element is not in its cache...
 - It will get the correct value from either another processor or the right location
 - And that other processor will now know it can't write because the dirty bit got cleared

CPU-0 reads byte at 0xdeadbeef...

Upper Level Cache & Memory...

CPU-0: I want to read 0xdeadbee0

Cache Controller: value of 0xdeadbee0 is 0xf00dd00dca11c003

CPU-0

| Address | Data | V | D |
|---------|------|---|---|
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

CPU-1

| Address | Data | V | D |
|---------|------|---|---|
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

CPU-1 reads byte at 0xdeadbeef...

Upper Level Cache & Memory...

CPU-1: I want to read 0xdeadbee0

Cache Controller: value of 0xdeadbee0 is 0xf00dd00dca11c003

CPU-0

| Address | Data | V | D |
|-----------------|-------------------------|----------|----------|
| deadbee0 | f00dd00dca11c003 | 1 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

CPU-1

| Address | Data | V | D |
|-----------------|-------------------------|----------|----------|
| deadbee0 | f00dd00dca11c003 | 1 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

CPU-1 writes data at 0xdeadbee0...

Upper Level Cache & Memory...

CPU-1: I want to start writing to 0xdeadbee0

CPU-0

| Address | Data | V | D |
|-----------------|-------------------------|----------|----------|
| deadbee0 | f00dd00dca11c003 | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

CPU-1

| Address | Data | V | D |
|-----------------|-------------------------|----------|----------|
| deadbee0 | cafef00dda016666 | 1 | 1 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

CPU-0 reads byte at 0xdeadbeef...

Upper Level Cache & Memory...

CPU-0: I want to read 0xdeadbee0

CPU-1: value of 0xdeadbee0 is 0xcafef00dda016666

CPU-0

| Address | Data | V | D |
|-----------------|-------------------------|----------|----------|
| deadbee0 | cafef00dda016666 | 1 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

CPU-1

| Address | Data | V | D |
|-----------------|-------------------------|----------|----------|
| deadbee0 | cafef00dda016666 | 1 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |
| | | 0 | 0 |

So Enter a New Miss Type:

Coherence

- A coherence miss occurs when two processes want to access the same data
 - Coherence misses are caused only by **writes**, not **reads**
 - The write will invalidate all **other** caches
- It means there is an **anti-pattern** that can cause performance artifacts on multiprocessors
 - Multiple processes reading to the same memory? Sure!
 - But if one starts writing to that memory...
 - The **other** processor will start getting coherency misses
 - But such misses also only go up to the shared cache:
Why multiprocessors, when possible, use a shared cache between all processors
 - Reasonably easy to avoid **with proper program structure**

Oh, and *Incoherence misses*

- What about when we have multiple processes running on the same processor
 - A modern x86 creates two "virtual" processors which share resources ("Hyperthreading")
- If those two processes have the same working set...
 - Great!
- If those two processes have different working sets...
 - This effectively acts like reducing the cache size with the corresponding increase in the miss rate

Virtual Memory Paging As A Cache...

- How virtual memory works we will cover later...
- But for now, its easy to model as a basic cache...
- Your program is given the illusion of as much RAM as it wants...
 - But only on the condition that it actually doesn't want it! :)
- Idea: Virtual memory can copy "pages" between RAM and disk
 - The main memory thus acts as a cache for the disk...

But What Are The Properties...

- Associativity: Effectively fully associative
- Replacement policy: Effectively full LRU
- Block size: 4kB or more
 - Some argue it should now be 64kB or 256kB these days
- Hit time...
 - Call in 0
- Miss penalty...
 - Latency to get a block from disk: 1ms or so for an SSD...
Or put in clock terms, a 1 GHz clock \rightarrow 1,000,000 clock cycles!

So what are the implications?

- As long as you don't really use it, Virtual Memory is great...
 - Basically as long as your **working set** fits in physical memory, virtual memory is great at handling a little extra...
- But as soon as your working set exceeds physical memory, your performance goes to crap...
 - The system starts **thrashing**: Repeatedly needing to copy data to and from disk...
 - Similar to **thrashing the cache** when your working set exceeds cache capacity
 - If you have a spinning disk, it really starts sounding like the computer is suffering....

Attacker Selected Input

- Alternatively, if the attacker can select the input...
 - The attacker can select **hard** input:
E.G. Traffic that causes ping-ponging in a direct mapped cache
- Nick's problem:
 - He had to cache IP addresses (32 bit values)
 - This is a network application for security
 - He only wants to store a small amount of information
 - On chip storage expensive (in this case, on an FPGA)
 - He had plenty of room to pipeline
 - Each network packet is independent: No need to forward
 - Misses are expensive
 - Requires processing the packet in software

Nick's Cache #1: Permutation Cache

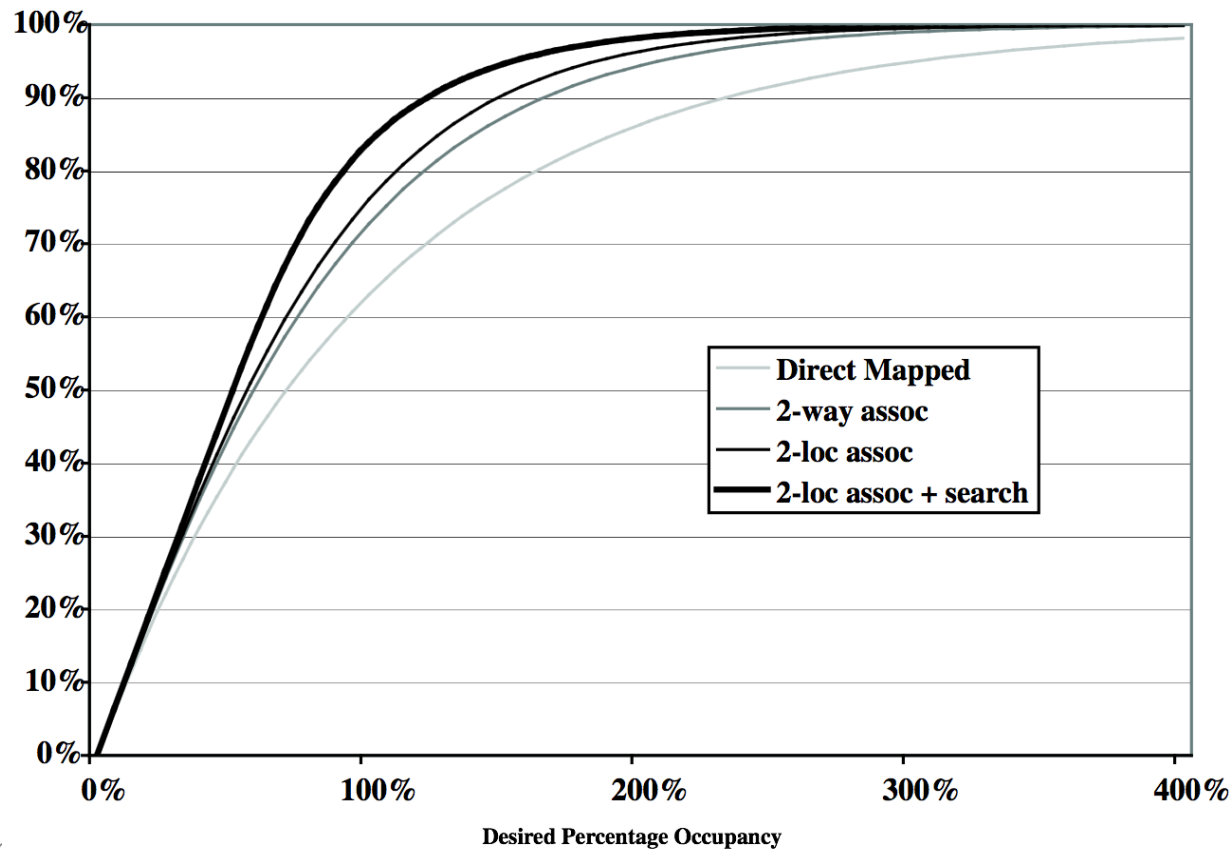
- Traditionally, you would hash the address
 - With a "salt" to randomize things
 - Attacker needs to break the hash function to predict whether two different addresses will match to the same location
 - But this requires storing the whole original IP for your tag
- So instead of a hash, use a 32b **keyed permutation**
 - Aka a 32b block cypher
- Now you can use a conventional tag/index approach
 - Requires only storing the tag -> space mattered in this application

Nick's Cache #2:

Location Associativity

- The fabric Nick used allowed "dual-ported" memories
 - Like your register file on your processor design: two independent read ports that operate at the same time
- Rather than using set associativity...
 - Instead do two **different** permutations (keys) and have one of two possible locations
- If X, Y, and Z map to the same location with one key...
 - They probably do not on the other key: fewer conflict misses
 - Even better, can probably move a value to further reduce conflict misses

Simulation... *Actual* Occupancy for Caching Random Elements



And Since the Permutation looks "random"...

- It is one of those cases where simulation matches reality...
 - Because I'm caching **random** values
- Most cache studies are much more program dependent
- But it does give a way of reducing **conflict** misses without having to increase associativity
 - Just need to calculate two permutations in parallel rather than just one