

CS 61C:
Great Ideas in Computer Architecture

Lecture 18:
Parallel Processing – SIMD

John Wawrzynek & Nick Weaver
<http://inst.eecs.berkeley.edu/~cs61c>

61C Survey

It would be nice to have a review lecture every once in a while, actually showing us how things fit in the bigger picture

Agenda

- **61C – the big picture**
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

61C Topics so far ...

- What we learned:
 1. Binary numbers
 2. C
 3. Pointers
 4. Assembly language
 5. Processor micro-architecture
 6. Pipelining
 7. Caches
 8. Floating point
- What does this buy us?
 - Promise: execution speed
 - Let's check!

Reference Problem

- Matrix multiplication


- Basic operation in many engineering, data, and imaging processing tasks
 - Ex:, Image filtering, noise reduction, ...
- Core operation in Neural Nets and Deep Learning
 - Image classification (cats ...)
 - Robot Cars
 - Machine translation
 - Fingerprint verification
 - Automatic game playing

- **dgemm**

- **d**ouble-precision floating-point **g**eneral **m**atrix-**m**ultiply
- Standard well-studied and widely used routine
- Part of Linpack/Lapack

2D-Matrices


- Square matrix of dimension $N \times N$
 - i indexes through rows
 - j indexes through columns



Matrix Multiplication

$$\mathbf{C} = \mathbf{A} * \mathbf{B}$$

$$C_{ij} = \sum_k (A_{ik} * B_{kj})$$



2D Matrix Memory Layout

a_{ij}

a_{00}	a_{01}	a_{02}	a_{03}
a_{10}	a_{11}	a_{12}	a_{13}
a_{20}	a_{21}	a_{22}	a_{23}
a_{30}	a_{31}	a_{32}	a_{33}

- $a[][]$ in C uses row-major
- Fortran uses column-major
- Our examples use column-major

Row-Major

Row	a_{13}	0x38
	a_{12}	0x30
	a_{11}	0x28
	a_{10}	0x20
	a_{03}	0x18
Row	a_{02}	0x10
	a_{01}	0x8
	a_{00}	0x0

$a_{ij} : a[i*N + j]$

Column-Major

Column	a_{31}	0x38
	a_{21}	0x30
	a_{11}	0x28
	a_{01}	0x20
	a_{30}	0x18
Column	a_{20}	0x10
	a_{10}	0x8
	a_{00}	0x0

$a_{ij} : a[i + j*N]$

dgemm Reference Code: Python

- Linear addressing, assumes “column-major” memory layout

```
def dgemm(N, a, b, c):  
    for i in range(N):  
        for j in range(N):  
            c[i+j*N] = 0  
            for k in range(N):  
                c[i+j*N] += a[i+k*N] * b[k+j*N]
```

N	Python [Mflops]
32	5.4
160	5.5
480	5.4
960	5.3

- 1 MFLOP = 1 Million floating-point operations per second (fadd, fmul)
- dgemm(N ...) takes $2 \cdot N^3$ flops

C

- $c = a * b$
- a, b, c are $N \times N$ matrices

```
// Scalar; P&H p. 226
void dgemm_scalar(int N, double *a, double *b, double *c) {
    for (int i=0; i<N; i++)
        for (int j=0; j<N; j++) {
            double cij = 0;
            for (int k=0; k<N; k++)
                //      a[i][k] * b[k][j]
                cij += a[i+k*N] * b[k+j*N];
            // c[i][j]
            c[i+j*N] = cij;
        }
}
```

Timing Program Execution

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(void) {
    // start time
    // Note: clock() measures execution time, not real time
    //       big difference in shared computer environments
    //       and with heavy system load
    clock_t start = clock();

    // task to time goes here:
    // dgemm(N, ...);

    // "stop" the timer
    clock_t end = clock();

    // compute execution time in seconds
    double delta_time = (double)(end-start)/CLOCKS_PER_SEC;
}
```

C versus Python

N	C [GFLOPS]	Python [GFLOPS]
32	1.30	0.0054
160	1.30	0.0055
480	1.32	0.0054
960	0.91	0.0053



Which other class gives you this kind of power?
We could stop here ... but why? Let's do better!

Agenda

- 61C – the big picture
- **Parallel processing**
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

Why Parallel Processing?

- CPU Clock Rates are no longer increasing
 - Technical & economic challenges
 - Advanced cooling technology too expensive or impractical for most applications
 - Energy costs are prohibitive
- Parallel processing is only path to higher speed
 - Compare airlines:
 - Maximum air-speed limited by economics
 - Use more and larger airplanes to increase throughput
 - (And smaller seats ...)

Using Parallelism for Performance

- Two basic approaches to parallelism:
 - Multiprogramming
 - run multiple independent programs in parallel
 - “Easy”
 - Parallel computing
 - run one program faster
 - “Hard”
- We’ll focus on parallel computing in the next few lectures

New-School Machine Structures (It's a bit more complicated!)

Software

Hardware

- Parallel Requests

Assigned to computer
e.g., Search “Katz”

- Parallel Threads

Assigned to core
e.g., Lookup, Ads

- Parallel Instructions

>1 instruction @ one time
e.g., 5 pipelined instructions

- Parallel Data

>1 data item @ one time
e.g., Add of 4 pairs of words

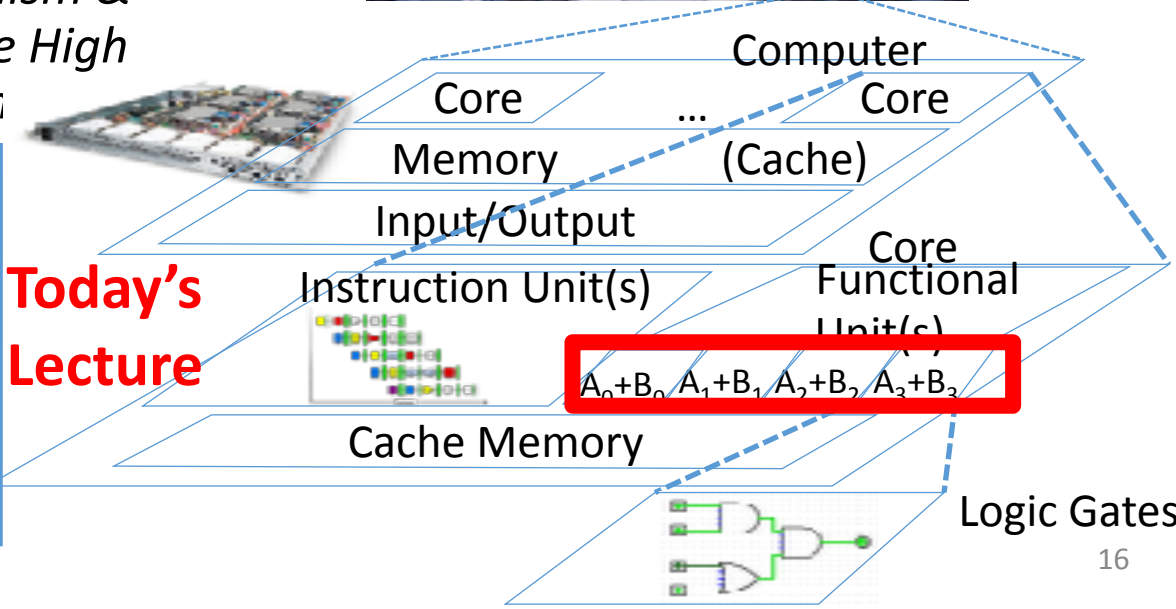
- Hardware descriptions

All gates @ one time

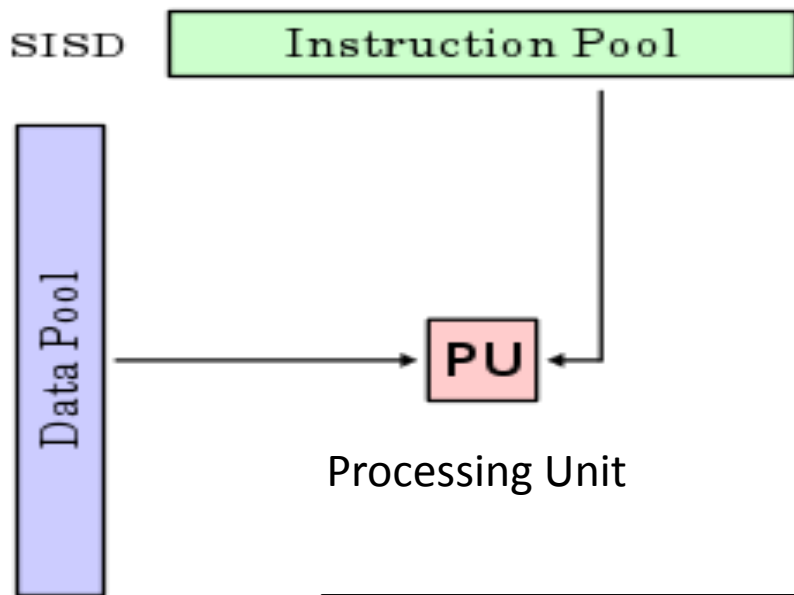
- Programming Languages

*Harness
Parallelism &
Achieve High
Perform*

Warehouse
Scale
Computer



Single-Instruction/Single-Data Stream (SISD)

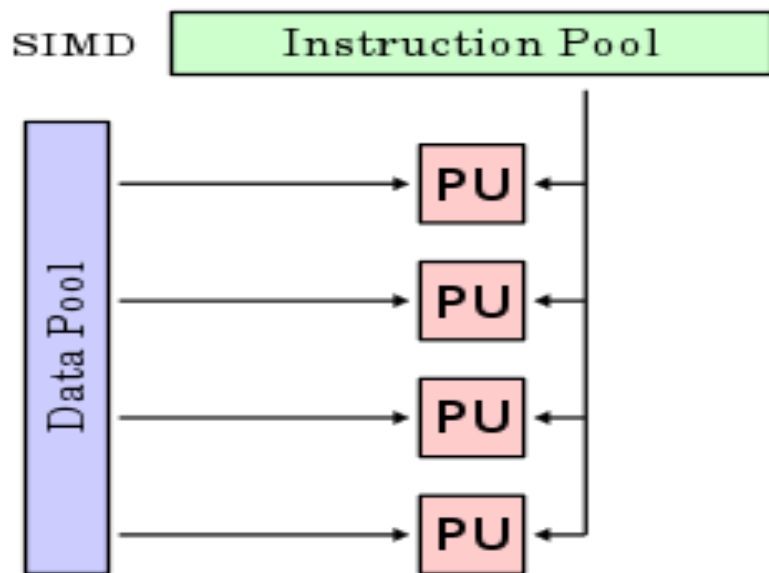


Processing Unit

- Sequential computer that exploits no parallelism in either the instruction or data streams. Examples of SISD architecture are traditional uniprocessor machines
 - E.g. Our RISC-V processor

This is what we did up to now in 61C

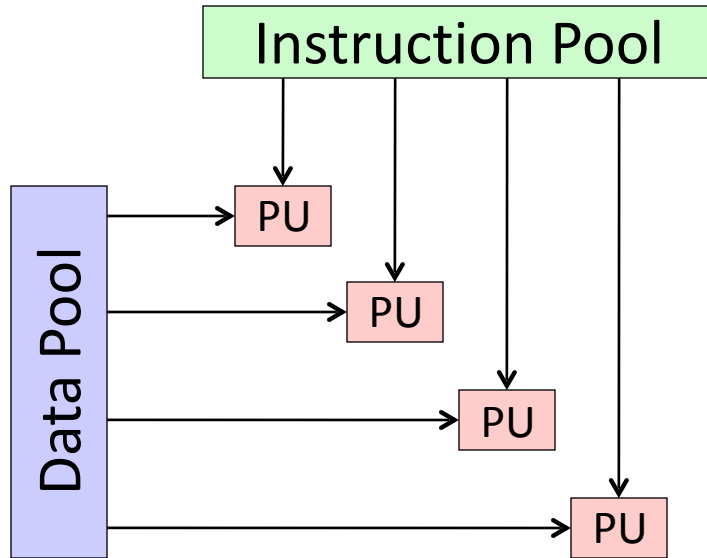
Single-Instruction/Multiple-Data Stream (SIMD or “sim-dee”)



- SIMD computer processes multiple data streams using a single instruction stream, e.g., Intel SIMD instruction extensions or NVIDIA Graphics Processing Unit (GPU)

Today's topic.

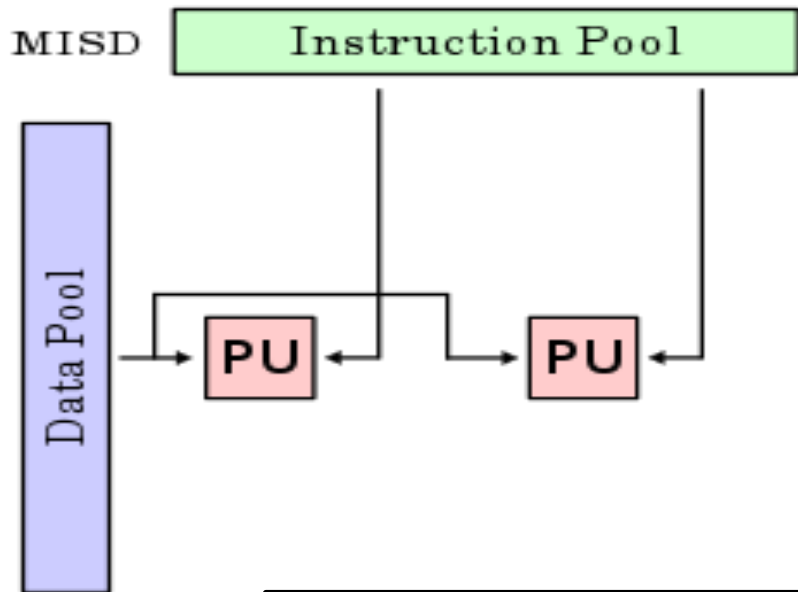
Multiple-Instruction/Multiple-Data Streams (MIMD or “mim-dee”)



- Multiple autonomous processors simultaneously executing different instructions on different data.
 - MIMD architectures include multicore and Warehouse-Scale Computers

Topic of Lecture 19 and beyond.

Multiple-Instruction/Single-Data Stream (MISD)



- Multiple-Instruction, Single-Data stream computer that processes multiple instruction streams with a single data stream.
 - Historical significance

This has few applications. Not covered in 61C.

Flynn* Taxonomy, 1966

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345 (Clovertown)

- SIMD and MIMD are currently the most common parallelism in architectures – usually both in same system!
- Most common parallel processing programming style: Single Program Multiple Data (“SPMD”)
 - Single program that runs on all processors of a MIMD
 - Cross-processor execution coordination using synchronization primitives

*Prof. Michael Flynn, Stanford



Patterson and Hennessy win Turing Award!



David A. Patterson, right, and John L. Hennessy in the early 1990s. The man won the Turing Award for their pioneering work on a computer chip design that is now used by most of the tech industry. Source: Getty

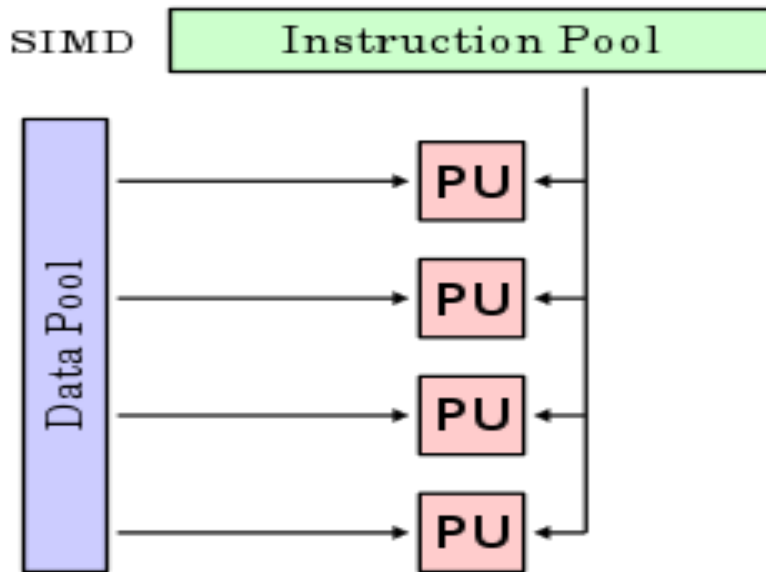


NEW YORK, NY, March 21, 2018 – [ACM](#), the Association for Computing Machinery, today named [John L. Hennessy](#), former President of Stanford University, and [David A. Patterson](#), retired Professor of the University of California, Berkeley, recipients of the 2017 ACM A.M. Turing Award for pioneering a systematic, quantitative approach to the design and evaluation of computer architectures with enduring impact on the microprocessor industry. Hennessy and Patterson created a systematic and quantitative approach to designing faster, lower power, and reduced instruction set computer (RISC) microprocessors.

Agenda


- 61C – the big picture
- Parallel processing
- **Single instruction, multiple data**
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

SIMD – “Single Instruction Multiple Data”




SIMD (Vector) Mode

SIMD Mode



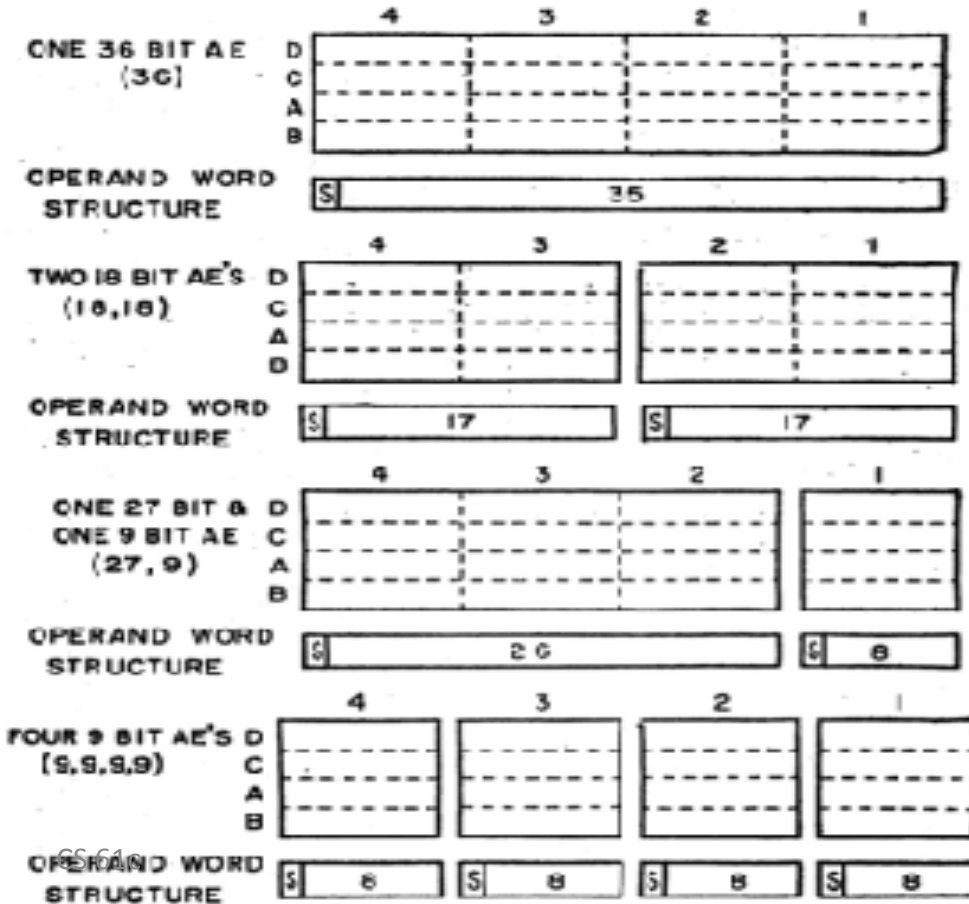
Scalar Mode



SIMD Applications & Implementations

- Applications
 - Scientific computing
 - Matlab, NumPy
 - Graphics and video processing
 - Photoshop, ...
 - Big Data
 - Deep learning
 - Gaming
 - ...
- Implementations
 - x86
 - ARM
 - RISC-V vector extensions

First SIMD Extensions: MIT Lincoln Labs TX-2, 1957



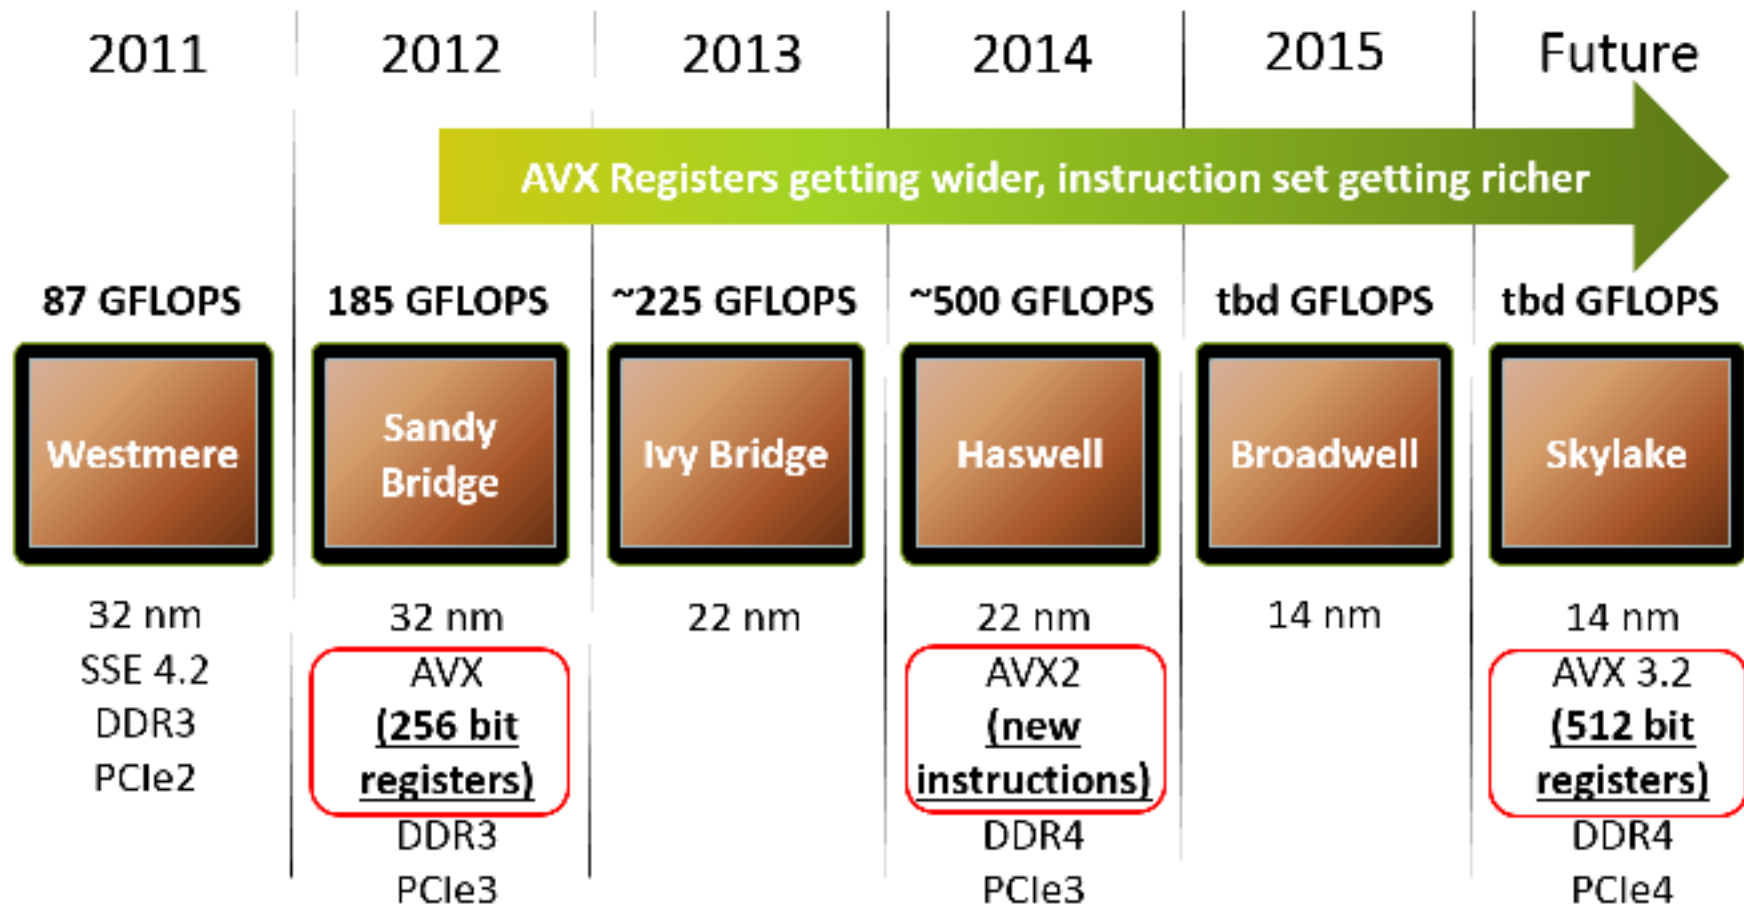
Intel x86 SIMD: Continuous Evolution

MMX 1997

1999	2000	2004	2006	2007	2008	2009	2010\11
SSE	SSE2	SSE3	SSSE3	SSE4.1	SSE4.2	AES-NI	AVX
70 instr Single-Precision Vectors Streaming operations	144 instr Double-precision Vectors 8/16/32 64/128-bit vector integer	13 instr Complex Data	32 instr Decode	47 instr Video Graphics building blocks Advanced vector instr	8 instr String/XML processing POP-Count CRC	7 instr Encryption and Decryption Key Generation	~100 new instr. ~300 legacy sse instr updated 256-bit vector 3 and 4- operand instructions

Intel Advanced Vector eXtensions

AVX also supported by AMD processors



Laptop CPU Specs

```
$ sysctl -a | grep cpu
```

```
hw.physicalcpu: 2
```

```
hw.logicalcpu: 4
```


```
machdep.cpu.brand_string:
```

```
Intel(R) Core(TM) i7-5557U CPU @ 3.10GHz
```


```
machdep.cpu.features: FPU VME DE PSE TSC MSR PAE MCE CX8 APIC SEP  
MTRR PGE MCA CMOV PAT PSE36 CLFSH DS ACPI MMX FXSR SSE SSE2 SS  
HTT TM PBE SSE3 PCLMULQDQ DTES64 MON DSCPL VMX EST TM2 SSSE3  
FMA CX16 TPR PDCM SSE4.1 SSE4.2 x2APIC MOVBE POPCNT AES PCID  
XSAVE OSXSAVE SEGLIM64 TSCTMR AVX1.0 RDRAND F16C
```

```
machdep.cpu.leaf7_features: SMEP ERMS RDWRFSGS TSC_THREAD_OFFSET  
BMI1 AVX2 BMI2 INVPCID SMAP RDSEED ADX IPT FPU_CSDS
```

AVX SIMD Registers



SIMD Data Types



(Now also AVX-512 available)

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- **SIMD matrix multiplication**
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

Problem

- Today's compilers can generate SIMD code
- But in some cases, better results by hand (assembly)
- We will study x86 (not using RISC-V as no vector hardware widely available yet)
 - Over 1000 instructions to learn ...
- Can we use the compiler to generate all non-SIMD instructions?

x86 SIMD “Intrinsics”



`_mmul_pd`

Technologies

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☒ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX-512
- ☐ KNC
- ☐ SVMML
- ☐ Other

Categories

- ☐ Application-Targeted
- ☐ Arithmetic
- ☐ Bit Manipulation
- ☐ Cast
- ☐ Compare

`_mm256d _mm256_mul_pd (__m256d a, __m256d b)`

Synopsis

`_mm256d _mm256_mul_pd (__m256d a, __m256d b)`

#include "immintrin.h"

Instruction: `vmulpd ymm, ymm, ymm`

CPUID Flag: AVX

← Intrinsic

← assembly instruction

Description

Multiply packed double-precision (64-bit) floating-point elements in `a` and `b`, and store the results in `dst`.

Operation

```
FOR j := 0 to 3
    i := j*64
    dst[i+63:i] := a[i+63:i] * b[i+63:i]
ENDFOR
dst[MAX:255] := 0
```

← 4 parallel multiplies

Performance

Architecture	Latency	Throughput
Haswell	5	0.5
Ivy Bridge	5	1
Sandy Bridge	5	1

← 2 instructions per clock cycle (CPI = 0.5)

<https://software.intel.com/sites/landingpage/IntrinsicsGuide/>

x86 Intrinsics AVX Data Types

Intrinsics: Direct access to assembly from C

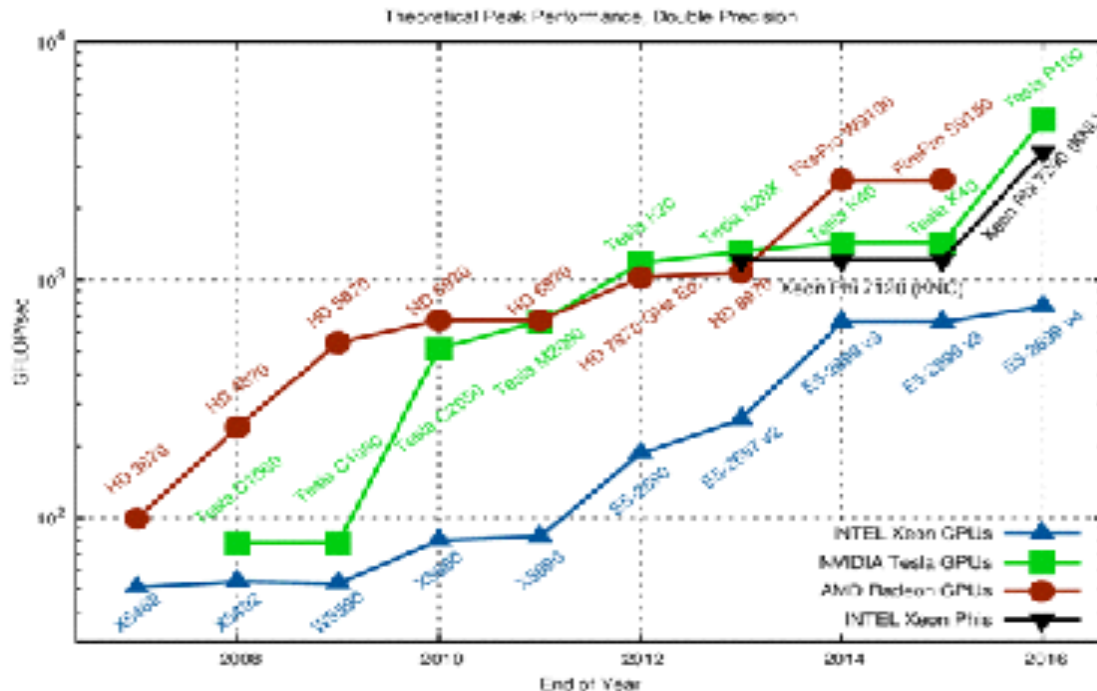
Type	Meaning
<code>__m256</code>	256-bit as eight single-precision floating-point values, representing a YMM register or memory location
<code>__m256d</code>	256-bit as four double-precision floating-point values, representing a YMM register or memory location
<code>__m256i</code>	256-bit as integers, (bytes, words, etc.)
<code>__m128</code>	128-bit single precision floating-point (32 bits each)
<code>__m128d</code>	128-bit double precision floating-point (64 bits each)

Intrinsics AVX Code Nomenclature

Marking	Meaning
[s/d]	Single- or double-precision floating point
[i/u]nnn	Signed or unsigned integer of bit size <i>nnn</i> , where <i>nnn</i> is 128, 64, 32, 16, or 8
[ps/pd/sd]	Packed single, packed double, or scalar double
epi32	Extended packed 32-bit signed integer
si256	Scalar 256-bit integer

Raw Double-Precision Throughput

Characteristic	Value
CPU	i7-5557U
Clock rate (sustained)	3.1 GHz
Instructions per clock (mul_pd)	2
Parallel multiplies per instruction	4
Peak double FLOPS	24.8 GFLOPS



<https://www.karlsruher.net/2013/06/cpu-gpu-and-mic-hardware-characteristics-over-time/>

Actual performance is lower because of overhead

Vectorized Matrix Multiplication


for i ...; **i+=4**

for j ...

Inner Loop:


```
__m256d c0 = {0,0,0,0};  
for (int k=0; k<N; k++) {  
    c0 = _mm256_fmadd_pd(  
        _mm256_load_pd(a+i+k*N),  
        _mm256_broadcast_sd(b+k+j*N),  
        c0);  
}  
_mm256_store_pd(c+i+j*N, c0);
```

i += 4



“Vectorized” dgemm

```
// AVX intrinsics; P&H p. 227
void dgemm_avx(int N, double *a, double *b, double *c) {
    // avx operates on 4 doubles in parallel
    for (int i=0; i<N; i+=4) {
        for (int j=0; j<N; j++) {
            // c0 = c[i][j]
            __m256d c0 = {0,0,0,0};
            for (int k=0; k<N; k++) {
                c0 = _mm256_add_pd(
                    c0, // c0 += a[i][k] * b[k][j]
                    _mm256_mul_pd(
                        _mm256_load_pd(a+i+k*N),
                        _mm256_broadcast_sd(b+k+j*N)));
            }
            _mm256_store_pd(c+i+j*N, c0); // c[i,j] = c0
        }
    }
}
```



Performance

N	Gflops	
	scalar	avx
32	1.30	4.56
160	1.30	5.47
480	1.32	5.27
960	0.91	3.64

- 4x faster
- But still << theoretical 25 GFLOPS!

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- **Loop unrolling**
- Memory access strategy - blocking
- And in Conclusion, ...

Loop Unrolling

- On high performance processors, optimizing compilers performs “loop unrolling” operation to expose more parallelism and improve performance:

```
for(i=0; i<N; i++)  
    x[i] = x[i] + s;
```

- Could become:

```
for(i=0; i<N; i+=4) {  
    x[i]      = x[i] + s;  
    x[i-1]    = x[i+1] + s;  
    x[i-2]    = x[i+2] + s;  
    x[i-3]    = x[i+3] + s;  
}
```

- Expose data-level parallelism for vector (SIMD) instructions or super-scalar multiple instruction issue
- Mix pipeline with unrelated operations to help with reduce hazards
- Reduce loop “overhead”
- Makes code size larger

Amdahl's Law* applied to **dgemm**

- Measured **dgemm** performance
 - Peak 5.5 GFLOPS
 - Large matrices 3.6 GFLOPS
 - Processor 24.8 GFLOPS
- Why are we not getting (close to) 25 GFLOPS?
 - Something else (not floating-point ALU) is limiting performance!
 - But what? Possible culprits:
 - Cache
 - Hazards
 - Let's look at both!

** Amdahl's Law states that the amount a speedup possible through parallelism is limited by the sequential (non-parallel) work.*

Pipeline Hazards – dgemm

```
// AVX intrinsics; P&H p. 227
void dgemm_avx(int N, double *a, double *b, double *c) {
    // avx operates on 4 doubles in parallel
    for (int i=0; i<N; i+=4) {
        for (int j=0; j<N; j++) {
            // c0 = c[i][j]
            __m256d c0 = {0,0,0,0};
            for (int k=0; k<N; k++) {
                c0 = _mm256_add_pd(
                    c0, // c0 += a[i][k] * b[k][j]
                    _mm256_mul_pd(
                        _mm256_load_pd(a+i+k*N),
                        _mm256_broadcast_sd(b+k+j*N)));
            }
            _mm256_store_pd(c+i+j*N, c0); // c[i,j] = c0
        }
    }
}
```



“add_pd” depends on result of “mult_pd” which depends on “load_pd”

Loop Unrolling

```
// Loop unrolling; P&H p. 352
```

```
const int UNROLL = 4;
```

```
void dgemm_unroll(int n, double *A, double *B, double *C) {
```

```
    for (int i=0; i<n; i+= UNROLL*4) {
```

```
        for (int j=0; j<n; j++) {
```

```
            __m256d c[4];  4 registers
```

```
            for (int x=0; x<UNROLL; x++)
```

```
                c[x] = _mm256_load_pd(C+i+x*4+j*n);
```

```
            for (int k=0; k<n; k++) {
```

```
                __m256d b = _mm256_broadcast_sd(B+k+j*n);
```

```
                for (int x=0; x<UNROLL; x++)
```

```
                    c[x] = _mm256_add_pd(c[x],  Compiler does the unrolling
```

```
                        _mm256_mul_pd(_mm256_load_pd(A+n*k+x*4+i), b));
```

```
            }
```

```
        for (int x=0; x<UNROLL; x++)
```

```
            _mm256_store_pd(C+i+x*4+j*n, c[x]);
```

```
    }
```

```
}
```

```
}
```

How do you verify that the generated code is actually unrolled?

Performance

N	Gflops		
	scalar	avx	unroll
32	1.30	4.56	12.95
160	1.30	5.47	19.70
480	1.32	5.27	14.50
960	0.91	3.64	6.91

WOW!

?

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- **Memory access strategy - blocking**
- And in Conclusion, ...

FPU versus Memory Access

- How many floating-point operations does matrix multiply take?
 - $F = 2 \times N^3$ (N^3 multiplies, N^3 adds)
- How many memory load/stores?
 - $M = 3 \times N^2$ (for A, B, C)
- Many more floating-point operations than memory accesses
 - $q = F/M = 2/3 * N$
 - Good, since arithmetic is faster than memory access
 - Let's check the code ...

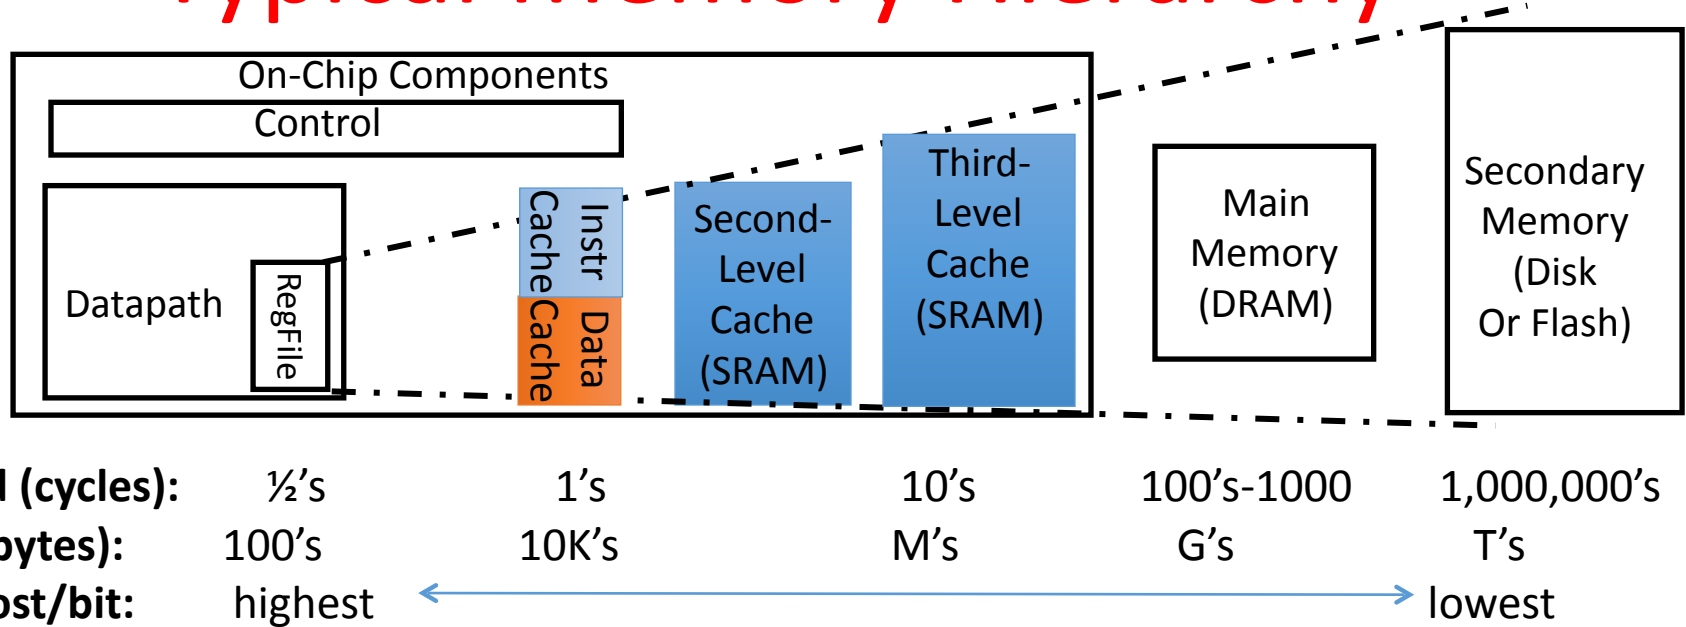
But memory is accessed repeatedly

Inner loop:

```
for (int k=0; k<N; k++) {  
    c0 = _mm256_add_pd(  
        c0, // c0 += a[i][k] * b[k][j]  
        _mm256_mul_pd(  
            _mm256_load_pd(a+i+k*N),  
            _mm256_broadcast_sd(b+k+j*N)));  
}
```

- $q = F/M = 1.6!$ (1.25 loads and 2 floating-point operations)

Typical Memory Hierarchy




- Where are the operands (A, B, C) stored?
- What happens as N increases?
- **Idea:** arrange that most accesses are to fast cache!

Blocking

- Idea:
 - Rearrange code to use values loaded in cache many times
 - Only “few” accesses to slow main memory (DRAM) per floating point operation
 - → throughput limited by FP hardware and cache, not slow DRAM
 - P&H, RISC-V edition p. 465

Blocking Matrix Multiply

(divide and conquer: sub-matrix multiplication)



Memory Access Blocking

```
// Cache blocking; P&H p. 555
const int BLOCKSIZE = 32;

void do_block(int n, int si, int sj, int sk, double *A, double *B, double *C) {
    for (int i=si; i<si+BLOCKSIZE; i+=UNROLL*4)
        for (int j=sj; j<sj+BLOCKSIZE; j++) {
            __m256d c[4];
            for (int x=0; x<UNROLL; x++)
                c[x] = _mm256_load_pd(C+i+x*4+j*n);
            for (int k=sk; k<sk+BLOCKSIZE; k++) {
                __m256d b = _mm256_broadcast_sd(B+k+j*n);
                for (int x=0; x<UNROLL; x++)
                    c[x] = _mm256_add_pd(c[x],
                                         _mm256_mul_pd(_mm256_load_pd(A+n+k+x*4+i), b));
            }
            for (int x=0; x<UNROLL; x++)
                _mm256_store_pd(C+i+x*4+j*n, c[x]);
        }
}

void dgemv_block(int n, double* A, double* B, double* C) {
    for(int sj=0; sj<n; sj+=BLOCKSIZE)
        for(int si=0; si<n; si+=BLOCKSIZE)
            for (int sk=0; sk<n; sk += BLOCKSIZE)
                do_block(n, si, sj, sk, A, B, C);
}
```

Performance

N	Gflops			
	scalar	avx	unroll	blocking
32	1.30	4.56	12.95	13.80
160	1.30	5.47	19.70	21.79
480	1.32	5.27	14.50	20.17
960	0.91	3.64	6.91	15.82

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

And in Conclusion, ...

- Approaches to Parallelism
 - SISD, SIMD, MIMD (next lecture)
- SIMD
 - One instruction operates on multiple operands simultaneously
- Example: matrix multiplication
 - Floating point heavy → exploit Moore's law to make fast