

CSI62

Operating Systems and Systems Programming

Lecture 2

Introduction to Processes


January 22nd, 2018

Profs. Anthony D. Joseph and Jonathan Ragan-Kelley

<http://cs162.eecs.Berkeley.edu>

Recall: What is an operating system?

- Special layer of software that provides application software access to hardware resources
 - Convenient abstraction of complex hardware devices
 - Protected access to shared resources
 - Security and authentication
 - Communication amongst logical entities



1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.2

Very Brief History of OS

- Several Distinct Phases:

Very Brief History of OS

- Several Distinct Phases:
 - Hardware Expensive, Humans Cheap
 - » Eniac, ... Multics



"I think there is a world market for maybe five computers." – *Thomas Watson, chairman of IBM, 1943*

1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.3

1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.4

Very Brief History of OS

- Several Distinct Phases:
 - Hardware Expensive, Humans Cheap
 - » Eniac, ... Multics



Thomas Watson was often called “the worlds greatest salesman” by the time of his death in 1956


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.5

Very Brief History of OS

- Several Distinct Phases:
 - Hardware Expensive, Humans Cheap
 - » Eniac, ... Multics
 - Hardware Cheaper, Humans Expensive
 - » PCs, Workstations, Rise of GUIs




1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.6

Very Brief History of OS

- Several Distinct Phases:
 - Hardware Expensive, Humans Cheap
 - » Eniac, ... Multics
 - Hardware Cheaper, Humans Expensive
 - » PCs, Workstations, Rise of GUIs
 - Hardware Really Cheap, Humans Really Expensive
 - » Ubiquitous devices, Widespread networking



1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.7

Very Brief History of OS

- Several Distinct Phases:
 - Hardware Expensive, Humans Cheap
 - » Eniac, ... Multics
 - Hardware Cheaper, Humans Expensive
 - » PCs, Workstations, Rise of GUIs
 - Hardware Really Cheap, Humans Really Expensive
 - » Ubiquitous devices, Widespread networking
- Rapid change in hardware leads to changing OS
 - Batch \Rightarrow Multiprogramming \Rightarrow Timesharing \Rightarrow Graphical UI \Rightarrow Ubiquitous Devices
 - Gradual migration of features into smaller machines
- Today
 - Small OS: 100K lines / Large: 10M lines (5M browser!)
 - 100-1000 people-years

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.8

OS Archaeology

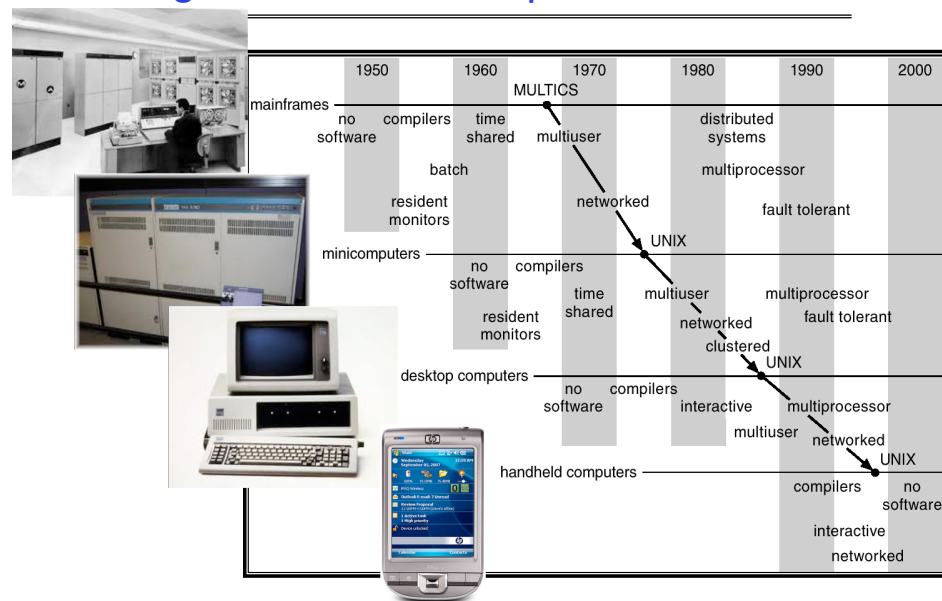
- Because of the cost of developing an OS from scratch, most modern OSes have a long lineage:
- Multics → AT&T Unix → BSD Unix → Ultrix, SunOS, NetBSD,...
- Mach (micro-kernel) + BSD → NextStep → XNU → Apple OS X, iPhone iOS, Watch OS
- MINIX → Linux → Android OS, Chrome OS, RedHat, Ubuntu, Fedora, Debian, Suse,...
- CP/M → QDOS → MS-DOS → Windows 3.1 → NT → 95 → 98 → 2000 → XP → Vista → 7 → 8 → 10 → Xbox One → ...

1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.9

Migration of OS Concepts and Features



1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.10

Today: Four Fundamental OS Concepts

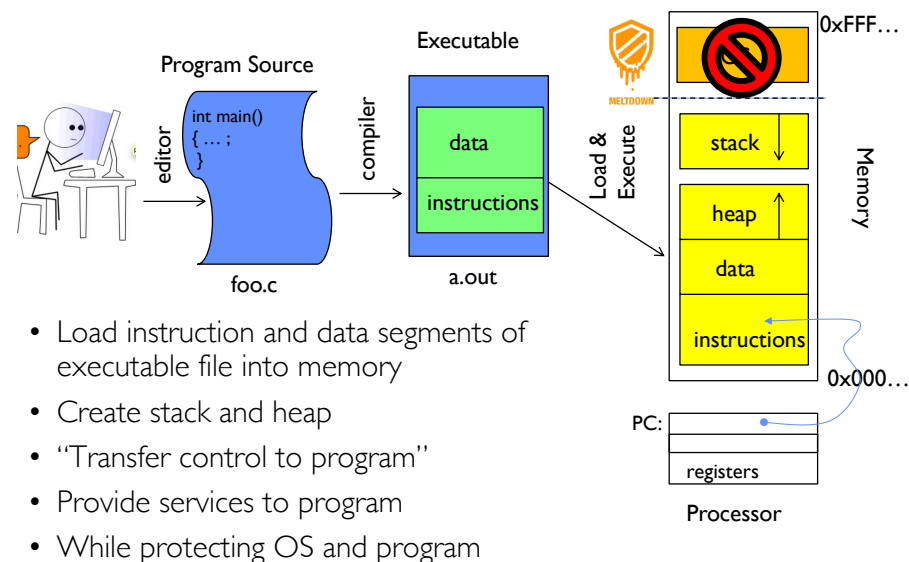
- Thread
 - Single unique execution context: fully describes program state
 - Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Programs execute in an *address space* that is distinct from the memory space of the physical machine
- Process
 - An instance of an executing program is a *process* consisting of an *address space* and one or more *threads of control*
- Dual mode operation / Protection
 - Only the "system" has the ability to access certain resources
 - The OS and the hardware are protected from user programs and user programs are isolated from one another by *controlling the translation* from program virtual addresses to machine physical addresses

1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.11

OS Bottom Line: Run Programs



- Load instruction and data segments of executable file into memory
- Create stack and heap
- "Transfer control to program"
- Provide services to program
- While protecting OS and program

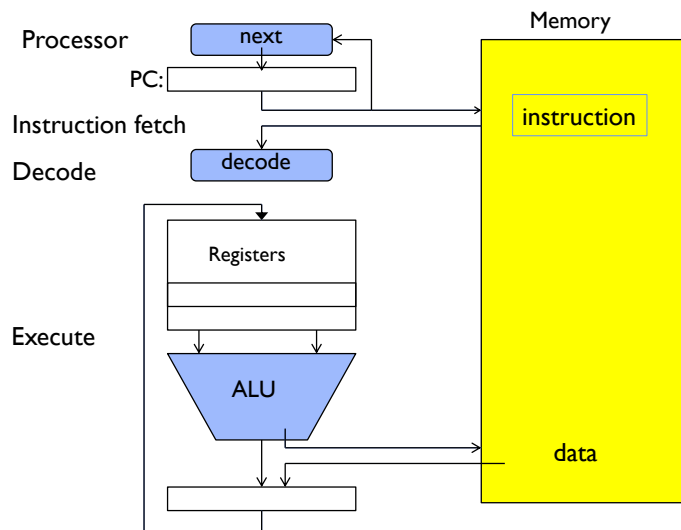
1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.12

Recall (61C): Instruction Fetch/Decode/Execute

The instruction cycle

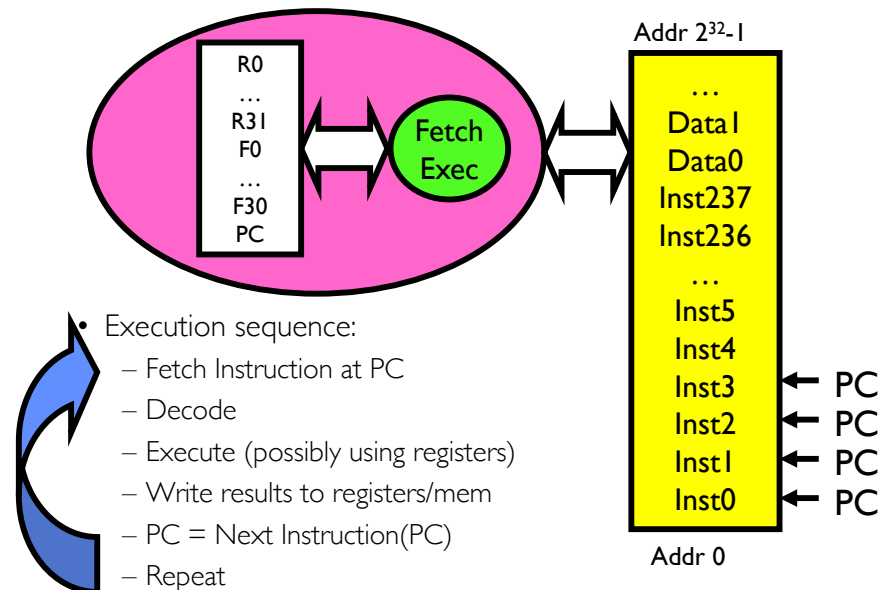


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.13

Recall (61C): What happens during program execution?



1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.14

First OS Concept: Thread of Control

- Certain registers hold the *context* of thread
 - Stack pointer holds the address of the top of stack
 - » Other conventions: Frame pointer, Heap pointer, Data
 - May be defined by the instruction set architecture or by compiler conventions
- Thread: Single unique execution context
 - Program Counter, Registers, Execution Flags, Stack
- A thread is executing on a processor when it is resident in the processor registers.
- PC register holds the address of executing instruction in the thread
- Registers hold the root state of the thread.
 - The rest is “in memory”


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.15

Second OS Concept: Program's Address Space

- Address space \Rightarrow the set of accessible addresses + state associated with them:
 - For a 32-bit processor there are $2^{32} = 4$ billion addresses
- What happens when you read or write to an address?
 - Perhaps nothing
 - Perhaps acts like regular memory
 - Perhaps ignores writes
 - Perhaps causes I/O operation
 - » (Memory-mapped I/O)
 - Perhaps causes exception (fault)

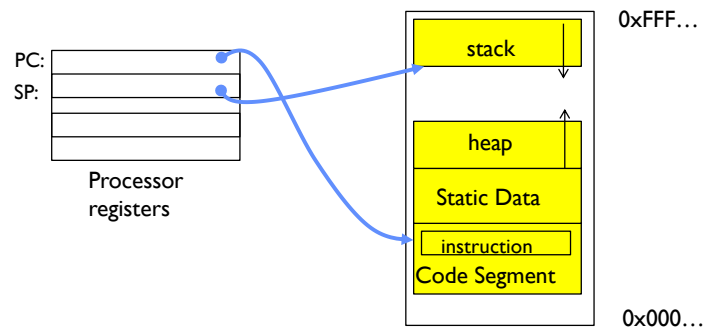


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.16

Address Space: In a Picture



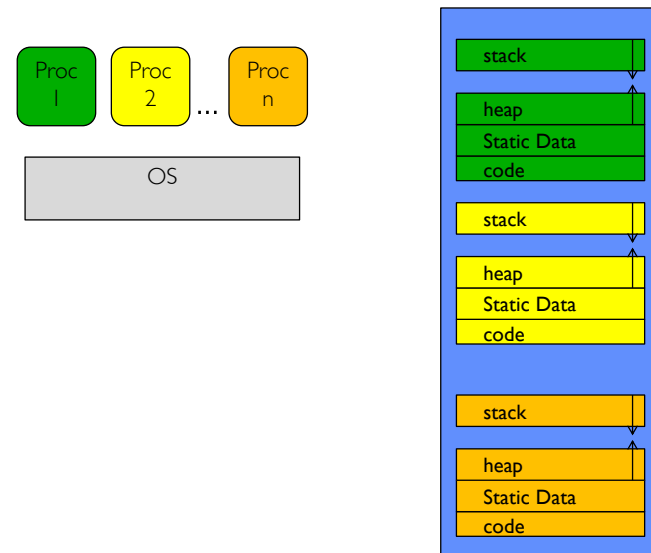
- What's in the code segment? Static data segment?
- What's in the Stack Segment?
 - How is it allocated? How big is it?
- What's in the Heap Segment?
 - How is it allocated? How big?

1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.17

Multiprogramming - Multiple Threads of Control



1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.18

Administrivia: Getting started


- Start homework 0 immediately ⇒ **Due next Monday (1/29)!**
 - cs162-xx account, Github account, registration survey
 - Vagrant and VirtualBox – VM environment for the course
 - » Consistent, managed environment on your machine
 - Get familiar with all the cs162 tools, submit to autograder via git
 - Homework slip days: **You have 3 slip days**
- **THIS Friday (1/26) is early drop day! Very hard to drop afterwards...**
- Should be going to section already!
- Group sign up form will be out after drop deadline
 - Work on finding groups ASAP: 4 people in a group!
 - Try to attend either same section or 2 sections by same TA

1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.19

How can we give the illusion of multiple processors?



- Assume a single processor. How do we provide the illusion of multiple processors?
 - Multiplex in time!
- Each virtual "CPU" needs a structure to hold:
 - Program Counter (PC), Stack Pointer (SP)
 - Registers (Integer, Floating point, others...?)
- How switch from one virtual CPU to the next?
 - Save PC, SP, and registers in current state block
 - Load PC, SP, and registers from new state block
- What triggers switch?
 - Timer, voluntary yield, I/O, other things

1/22/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 2.20

The Basic Problem of Concurrency

- The basic problem of concurrency involves resources:
 - Hardware: single CPU, single DRAM, single I/O devices
 - Multiprogramming API: processes think they have exclusive access to shared resources
- OS has to coordinate all activity
 - Multiple processes, I/O interrupts, ...
 - How can it keep all these things straight?
- Basic Idea: Use Virtual Machine abstraction
 - Simple machine abstraction for processes
 - Multiplex these abstract machines
- Dijkstra did this for the “THE system”
 - Few thousand lines vs 1 million lines in OS 360 (1K bugs)

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.21

Properties of this simple multiprogramming technique

- All virtual CPUs share same non-CPU resources
 - I/O devices the same
 - Memory the same
- Consequence of sharing:
 - Each thread can access the data of every other thread (good for sharing, bad for protection)
 - Threads can share instructions (good for sharing, bad for protection)
 - Can threads overwrite OS functions?
- This (unprotected) model is common in:
 - Embedded applications
 - Windows 3.1/Early Macintosh (switch only with yield)
 - Windows 95—ME (switch with both yield and timer)

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.22

Protection

- Operating System must protect itself from user programs
 - Reliability: compromising the operating system generally causes it to crash
 - Security: limit the scope of what processes can do
 - Privacy: limit each process to the data it is permitted to access
 - Fairness: each should be limited to its appropriate share of system resources (CPU time, memory, I/O, etc)
- It must protect User programs from one another
- Primary Mechanism: limit the translation from program address space to physical memory space
 - Can only touch what is mapped into process *address space*
- Additional Mechanisms:
 - Privileged instructions, in/out instructions, special registers
 - syscall processing, subsystem implementation
 - » (e.g., file access rights, etc)

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.23

Third OS Concept: Process

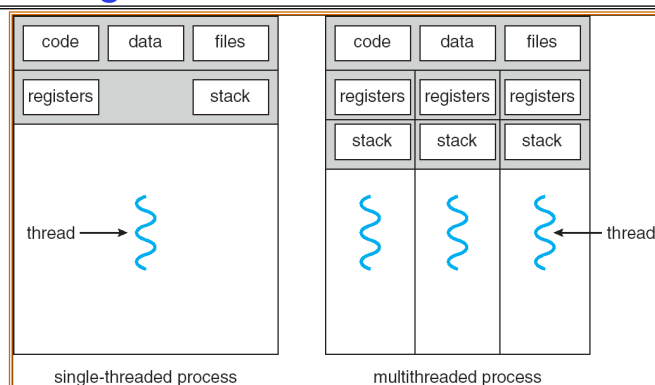
- **Process: execution environment with Restricted Rights**
 - **Address Space with One or More Threads**
 - Owns memory (address space)
 - Owns file descriptors, file system context, ...
 - Encapsulate one or more threads sharing process resources
- Why **processes**?
 - Protected from each other!
 - OS Protected from them
 - Processes provides memory protection
 - Threads more efficient than processes (later)
- Fundamental tradeoff between protection and efficiency
 - Communication easier *within* a process
 - Communication harder *between* processes
- Application instance consists of one or more processes

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.24

Single and Multithreaded Processes



- Threads encapsulate **concurrency**: “Active” component
- Address spaces encapsulate **protection**: “Passive” part
 - Keeps buggy program from trashing the system
- Why have multiple threads per address space?

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.25

Fourth OS Concept: Dual Mode Operation


- **Hardware** provides at least two modes:
 - “Kernel” mode (or “supervisor” or “protected”)
 - “User” mode: Normal programs executed
- What is needed in the hardware to support “dual mode” operation?
 - A bit of state (user/system mode bit)
 - Certain operations / actions only permitted in system/kernel mode
 - » In user mode they fail or trap
 - User → Kernel transition sets system mode AND saves the user PC
 - » Operating system code carefully puts aside user state then performs the necessary operations
 - Kernel → User transition *clears* system mode AND restores appropriate user PC
 - » return-from-interrupt

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.26

User/Kernel (Privileged) Mode



1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.27

Administrivia (Cont'd)

- Anthony's OH: **Wednesdays 1-3PM 465F Soda starting 1/31**
- Jonathan's OH : **Mondays 1-2:45PM 525 Soda starting 1/29**
- Avoid private Piazza posts – others have same question
- Three Free Online Textbooks:
 - Click on “Resources” link for a list of “Online Textbooks”
 - Can read O'Reilly books for free as long as on campus or VPN
 - » One book on Git, two books on C
- Webcast: <https://CalCentral.Berkeley.edu/> (CalNet sign in)
 - Webcast is ***NOT*** a replacement for coming to class!

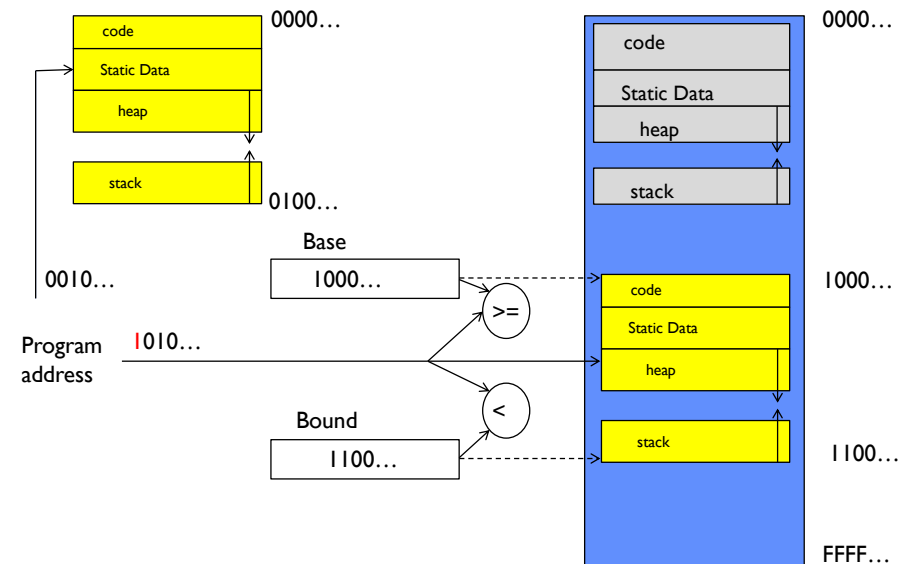
1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.28

5 min break

Simple Protection: Base and Bound (B&B)

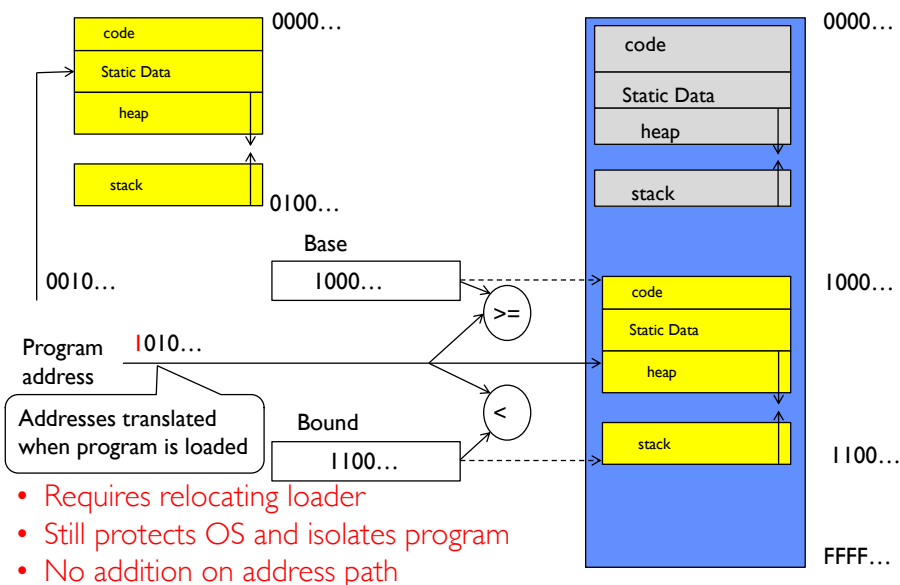


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.30

Simple Protection: Base and Bound (B&B)



- Requires relocating loader
- Still protects OS and isolates program
- No addition on address path


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.31

Another idea: Address Space Translation

- Program operates in an address space that is distinct from the physical memory space of the machine

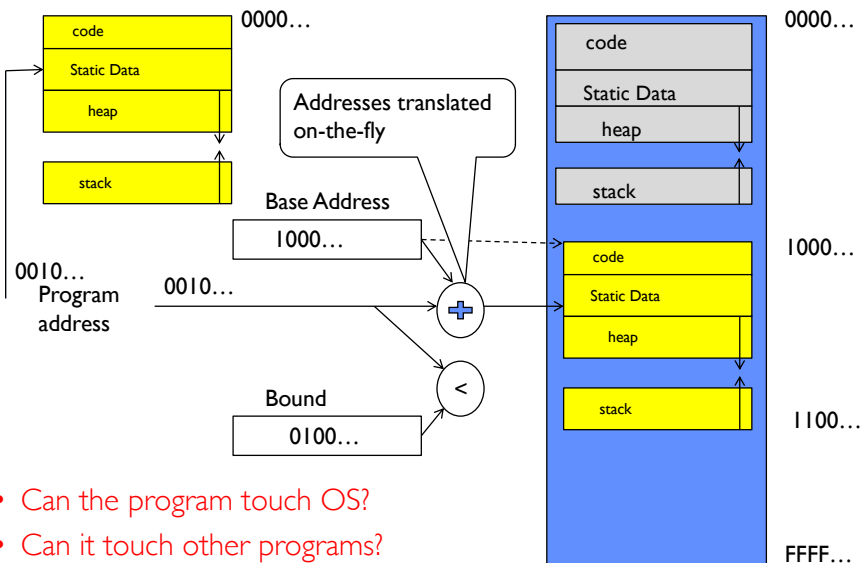


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.32

A simple address translation with Base and Bound



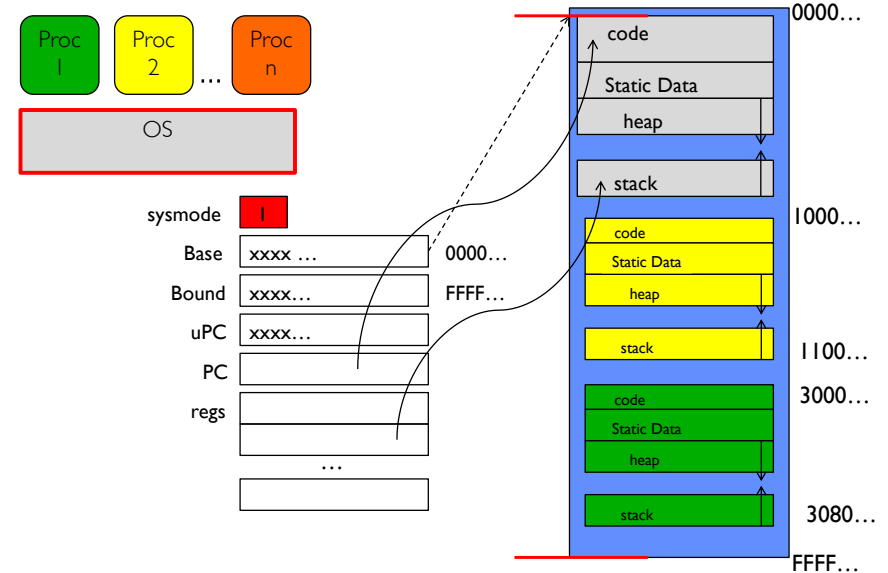
- Can the program touch OS?
- Can it touch other programs?

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.33

Tying it together: Simple B&B: OS loads process

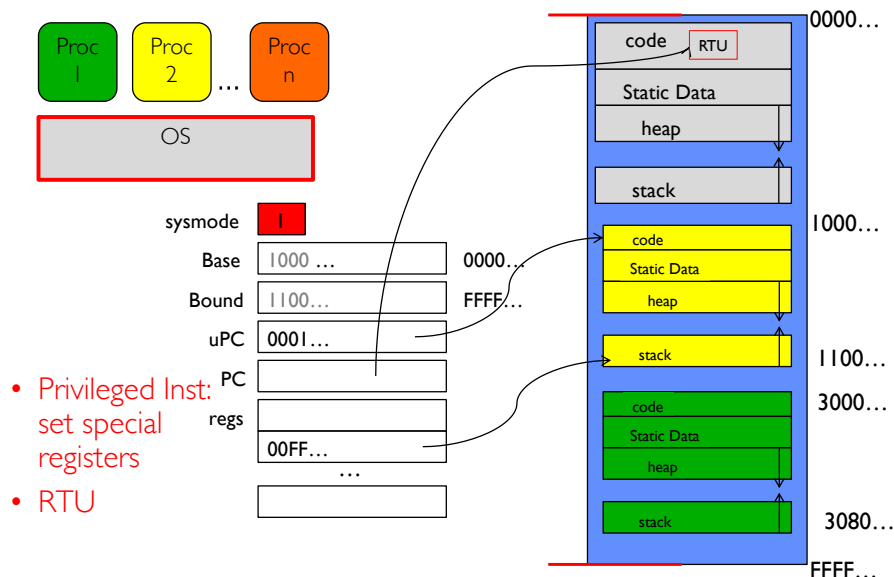


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.34

Simple B&B: OS gets ready to execute process



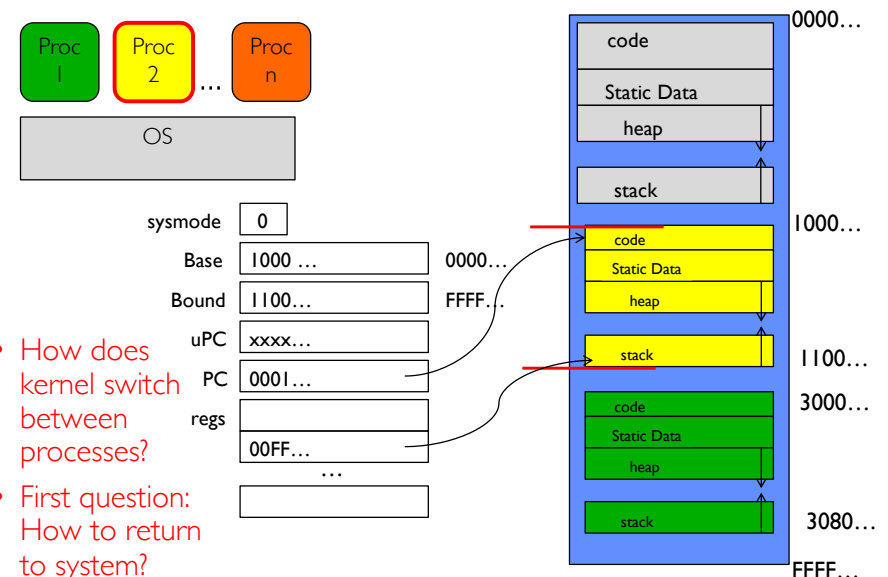
- Privileged Inst: set special registers
- RTU

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.35

Simple B&B: User Code Running



- How does kernel switch between processes?
- First question: How to return to system?

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

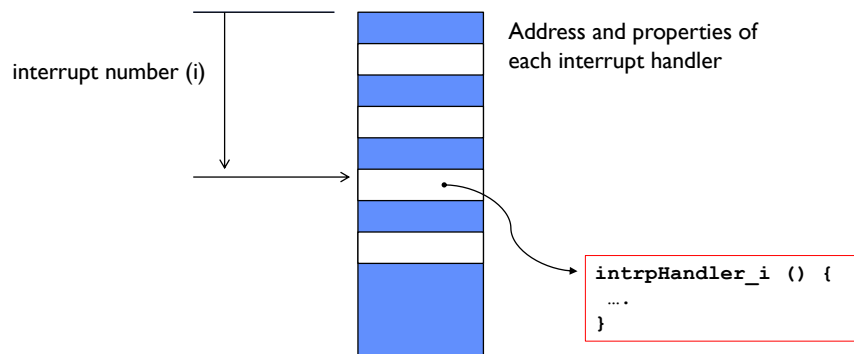
Lec 2.36

3 types of Mode Transfer

- Syscall
 - Process requests a system service, e.g., exit
 - Like a function call, but “outside” the process
 - Does not have the address of the system function to call
 - Like a Remote Procedure Call (RPC) – for later
 - Marshall the syscall id and args in registers and exec syscall
- Interrupt
 - External asynchronous event triggers context switch
 - e. g., Timer, I/O device
 - Independent of user process
- Trap or Exception
 - Internal synchronous event in process triggers context switch
 - e.g., Protection violation (segmentation fault), Divide by zero, ...
- All 3 are an UNPROGRAMMED CONTROL TRANSFER
 - Where does it go?

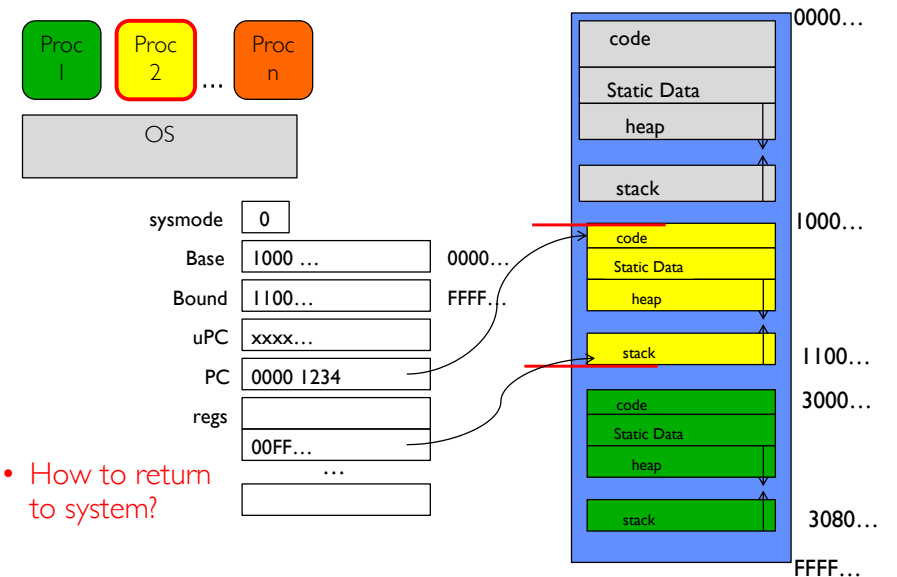
How do we get the system target address of the “unprogrammed control transfer?”

Interrupt Vector



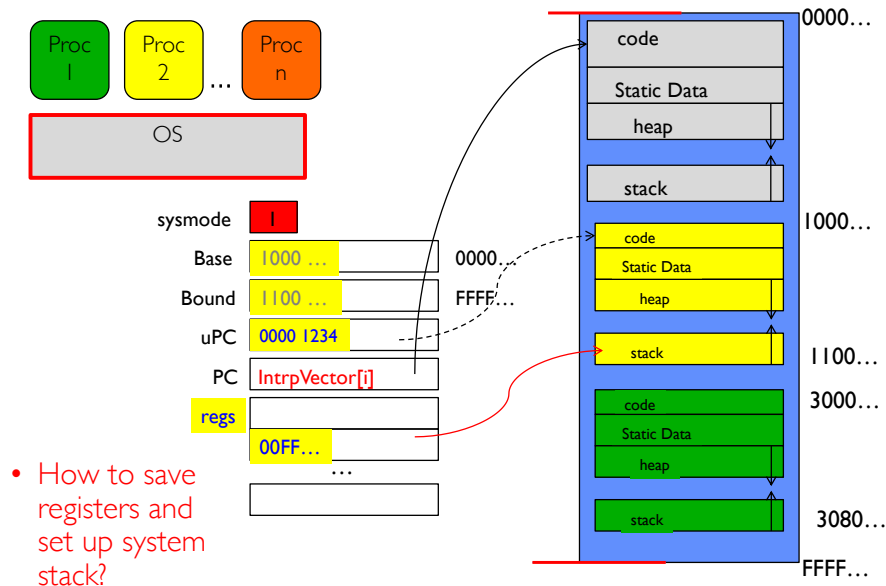
- Where else do you see this dispatch pattern?

Simple B&B: User => Kernel



- How to return to system?

Simple B&B: Interrupt

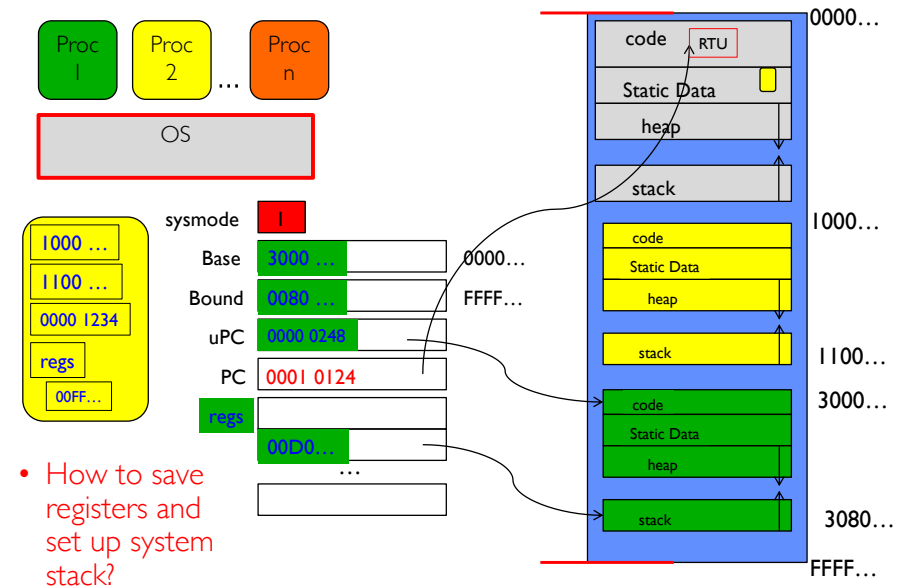


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.41

Simple B&B: Switch User Process

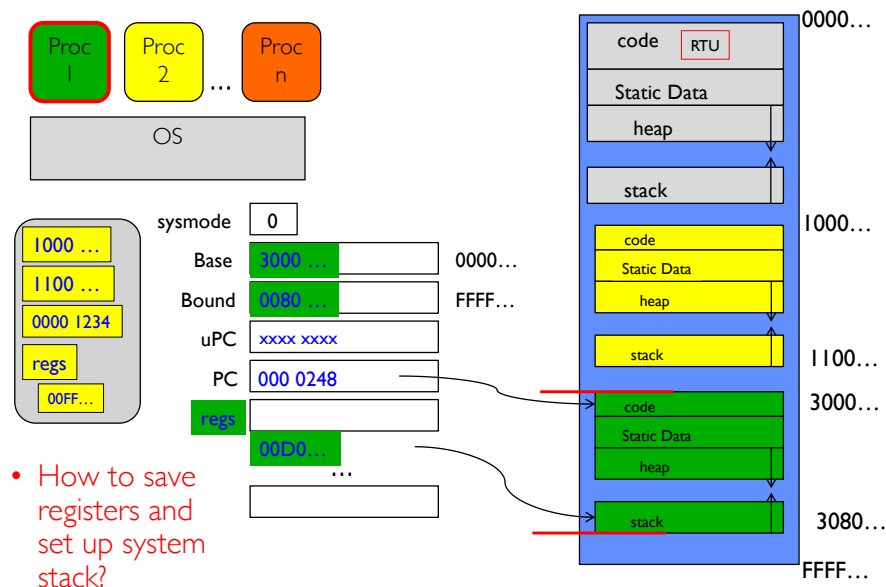


1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.42

Simple B&B: "resume"



1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.43

Conclusion: Four fundamental OS concepts

- Thread
 - Single unique execution context
 - Program Counter, Registers, Execution Flags, Stack
- Address Space with Translation
 - Programs execute in an *address space* that is distinct from the memory space of the physical machine
- Process
 - An instance of an executing program is a *process* consisting of an *address space* and one or more *threads of control*
- Dual Mode operation/Protection
 - Only the "system" has the ability to access certain resources
 - The OS and the hardware are protected from user programs and user programs are isolated from one another by *controlling the translation* from program virtual addresses to machine physical addresses

1/22/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 2.44