

CSI62

Operating Systems and Systems Programming

Lecture 3

Processes (cont'd), Fork

January 22nd, 2018

Profs. Anthony D. Joseph and Jonathan Ragan-Kelley

<http://cs162.eecs.berkeley.edu>

Recall: Four fundamental OS concepts

- Thread
 - Single unique execution context
 - Program Counter, Registers, Execution Flags, Stack
- Address Space w/ translation
 - Programs execute in an *address space* that is distinct from the memory space of the physical machine
- Process
 - An instance of an executing program is a *process consisting of an address space and one or more threads of control*
- Dual Mode operation/Protection
 - Only the “system” has the ability to access certain resources
 - The OS and the hardware are protected from user programs and user programs are isolated from one another by *controlling the translation* from program virtual addresses to machine physical addresses

1/24/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 3.2

Recall: 3 types of Mode Transfer

- Syscall
 - Process requests a system service, e.g., exit
 - Like a function call, but “outside” the process
 - Does not have the address of the system function to call
 - Like a Remote Procedure Call (RPC) – for later
 - Marshall the syscall id and args in registers and exec syscall
- Interrupt
 - External asynchronous event triggers context switch
 - e. g., Timer, I/O device
 - Independent of user process
- Trap or Exception
 - Internal synchronous event in process triggers context switch
 - e.g., Protection violation (segmentation fault), Divide by zero, ...
- All 3 are an UNPROGRAMMED CONTROL TRANSFER
 - Where does it go?

1/24/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 3.3


How do we get the system target address of the “unprogrammed control transfer?”

1/24/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 3.4

Interrupt Vector



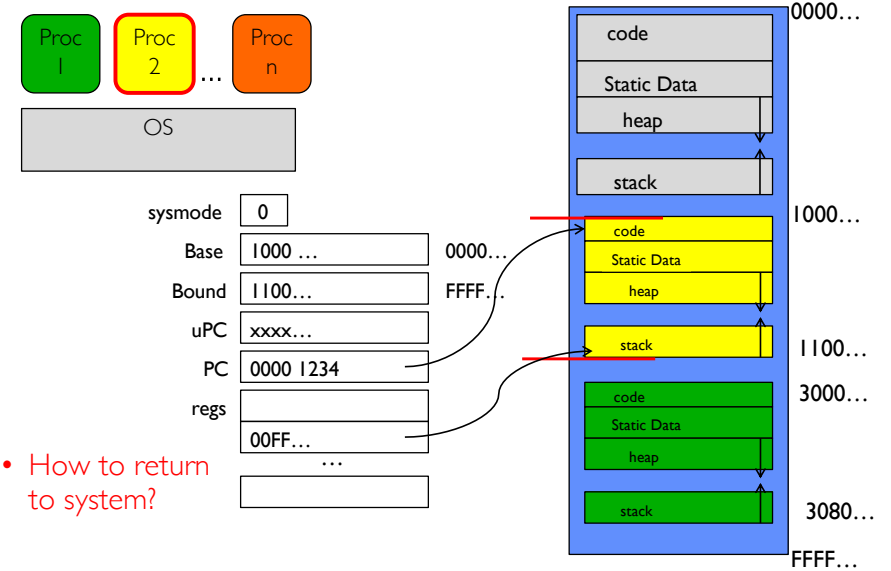
- Where else do you see this dispatch pattern?

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.5

Simple B&B: User => Kernel



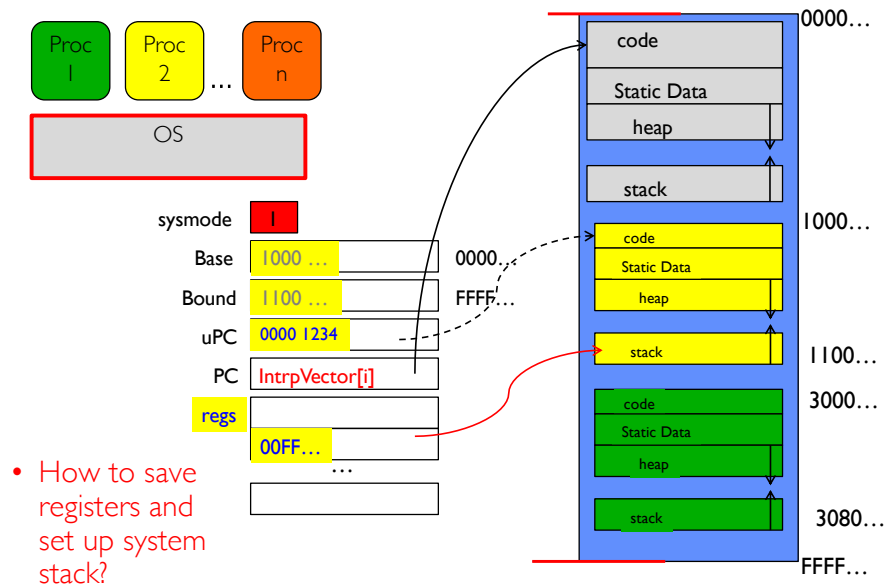
- How to return to system?

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.6

Simple B&B: Interrupt



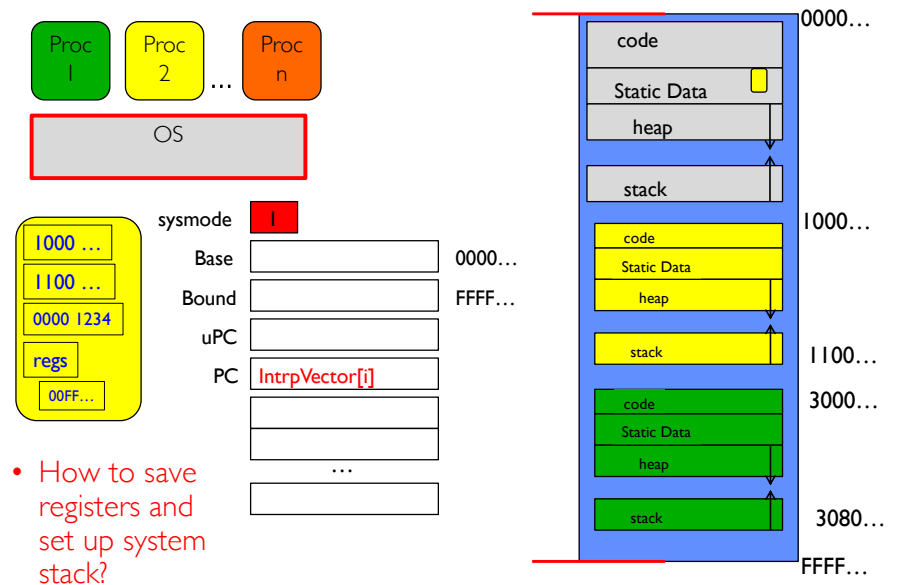
- How to save registers and set up system stack?

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.7

Simple B&B: Save state



- How to save registers and set up system stack?

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018


Lec 3.8

Process Control Block

(Assume single threaded processes for now)

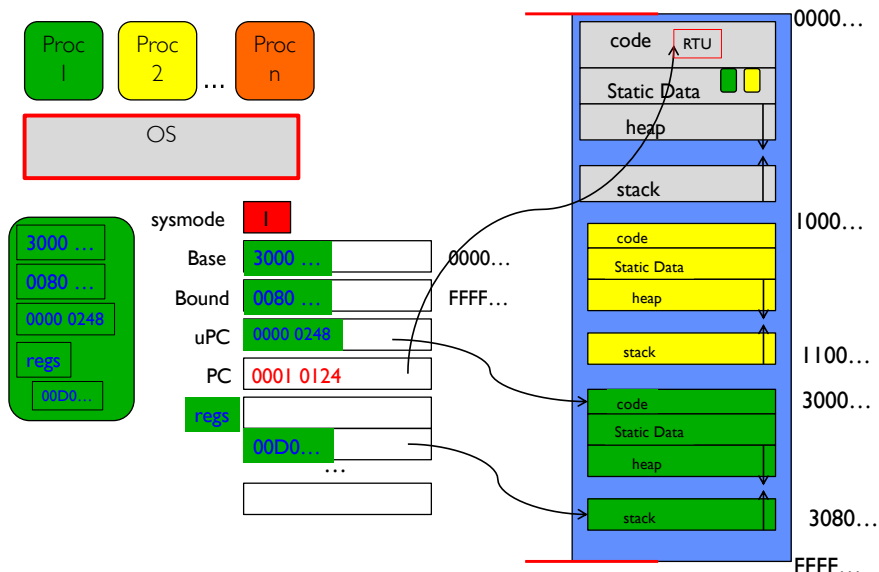
- Kernel represents each process as a process control block (PCB)
 - Status (running, ready, blocked, ...)
 - Registers, SP, ... (when not running)
 - Process ID (PID), User, Executable, Priority, ...
 - Execution time, ...
 - Memory space, translation tables, ...
- Kernel Scheduler maintains a data structure containing the PCBs
- Scheduling algorithm selects the next one to run

Recall: give the illusion of multiple processors?

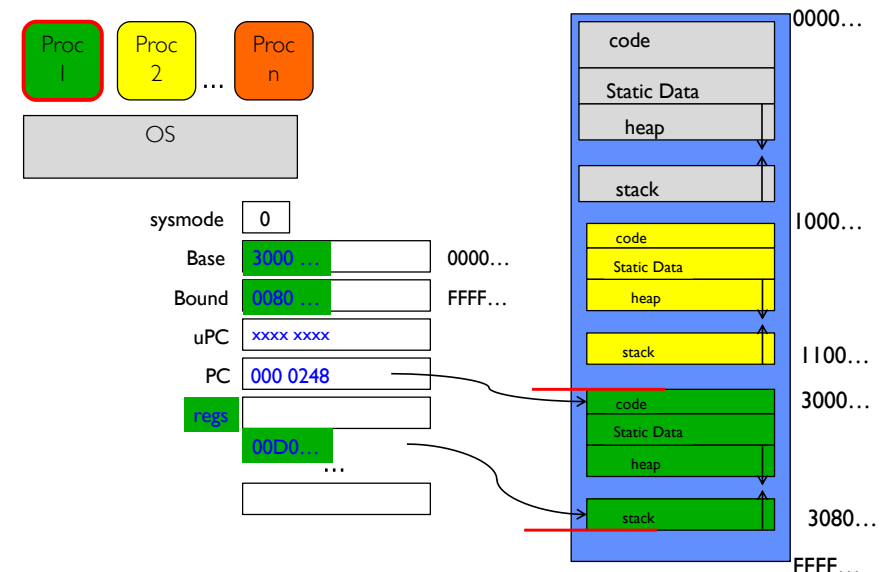


- Assume a single processor. How do we provide the *illusion* of multiple processors?
 - Multiplex in time!
 - Multiple “virtual CPUs”
- Each virtual “CPU” needs a structure to hold, i.e., **PCB**:
 - Program Counter (PC), Stack Pointer (SP)
 - Registers (Integer, Floating point, others...?)
- How switch from one virtual CPU to the next?
 - Save PC, SP, and registers in current **PCB**
 - Load PC, SP, and registers from new **PCB**
- What triggers switch?
 - Timer, voluntary yield, I/O, other things

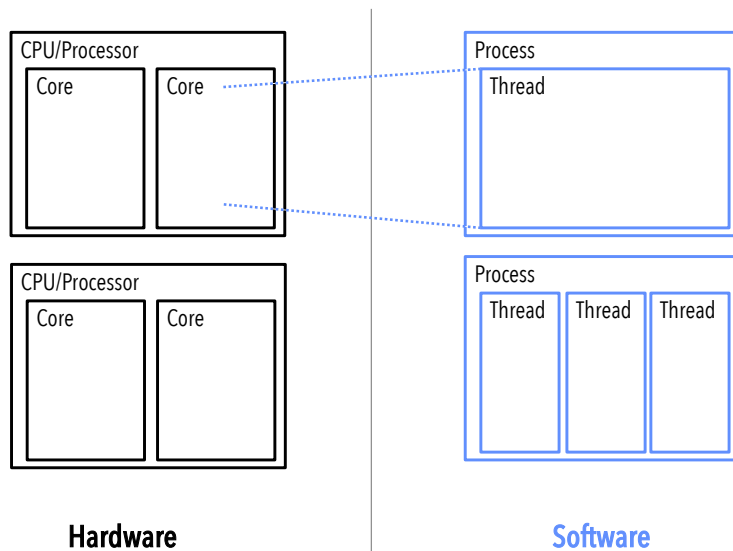
Simple B&B: Switch User Process



Simple B&B: “resume”



CPU, Processor, Core, Process, Thread!?



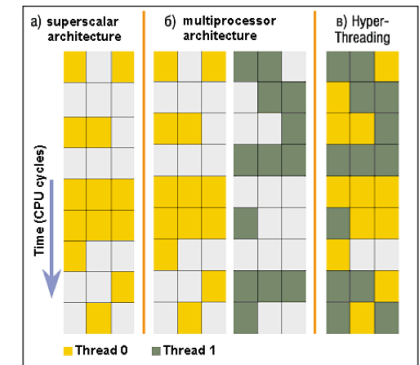
1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.13

Simultaneous MultiThreading/Hyperthreading

- Hardware technique
 - Superscalar processors can execute multiple instructions that are independent
 - Hyperthreading **duplicates register state** to make a second “thread,” allowing more instructions to run
- Can schedule each thread as if were separate CPU
 - But, sub-linear speedup!
- Original technique called “Simultaneous Multithreading”
 - <http://www.cs.washington.edu/research/smt/index.html>
 - SPARC, Pentium 4/Xeon (“Hyperthreading”), Power 5



Colored blocks show instructions executed

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.14

Scheduler

```

if ( readyProcesses(PCBs) ) {
    nextPCB = selectProcess(PCBs);
    run( nextPCB );
} else {
    run_idle_process();
}
    
```


- Scheduling: Mechanism for deciding which processes/threads receive the CPU
- Lots of different scheduling policies provide ...
 - Fairness or
 - Real-time guarantees or
 - Latency optimization or ..

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.15

Putting it together: web server

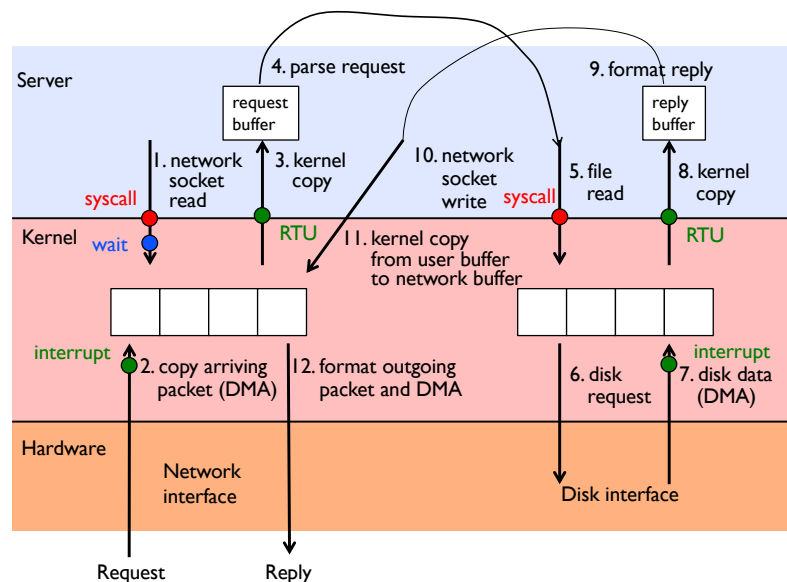


1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.16

Putting it together: web server



1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.17

Recall: 3 types of Kernel Mode Transfer


- Syscall
 - Process requests a system service, e.g., exit
 - Like a function call, but “outside” the process
 - Does not have the address of the system function to call
 - Like a Remote Procedure Call (RPC) – for later
 - Marshall the syscall ID and arguments in registers and execute syscall
- Interrupt
 - External asynchronous event triggers context switch
 - e.g., Timer, I/O device
 - Independent of user process
- Trap or Exception
 - Internal synchronous event in process triggers context switch
 - e.g., Protection violation (segmentation fault), Divide by zero, ...

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.18

Recall: User/Kernel (Privileged) Mode



1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.19

Implementing Safe Kernel Mode Transfers

- Important aspects:
 - Separate kernel stack
 - Controlled transfer into kernel (e.g., syscall table)
- Carefully constructed kernel code packs up the user process state and sets it aside
 - Details depend on the machine architecture
- Should be impossible for buggy or malicious user program to cause the kernel to corrupt itself

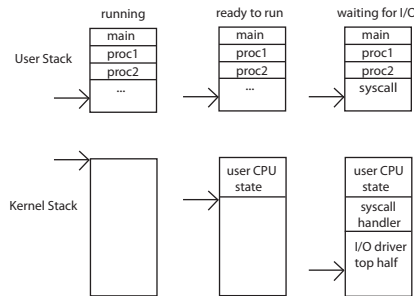
1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.20

Need for Separate Kernel Stacks

- Kernel needs space to work
- Cannot put anything on the user stack (Why?)
- Two-stack model
 - OS thread has interrupt stack (located in kernel memory) plus User stack (located in user memory)
 - Syscall handler copies user args to kernel space before invoking specific function (e.g., open)
 - Interrupts (???)



1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.21

Kernel System Call Handler

- Vector through well-defined syscall entry points!
 - Table mapping system call number to handler
- Locate arguments
 - In registers or on user (!) stack
- Copy arguments
 - From user memory into kernel memory
 - Protect kernel from malicious code evading checks
- Validate arguments
 - Protect kernel from errors in user code
- Copy results back
 - Into user memory

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.24

Hardware support: Interrupt Control

- Interrupt processing not visible to the user process:
 - Occurs between instructions, restarted transparently
 - No change to process state
 - What can be observed even with perfect interrupt processing?
- Interrupt Handler invoked with interrupts 'disabled'
 - Re-enabled upon completion
 - Non-blocking (run to completion, no waits)
 - Pack up in a queue and pass off to an OS thread for hard work
 - » wake up an existing OS thread

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.25

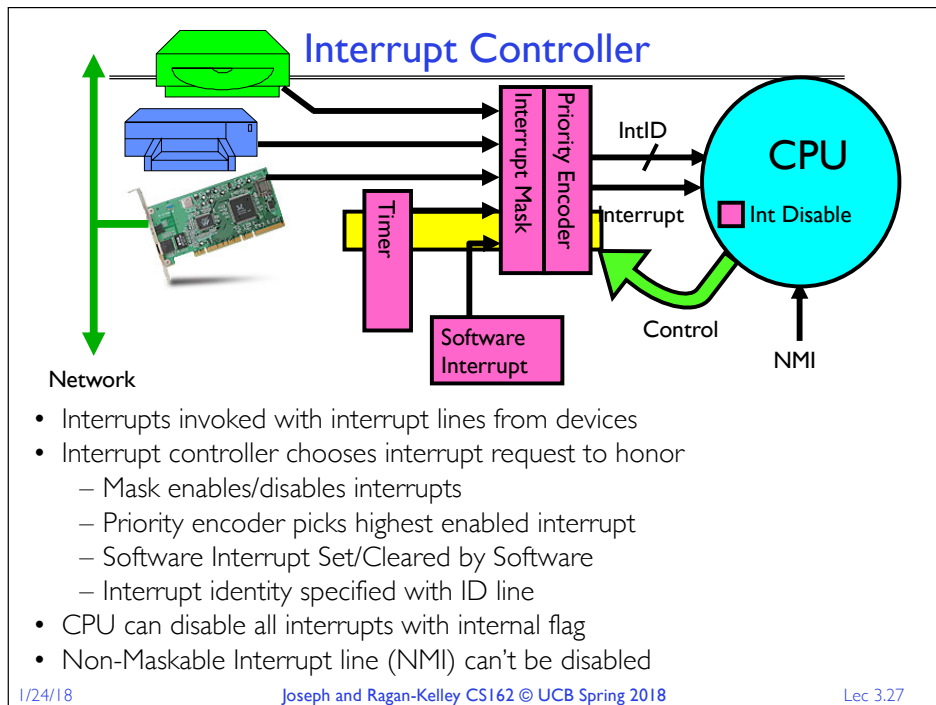
Hardware support: Interrupt Control

- OS kernel may enable/disable interrupts
 - On x86: CLI (disable interrupts), STI (enable)
 - Atomic section when select next process/thread to run
 - Atomic return from interrupt or syscall
- HW may have multiple levels of interrupt
 - Mask off (disable) certain interrupts, eg., lower priority
 - Certain Non-Maskable-Interrupts (NMI)
 - » e.g., kernel segmentation fault

1/24/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 3.26



How do we take interrupts safely?

- **Interrupt vector**
 - Limited number of entry points into kernel
- Kernel interrupt stack
 - Handler works regardless of state of user code
- Interrupt masking
 - Handler is non-blocking
- Atomic transfer of control
 - “Single instruction”-like to change:
 - » Program counter
 - » Stack pointer
 - » Memory protection
 - » Kernel/user mode
- Transparent restartable execution
 - User program does not know interrupt occurred

1/24/18 Joseph and Ragan-Kelley CS162 © UCB Spring 2018 Lec 3.28

Administrivia: Getting started

- **THIS** Friday (1/26) is early drop day! Very hard to drop afterwards...
- Get working on Homework 0 **due on Monday by 11:59PM!**
 - Get familiar with all the cs162 tools
 - Submit to autograder via git
- Participation: Attend section! Get to know your TA!
- Group sign up: now via autograder, then TA form (next week, after EDD)
 - Get finding groups of 4 people ASAP
 - Priority for same section; if cannot make this work, keep same TA

1/24/18 Joseph and Ragan-Kelley CS162 © UCB Spring 2018 Lec 3.29