

CSI62

Operating Systems and Systems Programming

Lecture 6

Concurrency (Continued), Thread and Processes

February 5th, 2018

Profs. Anthony D. Joseph and Jonathan Ragan-Kelley


<http://cs162.eecs.Berkeley.edu>

Recall: Use of Threads

- Version of program with Threads (loose syntax):

```
main() {  
    ThreadFork(ComputePI, "pi.txt");  
    ThreadFork(PrintClassList, "classlist.txt");  
}
```

- What does **ThreadFork()** do?
 - Start independent thread running given procedure
- What is the behavior here?
 - Now, you would actually see the class list
 - This *should* behave as if there are two separate CPUs




2/5/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 6.2

Recall: Memory Footprint: Two-Threads

- If we stopped this program and examined it with a debugger, we would see
 - Two sets of CPU registers
 - Two sets of Stacks
- Questions:
 - How do we position stacks relative to each other?
 - What maximum size should we choose for the stacks?
 - What happens if threads violate this?
 - How might you catch violations?



2/5/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 6.3

Recall: Dispatch Loop

- Conceptually, the dispatching loop of the operating system looks as follows:

```
Loop {  
    RunThread();  
    newTCB = ChooseNextThread();  
    SaveStateOfCPU(curTCB);  
    LoadStateOfCPU(newTCB);  
}
```

- This is an *infinite* loop
 - One could argue that this is all that the OS does
- Should we ever exit this loop???
 - When would that be?

2/5/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 6.4

Saving/Restoring state (often called “Context Switch”)

```
switch(tCur,tNew) {
    /* Unload old thread */
    TCB[tCur].regs.r7 = CPU.r7;
    ...
    TCB[tCur].regs.r0 = CPU.r0;
    TCB[tCur].regs.sp = CPU.sp;
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/

    /* Load and execute new thread */
    CPU.r7 = TCB[tNew].regs.r7;
    ...
    CPU.r0 = TCB[tNew].regs.r0;
    CPU.sp = TCB[tNew].regs.sp;
    CPU.retpc = TCB[tNew].regs.retpc;
    return; /* Return to CPU.retpc */
}
```

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.5

Switch Details (continued)

- What if you make a mistake in implementing switch?
 - Suppose you forget to save/restore register 32
 - Get intermittent failures depending on when context switch occurred and whether new thread uses register 32
 - System will give wrong result without warning
- Can you devise an exhaustive test to test switch code?
 - No! Too many combinations and inter-leavings
- Cautionary tale:
 - For speed, Topaz kernel saved one instruction in switch()
 - Carefully documented! Only works as long as kernel size < 1MB
 - What happened?
 - » Time passed, People forgot
 - » Later, they added features to kernel (no one removes features!)
 - » Very weird behavior started happening
 - Moral of story: Design for simplicity

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.6

Some Numbers

- Frequency of performing context switches: 10-100ms
 - Context switch time in Linux: 3-4 μ secs (Intel i7 & E5)
 - Thread switching faster than process switching (100 ns)
 - But switching across cores $\sim 2x$ more expensive than within-core
 - Context switch time increases sharply with size of working set*
 - Can increase 100x or more
- *The working set is subset of memory used by process in a time window
- Moral: context switching depends mostly on cache limits and the process or thread's hunger for memory

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.7

Some Numbers

- Many process are multi-threaded, so thread context switches may be either **within-process** or **across-processes**

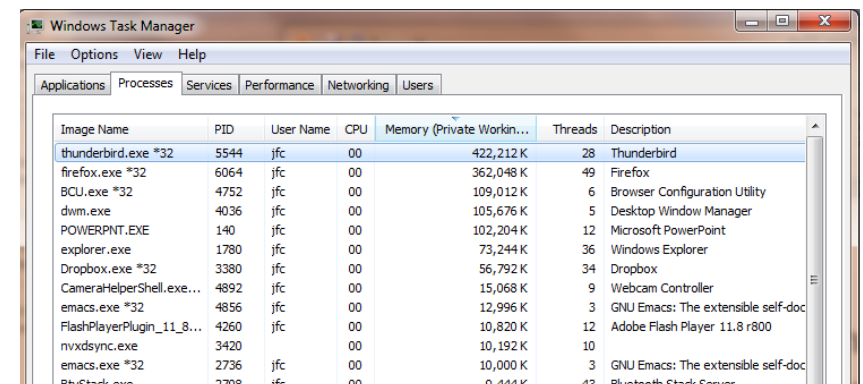



Image Name	PID	User Name	CPU	Memory (Private Workin...	Threads	Description
thunderbird.exe *32	5544	jfc	00	422,212 K	28	Thunderbird
firefox.exe *32	6064	jfc	00	362,048 K	49	Firefox
BCU.exe *32	4752	jfc	00	109,012 K	6	Browser Configuration Utility
dwm.exe	4036	jfc	00	105,676 K	5	Desktop Window Manager
POWERPNT.EXE	140	jfc	00	102,204 K	12	Microsoft PowerPoint
explorer.exe	1780	jfc	00	73,244 K	36	Windows Explorer
Dropbox.exe *32	3380	jfc	00	56,792 K	34	Dropbox
CameraHelperShell.exe...	4892	jfc	00	15,068 K	9	Webcam Controller
emacs.exe *32	4856	jfc	00	12,996 K	3	GNU Emacs: The extensible self-doc
FlashPlayerPlugin_11_8...	4260	jfc	00	10,820 K	12	Adobe Flash Player 11.8 r800
nvxdsync.exe	3420		00	10,192 K	10	
emacs.exe *32	2736	jfc	00	10,000 K	3	GNU Emacs: The extensible self-doc
BtvStack.exe	2708	jfc	00	9,444 K	43	Bluetooth Stack Server

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.8

What happens when thread blocks on I/O?



- What happens when a thread requests a block of data from the file system?
 - User code invokes a system call
 - Read operation is initiated
 - Run new thread/switch
- Thread communication similar
 - Wait for Signal/Join
 - Networking

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.9

External Events

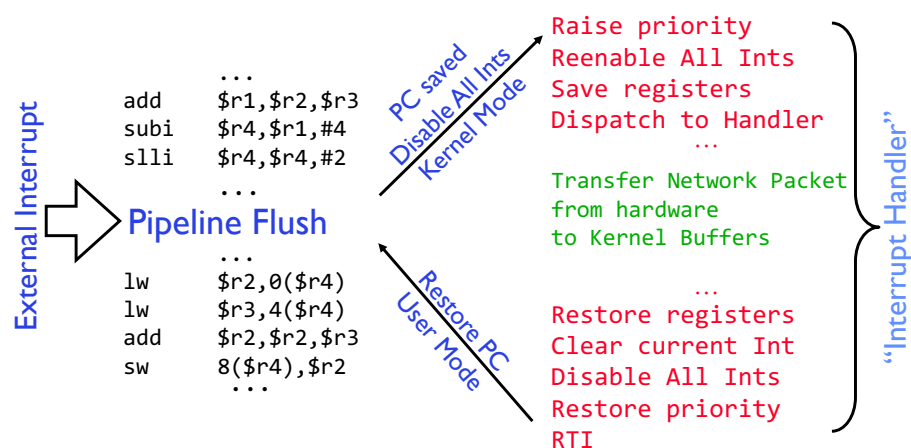
- What happens if thread never does any I/O, never waits, and never yields control?
 - Could the **ComputePI** program grab all resources and never release the processor?
 - » What if it didn't print to console?
 - Must find way that dispatcher can regain control!
- Answer: utilize external events
 - Interrupts: signals from hardware or software that stop the running code and jump to kernel
 - Timer: like an alarm clock that goes off every some milliseconds
- If we make sure that external events occur frequently enough, can ensure dispatcher runs

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.10

Example: Network Interrupt



- An interrupt is a hardware-invoked context switch
 - No separate step to choose what to run next
 - Always run the interrupt handler immediately


2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.11

Use of Timer Interrupt to Return Control

- Solution to our dispatcher problem
 - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:


```
TimerInterrupt() {
    DoPeriodicHouseKeeping();
    run_new_thread();
}
```

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.12

Thread Abstraction




- Illusion: Infinite number of processors

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.13

Thread Abstraction



- Illusion: Infinite number of processors
- Reality: Threads execute with variable speed
 - Programs must be designed to work with any schedule

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.14

Programmer vs. Processor View

Programmer's View	Possible Execution #1
.	.
.	.
.	.
x = x + 1;	x = x + 1;
y = y + x;	y = y + x;
z = x + 5y;	z = x + 5y;
.	.
.	.
.	.

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.15

Programmer vs. Processor View

Programmer's View	Possible Execution #1	Possible Execution #2
.	.	.
.	.	.
.	.	.
x = x + 1;	x = x + 1;	x = x + 1
y = y + x;	y = y + x;
z = x + 5y;	z = x + 5y;	thread is suspended
.	.	other thread(s) run
.	.	thread is resumed
.
		y = y + x
		z = x + 5y

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.16

Programmer vs. Processor View


Programmer's View	Possible Execution #1	Possible Execution #2	Possible Execution #3
.	.	.	.
.	.	.	.
.	.	.	.
$x = x + 1;$	$x = x + 1;$	$x = x + 1$	$x = x + 1$
$y = y + x;$	$y = y + x;$	$y = y + x$
$z = x + 5y;$	$z = x + 5y;$	thread is suspended
.	.	other thread(s) run	thread is suspended
.	.	thread is resumed	other thread(s) run
.	thread is resumed
		$y = y + x$
		$z = x + 5y$	$z = x + 5y$

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018


Lec 6.17

Possible Executions



a) One execution

b) Another execution



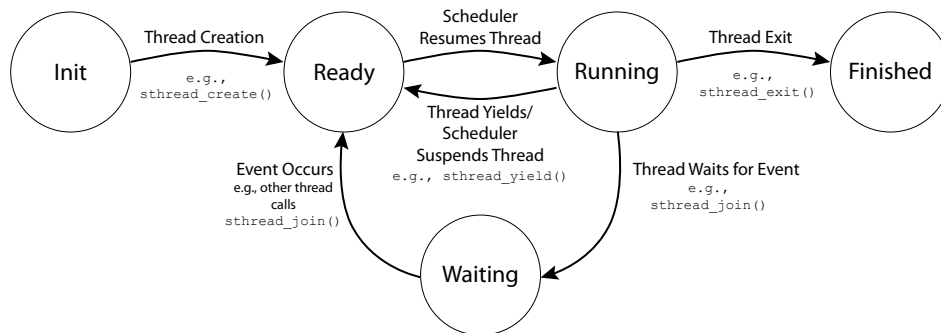
c) Another execution

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.18

Thread Lifecycle



2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.19

Administrivia

- Your section is your home for CS162
 - The TA needs to get to know you to judge participation
 - All design reviews will be conducted by your TA
 - You can attend alternate section by same TA, but try to keep the amount of such cross-section movement to a minimum
- First midterm: Wed Feb 28 6:30 – 8:30 PM
 - ROOM ASSIGNMENTS TBD
 - LET US KNOW DSP AND ACADEMIC CONFLICTS ASAP
 - 245 Li Ka Shing
 - 20 Barrows Hall
 - A1 Hearst Field Annex
 - 2060 Valley Life Sciences Building
 - 102 Wurster Hall

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.20

BREAK

Per Thread Descriptor (Kernel Supported Threads)

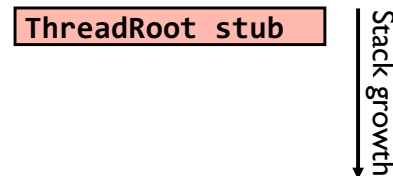
- Each Thread has a *Thread Control Block (TCB)*
 - Execution State: CPU registers, program counter (PC), pointer to stack (SP)
 - Scheduling info: state, priority, CPU time
 - Various Pointers (for implementing scheduling queues)
 - Pointer to enclosing process (PCB) – user threads
 - Etc (add stuff as you find a need)
- OS Keeps track of TCBs in “kernel memory”
 - In Array, or Linked List, or ...
 - I/O state (file descriptors, network connections, etc)

ThreadFork(): Create a New Thread

- **ThreadFork()** is a user-level procedure that creates a new thread and places it on ready queue
- Arguments to **ThreadFork()**
 - Pointer to application routine (**fcnPtr**)
 - Pointer to array of arguments (**fcnArgPtr**)
 - Size of stack to allocate
- Implementation
 - Sanity check arguments
 - Enter Kernel-mode and Sanity Check arguments again
 - Allocate new Stack and TCB
 - Initialize TCB and place on ready list (Runnable)


How do we initialize TCB and Stack?

- Initialize Register fields of TCB
 - Stack pointer made to point at stack
 - PC return address \Rightarrow OS (asm) routine **ThreadRoot()**
 - Two arg registers (a0 and a1) initialized to **fcnPtr** and **fcnArgPtr**, respectively
- Initialize stack data?
 - No. Important part of stack frame is in registers (ra)
 - Think of stack frame as just before body of **ThreadRoot()** really gets started



Initial Stack

How does Thread get started?




- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
 - This really starts the new thread

What does ThreadRoot() look like?

- `ThreadRoot()` is the root for the thread routine:


```
ThreadRoot() {
    DoStartupHousekeeping();
    UserModeSwitch(); /* enter user mode */
    Call fcnPtr(fcnArgPtr);
    ThreadFinish();
}
```

- Startup Housekeeping
 - Includes things like recording start time of thread
 - Other statistics
- Stack will grow and shrink with execution of thread
- Final return from thread returns into `ThreadRoot()` which calls `ThreadFinish()`
 - `ThreadFinish()` wake up sleeping threads



Multithreaded Processes

- Process Control Block (PCBs) points to multiple Thread Control Blocks (TCBs):



- Switching threads within a block is a simple thread switch
- Switching threads across blocks requires changes to memory and I/O address tables

Examples multithreaded programs

- Embedded systems
 - Elevators, planes, medical systems, smart watches
 - Single program, concurrent operations
- Most modern OS kernels
 - Internally concurrent because have to deal with concurrent requests by multiple users
 - But no protection needed within kernel
- Database servers
 - Access to shared data by many concurrent users
 - Also background utility processing must be done

Example multithreaded programs (con't)


- Network servers
 - Concurrent requests from network
 - Again, single program, multiple concurrent operations
 - File server, Web server, and airline reservation systems
- Parallel programming (more than one physical CPU)
 - Split program into multiple threads for parallelism
 - This is called Multiprocessing
- Some multiprocessors are actually uniprogrammed:
 - Multiple threads in one address space but one program at a time

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.29

A Typical Use Case



2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.30

Kernel Use Cases

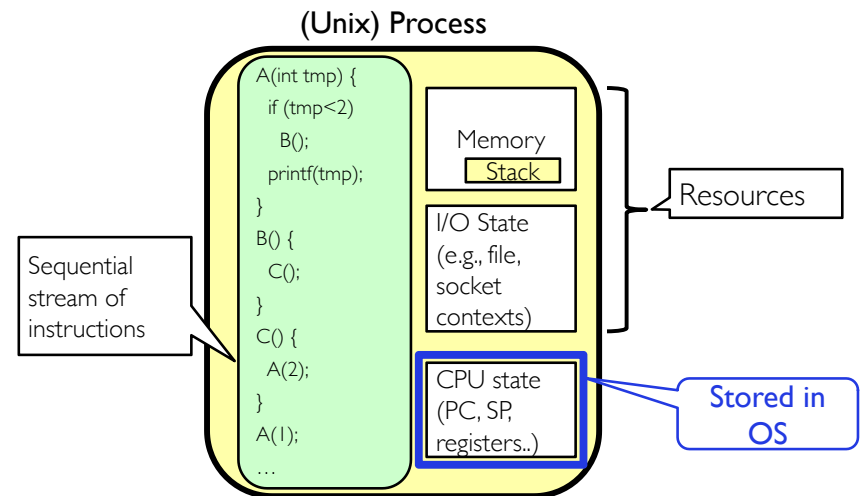
- Thread for each user process
- Thread for sequence of steps in processing I/O
- Threads for device drivers
- ...

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.31

Putting it Together: Process

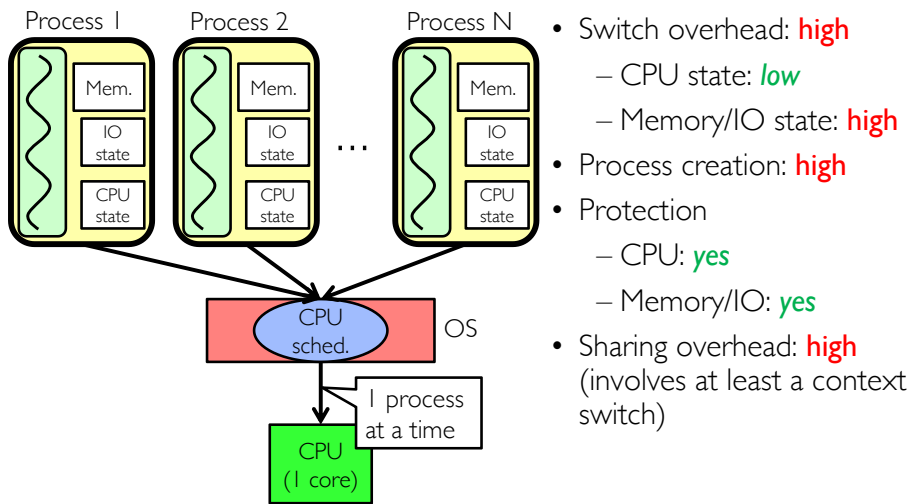


2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.32

Putting it Together: Processes

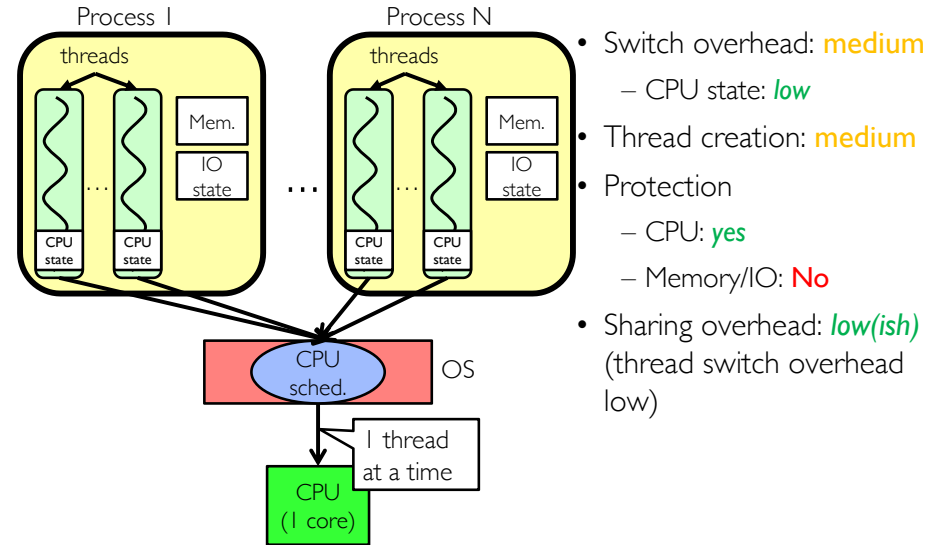


2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.33

Putting it Together: Threads



2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.34

Kernel versus User-Mode Threads

- We have been talking about kernel threads
 - Native threads supported directly by the kernel
 - Every thread can run or block independently
 - One process may have several threads waiting on different things
- Downside of kernel threads: a bit expensive
 - Need to make a crossing into kernel mode to schedule
- Lighter weight option: User Threads


2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.35

User-Mode Threads

- Lighter weight option:
 - User program provides scheduler and thread package
 - May have several user threads per kernel thread
 - User threads may be scheduled non-preemptively relative to each other (only switch on yield())
 - Cheap
- Downside of user threads:
 - When one thread blocks on I/O, all threads block
 - Kernel cannot adjust scheduling among all threads
 - Option: *Scheduler Activations*
 - » Have kernel inform user level when thread blocks...





2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018


Lec 6.36

Some Threading Models

Simple One-to-One Threading Model



Many-to-One



Many-to-Many

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.37

Threads in a Process


- Threads are useful at user-level: parallelism, hide I/O latency, interactivity
- Option A (early Java): user-level library, within a single-threaded process
 - Library does thread context switch
 - Kernel time slices between processes, e.g., on system call I/O
- Option B (SunOS, Linux/Unix variants): green threads
 - User-level library does thread multiplexing
- Option C (Windows): scheduler activations
 - Kernel allocates processors to user-level library
 - Thread library implements context switch
 - System call I/O that blocks triggers upcall
- Option D (Linux, MacOS, Windows): use kernel threads
 - System calls for thread fork, join, exit (and lock, unlock,...)
 - Kernel does context switching
 - Simple, but a lot of transitions between user and kernel mode

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.38

Putting it Together: Multi-Cores



- Switch overhead: **low** (only CPU state)
- Thread creation: **low**
- Protection
 - CPU: **yes**
 - Memory/IO: **No**
- Sharing overhead: **low** (thread switch overhead low, may not need to switch at all!)


2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.39

Recall: Simultaneous MultiThreading/Hyperthreading

- Hardware technique
 - Superscalar processors can execute multiple instructions that are independent
 - Hyperthreading duplicates register state to make a second “thread,” allowing more instructions to run
- Can schedule each thread as if were separate CPU
 - But, sub-linear speedup!
- Original called “Simultaneous Multithreading”
 - <http://www.cs.washington.edu/research/smt/index.html>
 - Intel, SPARC, Power (IBM)
 - A virtual core on AWS' EC2 is basically a hyperthread



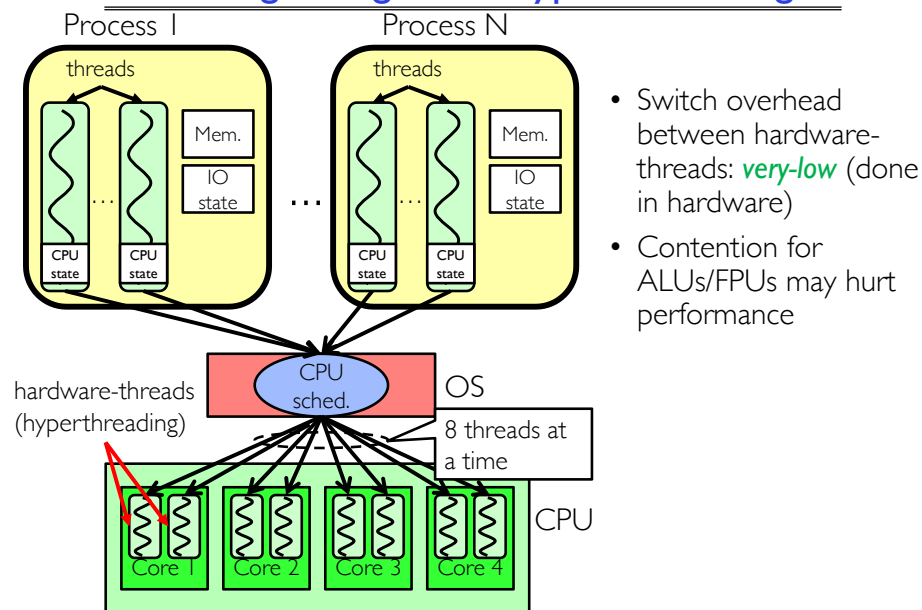
Colored blocks show instructions executed

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.40

Putting it Together: Hyper-Threading



2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.41

Classification

# threads Per AS:	# of addr spaces:	One	Many
One		MS/DOS, early Macintosh	Traditional UNIX
Many		Embedded systems (Geoworks, VxWorks, JavaOS, etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 10 Win NT to XP, Solaris, HP- UX, OS X

- Most operating systems have either
 - One or many address spaces
 - One or many threads per address space

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.42

Summary

- Processes have two parts
 - Threads (Concurrency)
 - Address Spaces (Protection)
- Various textbooks talk about **processes**
 - When this concerns concurrency, really talking about thread portion of a process
 - When this concerns protection, talking about address space portion of a process
- Concurrent threads are a very useful abstraction
 - Allow transparent overlapping of computation and I/O
 - Allow use of parallel processing when available
- Concurrent threads introduce problems when accessing shared data
 - Programs must be insensitive to arbitrary interleavings
 - Without careful design, shared variables can become completely inconsistent

2/5/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 6.43