

CSI62

Operating Systems and Systems Programming

Lecture 8

Locks, Semaphores, Monitors

February 12th, 2018

Profs. Anthony D. Joseph and Jonathan Ragan-Kelley

<http://cs162.eecs.Berkeley.edu>

Review: Too Much Milk Solution #3

- Here is a possible two-note solution:

<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) {\\X	if (noNote A) {\\Y
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	remove note B;
}	
remove note A;	

- Does this work? **Yes**. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At **X**:
 - If no note B, safe for A to buy,
 - Otherwise wait to find out what will happen
- At **Y**:
 - If no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

2/12/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 8.2

Review: Solution #3 discussion

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - » Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - » Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - » This is called “busy-waiting”
- There's a better way
 - Have hardware provide higher-level primitives than atomic load & store
 - Build even higher-level programming abstractions on this hardware support

2/12/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 8.3

Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock
 - **lock.Acquire()** – wait until lock is free, then grab
 - **lock.Release()** – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
milklock.Acquire();  
if (nomilk)  
    buy milk;  
milklock.Release();
```
- Once again, section of code between **Acquire()** and **Release()** called a “**Critical Section**”
- Of course, you can make this even simpler: suppose you are out of ice cream instead of milk
 - Skip the test since you always need more ice cream ;-)

2/12/18

Joseph and Ragan-Kelley CSI62 © UCB Spring 2018

Lec 8.4

Where are we going with synchronization?

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors Send/Receive
Hardware	Load/Store Disable Ints Test&Set Compare&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
 - Everything is pretty painful if only atomic primitives are load and store
 - Need to provide primitives useful at user-level

Goals for Today

- Explore several implementations of locks
- Continue with Synchronization Abstractions
 - Semaphores, Monitors, and Condition variables
- Very Quick Introduction to scheduling

How to Implement Locks?

- **Lock:** prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - » Important idea: all synchronization involves waiting
 - » Should sleep if waiting for a long time
- Atomic Load/Store: get solution like Milk #3
 - Pretty complex and error prone
- Hardware Lock instruction
 - Is this a good idea?
 - What about putting a task to sleep?
 - » How do you handle the interface between the hardware and scheduler?
 - Complexity?
 - » Done in the Intel 432 – each feature makes HW more complex and slow



Naïve use of Interrupt Enable/Disable

How can we build multi-instruction atomic operations?

- Recall: dispatcher gets control in two ways.
 - Internal: Thread does something to relinquish the CPU
 - External: Interrupts cause dispatcher to take CPU
- On a uniprocessor, can avoid context-switching by:
 - Avoiding internal events (although virtual memory tricky)
 - Preventing external events by disabling interrupts

Consequently, naïve Implementation of locks:

```
LockAcquire { disable Ints; }  
LockRelease { enable Ints; }
```

Naïve use of Interrupt Enable/Disable: Problems

Can't let user do this! Consider following:

```
LockAcquire();  
While(TRUE) {;
```

Real-Time system—no guarantees on timing!

- Critical Sections might be arbitrarily long

What happens with I/O or other important events?

- “Reactor about to meltdown. Help?”



Better Implementation of Locks by Disabling Interrupts

Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int value = FREE;
```



```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}  
  
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

New Lock Implementation: Discussion

- Why do we need to disable interrupts at all?
 - Avoid interruption between checking and setting lock value
 - Otherwise two threads could think that they both have lock

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

} Critical Section

- Note: unlike previous solution, the critical section (inside **Acquire()**) is very short
 - User of lock can take as long as they like in their own critical section: doesn't impact global machine behavior
 - Critical interrupts taken in time!

Interrupt Re-enable in Going to Sleep


- What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Enable Position →
Enable Position →
Enable Position →

How to Re-enable After Sleep()?

- In scheduler, since interrupts are disabled when you call sleep:
 - Responsibility of the next thread to re-enable ints
 - When the sleeping thread wakes up, returns to acquire and re-enables interrupts



2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.13

Atomic Read-Modify-Write Instructions

- Problems with previous solution:
 - Can't give lock implementation to users
 - Doesn't work well on multiprocessor
 - Disabling interrupts on all processors requires messages and would be very time consuming
- Alternative: **atomic instruction sequences**
 - These instructions read a value and write a new value atomically
 - Hardware is responsible for implementing this correctly
 - on both uniprocessors (not too hard)
 - and multiprocessors (requires help from cache coherence protocol)
 - Unlike disabling interrupts, can be used on both uniprocessors and multiprocessors

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.14

Examples of Read-Modify-Write

- `test&set (&address) {` `/* most architectures */`
 `result = M[address];` `/* return result from "address" and`
 `M[address] = 1;` `set value at "address" to 1 */`
 `return result;`
}
- `swap (&address, register) {` `/* x86 */`
 `temp = M[address];` `/* swap register's value to`
 `M[address] = register;` `value at "address" */`
 `register = temp;`
}
- `compare&swap (&address, reg1, reg2) {` `/* 68000 */`
 `if (reg1 == M[address]) {`
 `M[address] = reg2;`
 `return success;`
 }
 `else {`
 `return failure;`
 }
}

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.15

Implementing Locks with test&set

- Another flawed, but simple solution:
`int value = 0; // Free`
`Acquire() {`
 `while (test&set(value)); // while busy`
}
- `Release() {`
 `value = 0;`
}
- Simple explanation:
 - If lock is free, test&set reads 0 and sets value=1, so lock is now busy
It returns 0 so while exits
 - If lock is busy, test&set reads 1 and sets value=1 (no change)
It returns 1, so while loop continues
 - When we set value = 0, someone else can get lock
- Busy-Waiting:** thread consumes cycles while waiting

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.16

Problem: Busy-Waiting for Lock

- Positives for this solution
 - Machine can receive interrupts
 - User code can use this lock
 - Works on a multiprocessor
- Negatives
 - This is very inefficient as thread will consume cycles waiting
 - Waiting thread may take cycles away from thread holding lock (no one wins!)
 - Priority Inversion**: If busy-waiting thread has higher priority than thread holding lock \Rightarrow no progress!
- Priority Inversion problem with original Martian rover
- For semaphores and monitors, waiting thread may wait for an arbitrary long time!
 - Thus even if busy-waiting was OK for locks, definitely not ok for other primitives
 - Homework/exam solutions should avoid busy-waiting!



2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.17

Better Locks using test&set

- Can we build test&set locks without busy-waiting?
 - Can't entirely, but can minimize!
 - Idea: only busy-wait to atomically check lock value

```
int guard = 0;
int value = FREE;
```



```
Acquire() {
    // Short busy-wait time
    while (test&set(guard));
    if (value == BUSY) {
        put thread on wait queue;
        go to sleep() & guard = 0;
    } else {
        value = BUSY;
        guard = 0;
    }
}

Release() {
    // Short busy-wait time
    while (test&set(guard));
    if anyone on wait queue {
        take thread off wait queue
        Place on ready queue;
    } else {
        value = FREE;
    }
    guard = 0;
}
```

- Note: sleep has to be sure to reset the guard variable
 - Why can't we do it just before or just after the sleep?

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.18

Locks using Interrupts vs. test&set

Compare to "disable interrupt" solution

```
int value = FREE;
```



```
Acquire() {
    disable interrupts;
    if (value == BUSY) {
        put thread on wait queue;
        Go to sleep();
        // Enable interrupts?
    } else {
        value = BUSY;
    }
    enable interrupts;
}

Release() {
    disable interrupts;
    if (anyone on wait queue) {
        take thread off wait queue
        Place on ready queue;
    } else {
        value = FREE;
    }
    enable interrupts;
}
```

Basically replace

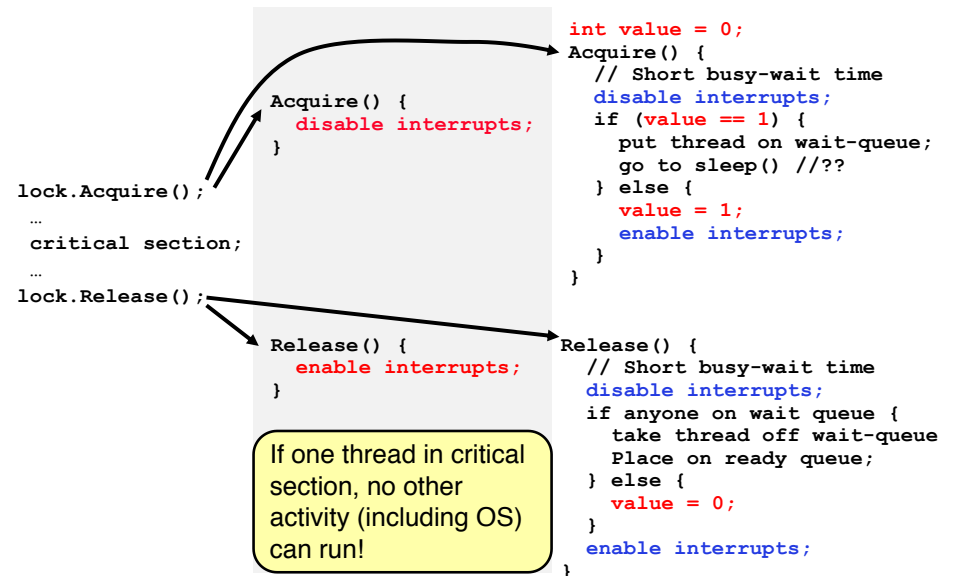
- disable interrupts** \rightarrow **while (test&set(guard));**
- enable interrupts** \rightarrow **guard = 0;**

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.19

Recap: Locks using interrupts

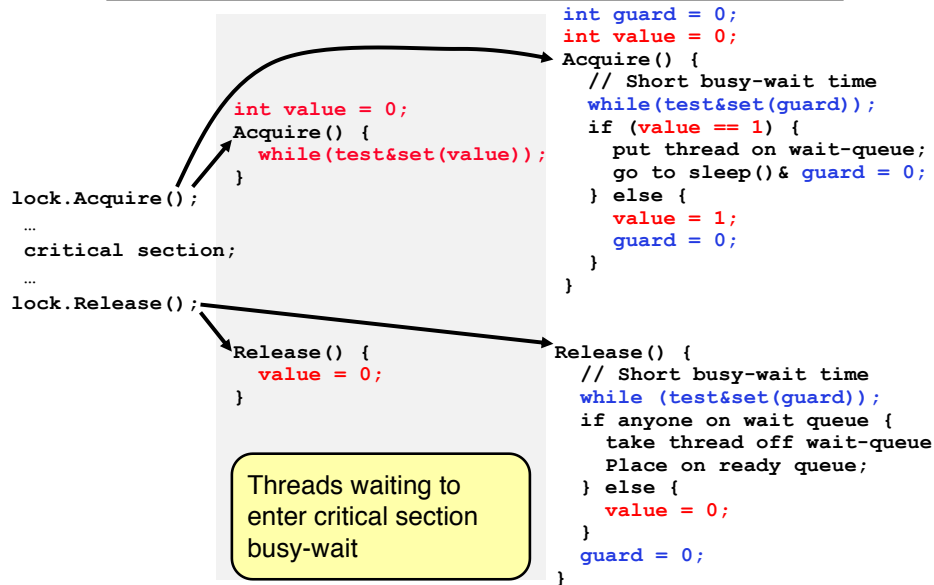


2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.20

Recap: Locks using test & wait



2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.21

Administrivia

- Midterm I **Wednesday 2/28 6:30-8:30PM**
- Homework I due **today 11:59PM**
- Project I Design Document due **Wednesday 2/14 11:59PM**
- Project I Design reviews upcoming
 - High-level discussion of your approach
 - » What will you modify?
 - » What algorithm will you use?
 - » How will things be linked together, etc.
 - » Do not need final design (complete with all semicolons!)
 - You *will* be asked about testing
 - » Understand testing framework
 - » Are there things you are doing that are not tested by tests we give you?
- *Do your own work!*
 - Please do not try to find solutions from previous terms
 - We will be on the look out for anyone doing this...today

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.22

BREAK

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.23

Higher-level Primitives than Locks

- Goal of last couple of lectures:
 - What is right abstraction for synchronizing threads that share memory?
 - Want as high a level primitive as possible
- Good primitives and practices important!
 - Since execution is not entirely sequential, really hard to find bugs, since they happen rarely
 - UNIX is pretty stable now, but up until about mid-80s (10 years after started), systems running UNIX would crash every week or so – concurrency bugs
- Synchronization is a way of coordinating multiple concurrent activities that are using shared state
 - This lecture and the next presents a some ways of structuring sharing

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.24

Semaphores



- Semaphores are a kind of generalized lock
 - First defined by Dijkstra in late 60s
 - Main synchronization primitive used in original UNIX
- Definition: a Semaphore has a non-negative integer value and supports the following two operations:
 - **P()**: an atomic operation that waits for semaphore to become positive, then decrements it by 1
 - » Think of this as the wait() operation
 - **V()**: an atomic operation that increments the semaphore by 1, waking up a waiting P, if any
 - » Think of this as the signal() operation
 - Note that **P()** stands for “proberen” (to test) and **V()** stands for “verhogen” (to increment) in Dutch


2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.25

Semaphores Like Integers Except

- Semaphores are like integers, except
 - No negative values
 - Only operations allowed are P and V – can't read or write value, except to set it initially
 - Operations must be atomic
 - » Two P's together can't decrement value below zero
 - » Similarly, thread going to sleep in P won't miss wakeup from V – even if they both happen at same time
- Semaphore from railway analogy
 - Here is a semaphore initialized to 2 for resource control:



2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.26

Two Uses of Semaphores

Mutual Exclusion (initial value = 1)

- Also called “Binary Semaphore”.
- Can be used for mutual exclusion:

```
semaphore.P();  
// Critical section goes here  
semaphore.V();
```

Scheduling Constraints (initial value = 0)

- Allow thread 1 to wait for a signal from thread 2, i.e., thread 2 **schedules** thread 1 when a given **event** occurs
- Example: suppose you had to implement ThreadJoin which must wait for thread to terminate:


```
Initial value of semaphore = 0  
ThreadJoin {  
    semaphore.P();  
}  
ThreadFinish {  
    semaphore.V();  
}
```

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.27

Producer-Consumer with a Bounded Buffer



- Problem Definition
 - Producer puts things into a shared buffer
 - Consumer takes them out
 - Need synchronization to coordinate producer/consumer
- Don't want producer and consumer to have to work in lockstep, so put a fixed-size buffer between them
 - Need to synchronize access to this buffer
 - Producer needs to wait if buffer is full
 - Consumer needs to wait if buffer is empty
- Example 1: GCC compiler
 - **cpp | cc1 | cc2 | as | ld**
- Example 2: Coke machine
 - Producer can put limited number of Cokes in machine
 - Consumer can't take Cokes out if machine is empty



2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.28

Correctness constraints for solution

- Correctness Constraints:
 - Consumer must wait for producer to fill buffers, if none full (scheduling constraint)
 - Producer must wait for consumer to empty buffers, if all full (scheduling constraint)
 - Only one thread can manipulate buffer queue at a time (mutual exclusion)
- Remember why we need mutual exclusion
 - Because computers are stupid
 - Imagine if in real life: the delivery person is filling the machine and somebody comes up and tries to stick their money into the machine
- General rule of thumb:
Use a separate semaphore for each constraint
 - Semaphore fullBuffers; // consumer's constraint
 - Semaphore emptyBuffers; // producer's constraint
 - Semaphore mutex; // mutual exclusion

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.29

Full Solution to Bounded Buffer

```
Semaphore fullSlots = 0; // Initially, no coke
Semaphore emptySlots = bufSize; // Initially, num empty slots
Semaphore mutex = 1; // No one using machine

Producer(item) {
    emptySlots.P(); // Wait until space
    mutex.P(); // Wait until machine free
    Enqueue(item);
    mutex.V(); // Tell consumers there is
    fullSlots.V(); // more coke
}

Consumer() {
    fullSlots.P(); // Check if there's a coke
    mutex.P(); // Wait until machine free
    item = Dequeue();
    mutex.V(); // tell producer need more
    emptySlots.V();
    return item;
}
```

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.30

Discussion about Solution

Why asymmetry?

- Producer does: emptySlots.P(), fullSlots.V()
- Consumer does: fullSlots.P(), emptySlots.V()

Decrease # of
empty slots

Increase # of
occupied slots

Decrease # of
occupied slots

Increase # of
empty slots

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.31

Discussion about Solution (cont'd)

Is order of P's important?

Is order of V's important?

What if we have 2 producers or 2 consumers?

```
Producer(item) {
    mutex.P();
    emptySlots.P();
    Enqueue(item);
    mutex.V();
    fullSlots.V();
}

Consumer() {
    fullSlots.P();
    mutex.P();
    item = Dequeue();
    mutex.V();
    emptySlots.V();
    return item;
}
```

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.32

Motivation for Monitors and Condition Variables


- Semaphores are a huge step up; just think of trying to do the bounded buffer with only loads and stores
 - Problem is that semaphores are dual purpose:
 - » They are used for both mutex and scheduling constraints
 - » Example: the fact that flipping of P's in bounded buffer gives deadlock is not immediately obvious. How do you prove correctness to someone?
- Cleaner idea: Use *locks* for mutual exclusion and *condition variables* for scheduling constraints
- Definition: **Monitor**: a **lock** and zero or more **condition variables** for managing concurrent access to shared data
 - Some languages like Java provide this natively
 - Most others use actual locks and condition variables

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.33

Monitor with Condition Variables



- Lock**: the lock provides mutual exclusion to shared data
 - Always acquire before accessing shared data structure
 - Always release after finishing with shared data
 - Lock initially free
- Condition Variable**: a queue of threads waiting for something *inside* a critical section
 - Key idea: make it possible to go to sleep inside critical section by atomically releasing lock at time we go to sleep
 - Contrast to semaphores: Can't wait inside critical section

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.34

Simple Monitor Example (version 1)

- Here is an (infinite) synchronized queue

```
Lock lock;
Queue queue;

AddToQueue(item) {
    lock.Acquire();           // Lock shared data
    queue.enqueue(item);      // Add item
    lock.Release();           // Release Lock
}

RemoveFromQueue() {
    lock.Acquire();           // Lock shared data
    item = queue.dequeue();    // Get next item or null
    lock.Release();           // Release Lock
    return(item);             // Might return null
}
```
- Not very interesting use of "Monitor"
 - It only uses a lock with no condition variables
 - Cannot put consumer to sleep if no work!

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.35

Condition Variables

- How do we change the RemoveFromQueue() routine to wait until something is on the queue?
 - Could do this by keeping a count of the number of things on the queue (with semaphores), but error prone
- Condition Variable**: a queue of threads waiting for something *inside* a critical section
 - Key idea: allow sleeping inside critical section by atomically releasing lock at time we go to sleep
 - Contrast to semaphores: Can't wait inside critical section
- Operations:
 - Wait(&lock)**: Atomically release lock and go to sleep. Re-acquire lock later, before returning.
 - Signal()**: Wake up one waiter, if any
 - Broadcast()**: Wake up all waiters
- Rule: Must hold lock when doing condition variable ops!
 - In Birrell paper, he says can perform signal() outside of lock – IGNORE HIM (this is only an optimization)

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.36

Complete Monitor Example (with condition variable)

- Here is an (infinite) synchronized queue

```
Lock lock;
Condition dataready;
Queue queue;

AddToQueue(item) {
    lock.Acquire();           // Get Lock
    queue.enqueue(item);      // Add item
    dataready.signal();       // Signal any waiters
    lock.Release();           // Release Lock
}

RemoveFromQueue() {
    lock.Acquire();           // Get Lock
    while (queue.isEmpty()) { // If nothing, sleep
        dataready.wait(&lock);
    }
    item = queue.dequeue();   // Get next item
    lock.Release();           // Release Lock
    return(item);
}
```

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.37

Summary (1/2)

- Important concept: **Atomic Operations**
 - An operation that runs to completion or not at all
 - These are the primitives on which to construct various synchronization primitives
- Talked about hardware atomicity primitives:
 - Disabling of Interrupts, test&set, swap, compare&swap, conditional
- Showed several constructions of Locks
 - Must be very careful not to waste/tie up machine resources
 - » Shouldn't disable interrupts for long
 - » Shouldn't spin wait for long
 - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.38

Summary (2/2)

- Semaphores**: Like integers with restricted interface
 - Two operations:
 - » **P()**: Wait if zero; decrement when becomes non-zero
 - » **V()**: Increment and wake a sleeping task (if exists)
 - » Can initialize value to any non-negative value
 - Use separate semaphore for each constraint
- Monitors**: A lock plus one or more condition variables
 - Always acquire lock before accessing shared data
 - Use condition variables to wait inside critical section
 - » Three Operations: **Wait()**, **Signal()**, and **Broadcast()**

2/12/18

Joseph and Ragan-Kelley CS162 © UCB Spring 2018

Lec 8.39